

AWS CDK - Best Practices From The Trenches | Ran Isenberg | Conf42 Cloud Native 2023



Conf42

3.29K subscribers

Subscribe



Like



612 views Premiered Apr 30, 2023 [Conf42 Cloud Native 2023](#)



Ran The Builder
Build Serverless Services

AWS CDK – BEST PRACTICES FROM THE TRENCHES

RAN ISENBERG, PRINCIPAL SOFTWARE ARCHITECT



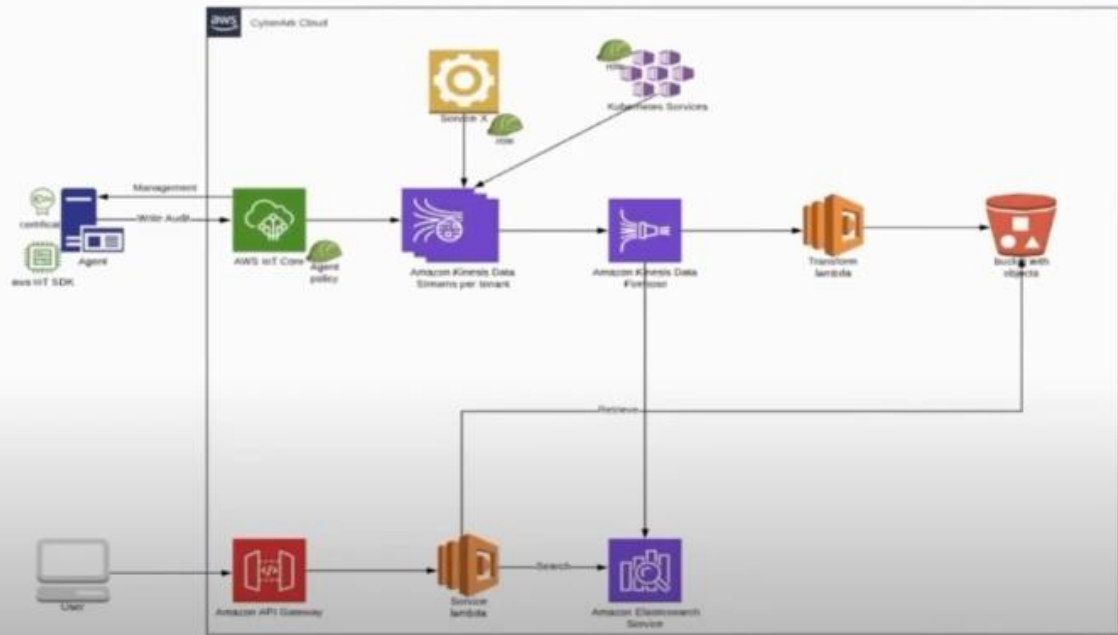
CYBERARK

"AWS CDK LETS YOU BUILD
RELIABLE, SCALABLE, COST-
EFFECTIVE APPLICATIONS IN THE
CLOUD WITH THE
CONSIDERABLE EXPRESSIVE
POWER OF A PROGRAMMING
LANGUAGE" – AWS DOCS



#CDKDay

USE CASE – AUDIT SERVICE POC



This POC is an ETL log service that transforms the streaming data with a Lambda before saving it off into a database.



With great power comes great responsibility.

Uncle ben, 2002, Spiderman

AGENDA

- CDK App guidelines
- Constructs guidelines
- CI/CD guidelines
- Security
- Resilience
- General development tips
- Summary

Quality Management
Smart Expectations
Legal Results
BEST PRACTICES
Requirements
Method Complying
Self-Assessment Balance
Transparent Organizations
Documenting Strategic
Data-Driven

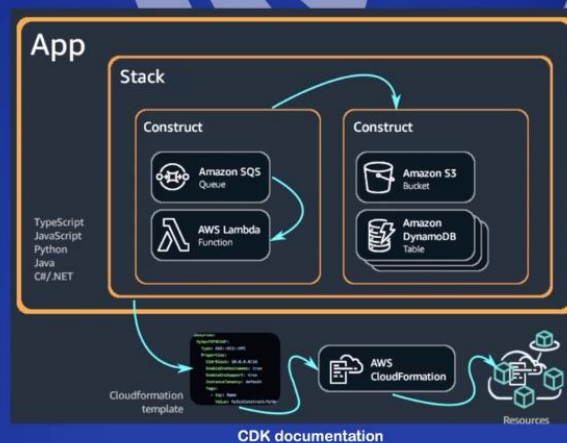
- Principal Software Architect @CyberArk
- AWS Community Builder
- Owner & Blogger @RanTheBuilder.Cloud



BEST PRACTICES

CDK APP GUIDELINES

- One business domain
- One repository & CI/CD pipeline
- Maintained by one team
- One CDK application & stack
- Small blast radius




CDK APP Guidelines

- When to split to a new application & repository:
 1. Different team will maintain the new application
 2. Different business domain
- Don't over split! Balance is key
- Multiple repositories:
 - Increase development complexity of new cross repo features
 - Share deployment time parameters (SSM/CloudMap)

PROJECT STRUCTURE GUIDELINES

- IaC and business domain code together
- One CI/CD pipeline
- Tests:
 - Unit / integration / e2e
 - Security & CDK infra




It is recommended to have 3 folders, a CDK folder, a main service folder for the business domain logic as Lambdas, and Tests folder

CDK Template Project



- Self service
- Internal training
- Reduce cognitive load
- Jump start into SaaS development
- Organization level: same tools, CI/CD pipeline, tech stack



- <https://github.com/ran-isenberg/aws-lambda-handler-cookbook>

You can create a CDK Template project that teams can use in a self-service manner to get all the common things having the best practice patterns and they can just start writing the main business logic code.

A green road sign with a white border and the text 'BEST PRACTICES' in white capital letters.

CONSTRUCT GUIDELINES

Stack/Construct Composition

- Don't define all resources in the stack
 - Use constructs
 - Exception - Lambda layer used in multiple constructs
- Constructs are easy to share

Shareable Constructs

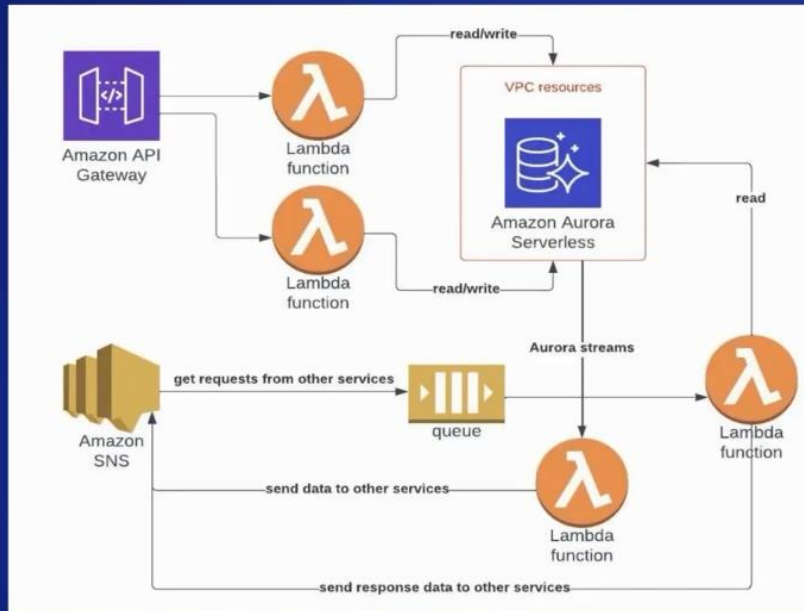
- Platform engineers own & maintain
- Pros:
 - Secure, cost effective, tested constructs
 - Save time for developers
- Cons:
 - Versioned
 - Can cause breakage/resource deletion on upgrades

We have a versioned Python library of CDK constructs that our developers can pick and choose from.

Shareable Constructs

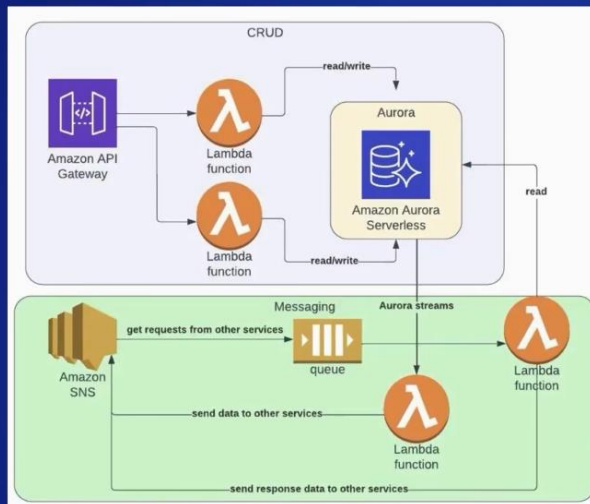
- Internal library of common constructs
 - WAF rules for API Gateway/CloudFront distributions.
 - SNS -> SQS pattern with encryption at REST
 - AWS AppConfig dynamic configuration construct.
 - Datadog logs shipper/log PII sanitizer
- External resources:
 - <https://constructs.dev>
 - Serverless land
 - cdkpatterns.com
 - <https://aws.amazon.com/solutions/constructs>

Business Domain Driven Constructs



How can we split this design into smaller separate constructs to enable sharing in future?

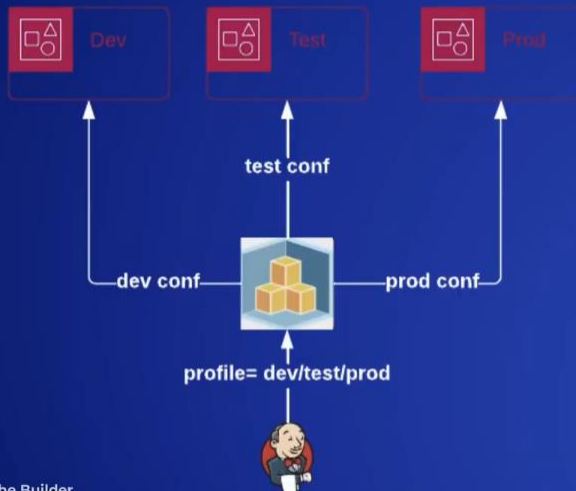
Business Domain Driven Constructs



BEST PRACTICES

CI/CD GUIDELINES

MODEL YOUR CI/CD PIPELINE STAGES IN CODE



We use Jenkins to set the environment variables and inject them into the CDK for different configurations and accounts

Model Your CI/CD Pipeline Stages in Code

- Why multiple accounts?
 - Account breach smaller blast radius
 - AWS resource quota limits
- How to model stages in CDK?
 - Environment variables
 - 'if statements' for the win
 - Apply different configuration

Model Your CI/CD Pipeline Stages in Code

```
profile = os.getenv('PROFILE')
table = dynamodb.Table(
    self,
    table_id,
    table_name=table_id,
    partition_key=dynamodb.Attribute(name='order_id', type=dynamodb.AttributeType.STRING),
    billing_mode=dynamodb.BillingMode.PAY_PER_REQUEST,
    point_in_time_recovery=False if profile == 'dev' else True,
    removal_policy=RemovalPolicy.DESTROY if profile == 'dev' else RemovalPolicy.RETAIN,
)
```

BEST PRACTICES

SECURITY GUIDELINES

Secrets in CDK

- NEVER write hardcoded secrets in CDK or config files
- Store as GitHub/Jenkins/pipeline secret
 - Inject to CDK as environment variable/parameter during deploy
- Deploy secrets:
 - AWS Secrets Manager
 - SSM parameter store encrypted string
- Consume in Lambda from SSM/Secrets manager:
 - Secret name as lambda env. variable

Some Security Defaults Are Not Good Enough

AWS News Blog

Amazon S3 Encrypts New Objects By Default

by Sébastien Stormacq | on 05 JAN 2023 | in Amazon Simple Storage Service (S3),

DynamoDB encryption at rest

PDF | RSS

All user data stored in Amazon DynamoDB is fully encrypted at rest.

- What about SNS encryption at rest?
 - Disabled by default
- Security defaults differ by service
- AWS sets better defaults over time

Some Security Defaults Are Not Good Enough

- Shared responsibility model
- Don't expect AWS to do all the work for you
- Enable security best practices for all resources
- Security Review, scheduled PT
- Run CDK security tests – CDK nag

AWS CDK Security Tests

```
from aws_cdk import App, Aspects
from cdk_nag import AwsSolutionsChecks, HIPAASecurityChecks
from cdk.my_service.service_stack import ServiceStack

def test_cdk_nag_default():
    app = App()

    service_stack = ServiceStack(app, 'service-test')
    Aspects.of(service_stack).add(AwsSolutionsChecks(verbose=True))

def test_cdk_nag_hipaa():
    app = App()

    service_stack = ServiceStack(app, 'service-test')
    Aspects.of(service_stack).add(HIPAASecurityChecks(verbose=True))
```

This tool is called **CDK Nag** for CDK security tests before push/deploy

Write Your Own IAM Policies

```
def _build_db(self, id_prefix: str, my_role: iam.Role) -> dynamodb.Table:
    table_id = f'{id_prefix}{constants.TABLE_NAME}'
    table = dynamodb.Table(
        self,
        table_id,
        table_name=table_id,
        partition_key=dynamodb.Attribute(name='order_id', type=dynamodb.AttributeType.STRING),
        billing_mode=dynamodb.BillingMode.PAY_PER_REQUEST,
        point_in_time_recovery=True,
        removal_policy=RemovalPolicy.DESTROY,
    )
    table.grant_read_write_data(my_role)
    return table
```

Grants: BatchGetItem, GetRecords, GetShardIterator, Query, GetItem, Scan, BatchWriteItem, PutItem, UpdateItem, DeleteItem, DescribeTable

You always want to assign the least amount of privileges using the CDK to write your own IAM rules and policies

Write Your Own IAM Policies

```
def _build_lambda_role(self, db: dynamodb.Table) -> iam.Role:
    return iam.Role(
        self,
        'ServiceRole',
        assumed_by=iam.ServicePrincipal('lambda.amazonaws.com'),
        inline_policies={
            'dynamodb_db':
                iam.PolicyDocument(statements=[
                    iam.PolicyStatement(actions=['dynamodb:PutItem', 'dynamodb:GetItem'], resources=[db.table_arn],
                                         effect=iam.Effect.ALLOW)
                ])
    })
```

- Grants only GetItem, PutItem
- Prefer least privilege method – assign only what you need, no more, no less
- Better developers understand IAM policies



BEST PRACTICES

RESILIENCE GUIDELINES

Changing Logical ID is Dangerous

- **Unique resource ID**
- **Innocent refactor can be hazardous:**
 - Stateful logical ids must NEVER change
 - Cross account trust role can break
- **Critical resources can get deleted due to bugs**
- **Write CDK unit tests**

Changing logical IDs mean the CDK will delete that resource and recreate it, this is not what we want when redeploying

CDK Unit Tests

```
from aws_cdk import App
from aws_cdk.assertions import Template
from cdk.my_service.service_stack import ServiceStack

def test_synthesizes_properly():
    app = App()

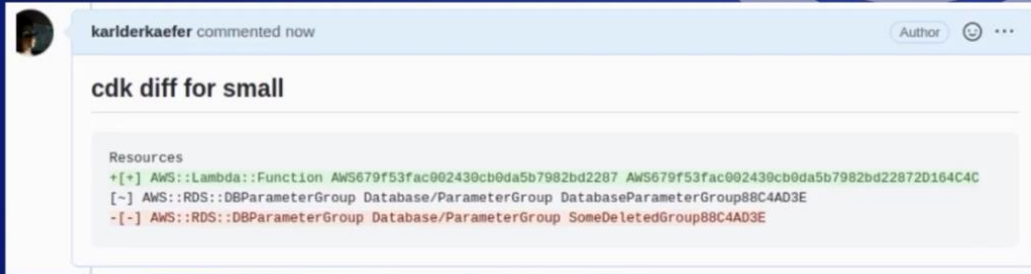
    service_stack = ServiceStack(app, 'service-test')

    # Prepare the stack for assertions.
    template = Template.from_stack(service_stack)

    # verify that we have one API GW, that is it not deleted by mistake
    template.resource_count_is('AWS::ApiGateway::RestApi', 1)
    table = template.find_resources('AWS::DynamoDB::Table')
    # assert table's key matches the logical id
```

The CDK unit tests should run before push/deploy of your code.

CHANGES VISIBILITY



This **cdk diff** tool is OSS and allows you to see the changes made to the code before push/deploy

Backups

- Retain policy – RETAIN in production
 - Restore vs. lose customer data forever
- Backup your stateful resources:
 - DynamoDB point in time
 - AWS Backup

GENERAL DEVELOPMENT GUIDELINES

General Development Tips

- Console first approach
- CFN low level FTW
- Tag it!
- CDK code maintainability > abstraction
 - Avoid “cool” factory methods
 - Keep it simple
 - IaC must be readable and easy to follow

CFN is Cloud Formation and we can go lower than the CDK to implement specific functionalities

Summary

- With great power comes great responsibility
- Shared responsibility model
- Enforce best practices in organization:
 - CDK App, stack & construct guidelines
 - Share constructs
 - CDK Template self service
 - Security
 - Resilience



Ran The Builder
Build Serverless Services



CYBERARK

THANK YOU!



@RANISENBERG



[HTTPS://WWW.RANTHEBUILDER.CLOUD](https://www.ranthebuilder.cloud)



@ISENBERGRAN