

DAT333

Building highly resilient applications with Amazon DynamoDB

Jeff Duffy

Product Manager,
Amazon DynamoDB
Amazon Web Services

Tom Skinner

Director,
Measurement Infrastructure
Amazon Ads

Richard Edwards III

Principal Software Engineer
Amazon Ads

AWS re:Invent 2023 - Building highly resilient applications with Amazon DynamoDB (DAT333)



AWS Events
100K subscribers

Subscribe



3



Share



Clip



Save



120 views Dec 3, 2023 #AWSreInvent #AWSreInvent2023

Join this session to explore how resiliency features of Amazon DynamoDB help you build scalable, reliable applications. Learn how to prepare for the unexpected with DynamoDB capabilities like redundant storage, automatic throughput scaling, and multi-active, multi-Region data replication to achieve your business continuity goals at scale. Additionally, Amazon Advertising shares why they chose to migrate to DynamoDB for their most critical workloads and discusses their technical and architectural approaches to achieving the highest level of resiliency.

Learn more about AWS re:Invent at <https://go.aws/46iuzGv>.

Subscribe:

More AWS videos: <http://bit.ly/203zs75>

More AWS events videos: <http://bit.ly/316g9t4>

ABOUT AWS

Amazon Web Services (AWS) hosts events, both online and in-person, bringing the cloud computing community together to connect, collaborate, and learn from AWS experts.

AWS is the world's most comprehensive and broadly adopted cloud platform, offering over 200 fully featured services from data centers globally. Millions of customers—including the fastest-growing startups, largest enterprises, and leading government agencies—are using AWS to lower costs, become more agile, and innovate faster.

Agenda

What resilience is

How DynamoDB helps you build for resilience

Why Amazon Ads chose DynamoDB

Deep dive on how Amazon Ads builds for resilience

Resilience

The ability to adjust to change

Infrastructure
failure

Varying
demand

System
modifications

Resilience components

Disaster recovery (DR)

Focuses on
entire workload

Respond **after** failure

High availability (HA)

Focuses on workload
components

Respond **during** failure

Resilience targets



Recovery point objective (RPO)

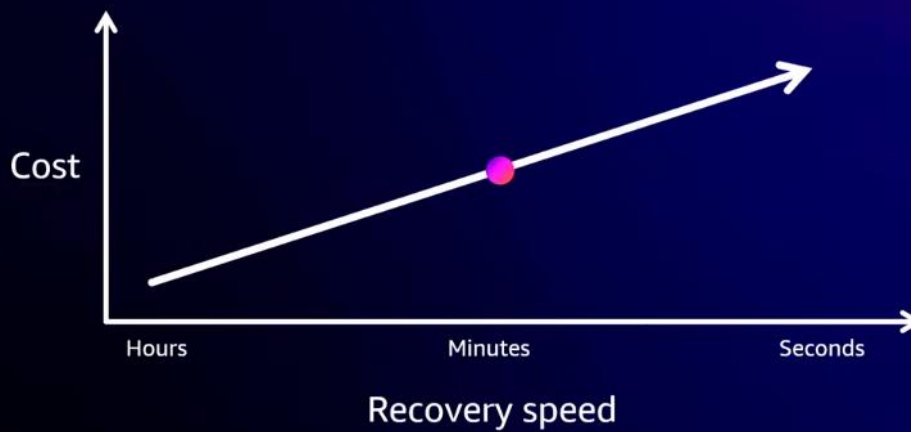
How much data are we
willing to lose or recreate?



Recovery time objective (RTO)

How quickly must we
return to operations?

Higher resilience usually means higher cost



Resilience strategies

Strategy	RPO	RTO	Cost
Backup and restore	Hours	Hours	Lowest
Pilot light	10s of minutes	10s of minutes	Lower
Warm standby	Single minutes	Single minutes	Higher
Active/active	Zero	Zero	Highest

We have 4 strategies in our well architected framework

AWS Global Infrastructure

32 Regions

102 Availability Zones

Regional API endpoints

Resilience shared responsibility model

Customer
Responsibility
for resilience **in**
the cloud

AWS
Responsibility
for resilience **of**
the cloud

Continuous testing of critical infrastructure

Workload architecture

Change management and
operational resilience

Observability and failure management

Network, quotas, and constraints

Hardware and services

Compute

Storage

Database

Networking

AWS Global Infrastructure

Regions

Availability Zones

Edge locations

DynamoDB resilience – Foundational




Serverless

3-AZ redundancy

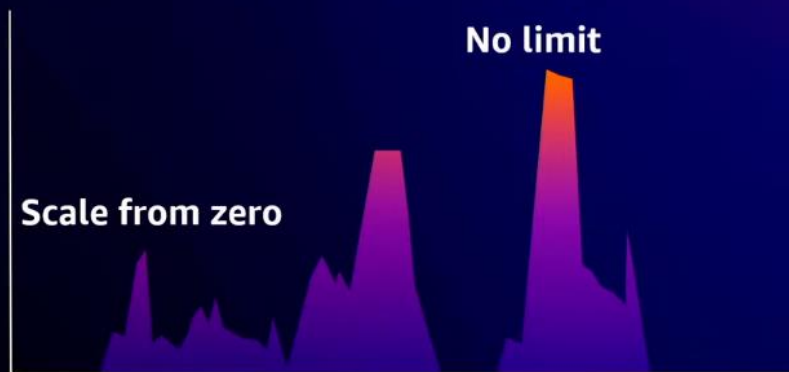
Zero-downtime updates

DynamoDB resilience – Capacity management



Auto Scaling for provisioned capacity mode

DynamoDB resilience – Capacity management



On-demand capacity mode

DynamoDB resilience – Recovery

Point-in-time recovery (PITR)



Scheduled backups



DynamoDB resilience – Global tables



Multi-active, multi-Region

99.999% availability

No failover required

DynamoDB multi-Region resilience strategies

Strategy	RPO	RTO	DynamoDB features
Backup and restore	Hours	Hours	Scheduled backups PITR
Pilot light	10s of minutes	10s of minutes	Global tables Minimum provisioned scaling or on-demand
Warm standby	Single minutes	Single minutes	Global tables Fully provisioned scaling or on-demand
Active/active	Zero	Zero	Multiple Region writes Fully provisioned scaling or on-demand

Ads measurement

Tom Skinner

Amazon Ads




Advertising analytics – Lifecycle



Attribution data flow



Workload demands



amazon ads



Ads measurement


4+PB
DDB

100B+
events/
day

90+MM
RCUs

5+MM
WCUs

Legacy architecture



Legacy architecture challenges



High operational load

Availability concerns and difficult disaster recovery

Inflexible scaling

Migration requirements



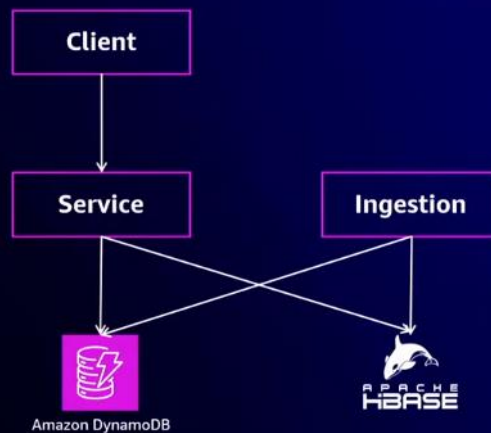
Seamless to our customers

Same or lower latency

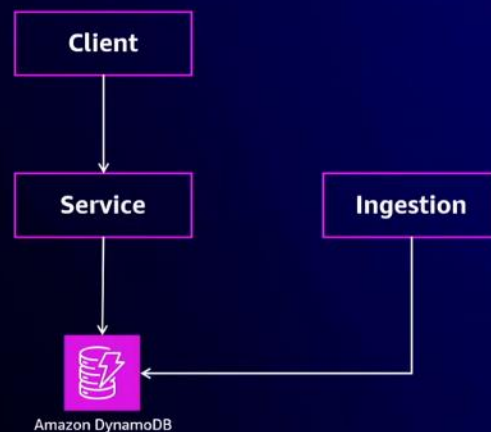
Managed solution

Less than a month

Migration dual write



Final architecture



Results



Service availability increased to five 9s (99.999%)


Developers productive in 2 weeks

Reduced ticket load by 40%

Cost-neutral


Principal Engineer Rich Edwards

The challenge: Internet-scale join



$$100\text{B+ events/day} \times 14 \text{ days} = 1.4 \text{ T+ events}$$

The challenge: Internet-scale join



$$100\text{B+ events/day} \times 14 \text{ days} = 1.4 \text{ T+ events}$$

Functional requirements



Low operational load

High availability (lower RTO)

Fast disaster recovery (no data loss)

Support all existing workloads

Our solution



HBase operational issues


Availability Zone redundancy

Disaster recovery


Query load balancing

Operational cluster maintenance

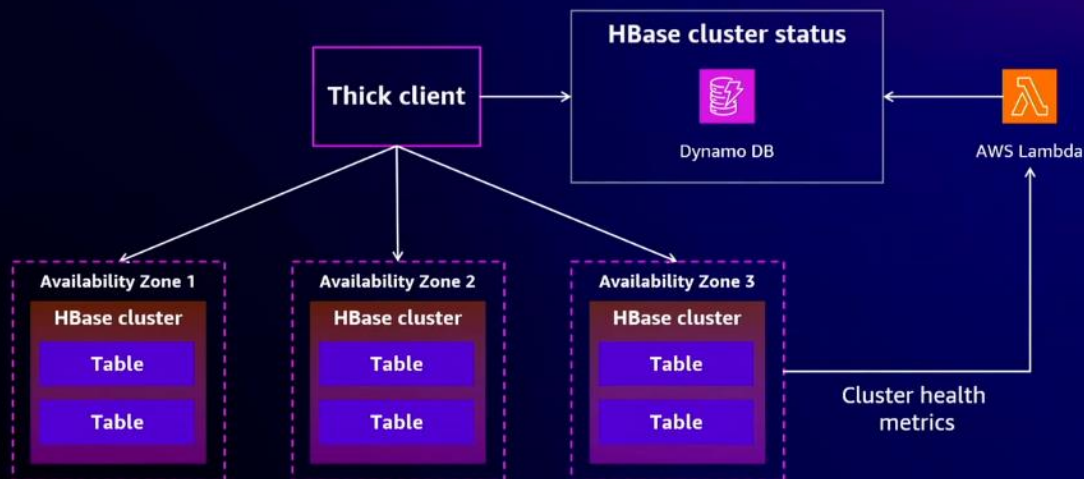
Availability Zone redundancy



Query load balancing



Cluster management



Our HBase pain points

Cluster ownership overhead


- Custom HBase builds
- Ramping SDEs on HBase and its internals
- Managing HBase replication
- Cluster maintenance



Our HBase pain points

Low elasticity

- Adding nodes is easy, removing nodes is challenging



DynamoDB solved our pain points

Fully managed

- No OS/security patching
- No redundancy management

Elastic

- Full auto scaling



Our DynamoDB design considerations

These all impact resiliency together

Table
structure

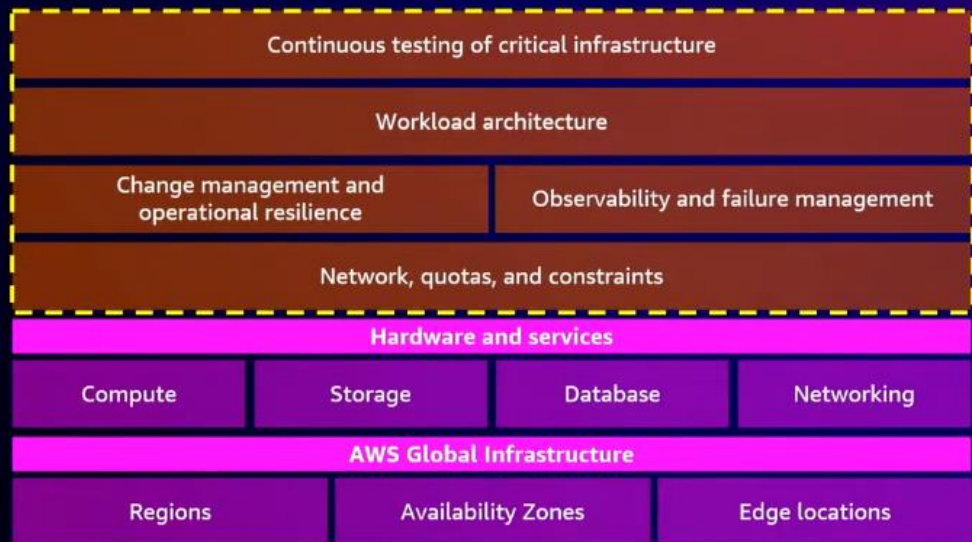
Throughput

Table
management


Resilience shared responsibility model

Customer
Responsibility
for resilience **in**
the cloud


AWS
Responsibility
for resilience **of**
the cloud



Customer responsibility



Customer responsibility



Our DynamoDB design considerations

Table
structure

Throughput

Table
management

Time series query use cases

Query all sorted ad interaction data (20% – workload)

Query relevant sorted ad interaction data (80% – workload)

Partition key: ID = 1, Sort key: = Attribute, Timestamp



Time series schema types

Fully sorted

Partially sorted (our HBase schema)

Scatter sorted

Table structure

HBase partially sorted table layout

Partition key	Sort key	Payload
ID	Attribute#timestamp	Blob

DynamoDB partially sorted table layout

Partition key	Sort key	Payload
ID	Attribute#timestamp	Blob

DynamoDB fully sorted table layout

Partition key	Sort key	Payload
ID	Timestamp#attribute	Blob

DynamoDB scatter sorted table layout

Partition key	Sort key	Payload
Id#attribute	Timestamp	Blob

Table structure – Partially sorted

Partially sorted table layout

Partition key	Sort key	Payload
id	attribute#timestamp	Blob

Partition key: ID = 1



Table structure – Query view

Query

id = 1,
attr = 1234



User

Partially sorted table layout

Partition key	Sort key	Payload
id	attribute#timestamp	Blob

Query result

id=1,
attr = 1234,
timestamp, blob

Table structure – Query view

Query

id = 1,
attr = 1234
... 345 ... N



User

Partially sorted table layout

Partition key	Sort key	Payload
id	attribute#timestamp	Blob

High throughput risk

Query result

id=1,
attr =1234,
timestamp, blob

Table structure – Fully sorted

Fully sorted table layout

Partition key	Sort key	Payload
id	timestamp#attribute	Blob

Partition Key: ID = 1

Oct 1st

Oct 31st

Attribute
1234

.....

Attribute
345

.....

Attribute
345

.....

Attribute
1234

Table structure – Query view

Query

id = 1
attr = 1234
... 345 ... N



User

DynamoDB fully sorted table layout

Partition key	Sort key	Payload
id	timestamp#attribute	Blob

Throughput risk

Query result

id=1,
timestamp, blob

Table structure – Scatter sorted

Scatter sorted table layout

Partition key	Sort key	Payload
id#attribute	timestamp	Blob

Partition Key: ID = 1, Attribute = 1234



Partition Key: ID = 1, Attribute = 345



Table structure – Query view

Query

id = 1,
attr = 1234
... 345 ... N



User

DynamoDB scatter sorted table layout

Partition key	Sort key	Payload
id#attribute	timestamp	Blob

No throughput risk

Query result

id=1,
attr = 1234,
timestamp, blob

Time series query use cases

Query all sorted ad interaction data (20% – workload)

Query relevant sorted ad interaction data (80% – workload)

Table structure – Scatter sorted selected

HBase partially sorted table layout

Partition key	Sort key	Payload
id	attribute#timestamp	Blob

~~DynamoDB partially sorted table layout~~

Partition key	Sort key	Payload
id	attribute#timestamp	Blob

~~DynamoDB fully sorted table layout~~

Partition key	Sort key	Payload
id	timestamp#attribute	Blob

DynamoDB scatter sorted table layout

Partition key	Sort key	Payload
id#attribute	timestamp	Blob

Table structure

DynamoDB scatter sorted table layout

Partition key	Sort key	Payload
id#attribute	timestamp	Blob

DynamoDB GSI fully sorted table layout

Partition key	Sort key	Payload
id	timestamp#attribute	Blob

Partition: ID = 1, Attribute = 1234



Partition: ID = 1, Attribute = 345



Table structure – Query amplification

Query

id = 1,
attr = 1234
... 345 ... N



User

DynamoDB scatter sorted table layout

Partition key	Sort key	Payload
id#attribute	timestamp	Blob

No DDB throughput risk

Query result

id=1,
attr = 1234,
timestamp, blob

Table structure – Query amplification

Query

id = 1,
attr = 1234
... 345 ... N



User

Application



Table structure – Query amplification



Table structure – Query amplification




Table structure – Fully sorted selected

HBase table layout		
Primary key	Sort key	Payload
id	attribute#timestamp	Blob

DynamoDB partially sorted table layout		
Primary key	Sort key	Payload
id	attribute#timestamp	Blob

DynamoDB fully sorted table layout		
Primary key	Sort key	Payload
id	timestamp#attribute	Blob

DynamoDB scatter sorted table layout		
Primary key	Sort key	Payload
id#attribute	timestamp	Blob

Our DynamoDB design considerations

Table
structure

Throughput

Table
management

Default throughput

DynamoDB supports per partition

3,000 RCUs (strongly
consistent)/6,000 RCUs
(eventually consistent)

1,000 WCUs

DynamoDB table

Partition key	Sort key	Payload
Id	timestamp#attr	Blob

Static replication throughput

This example supports per partition

9,000 RCUs (strongly
consistent)/18,000 RCUs
(eventually consistent)

1,000 WCUs

DynamoDB table

Partition key	Sort key	Payload
id#0	timestamp#attr	Blob
id#1	timestamp#attr	Blob
id#2	timestamp#attr	Blob

GSI replication throughput

This example supports per partition

9,000 RCUs (strongly consistent)/36,000 RCUs (eventually consistent)

1,000 WCUs

Primary table

Partition key	Sort key	Payload
id#0	timestamp#attribute	Blob
id#1	timestamp#attribute	Blob
id#2	timestamp#attribute	Blob

Read replica (GSI table)

Partition key	Sort key	Payload
id#0	timestamp#attribute	Blob
id#1	timestamp#attribute	Blob
id#2	timestamp#attribute	Blob

GT replication throughput

This example supports per partition

9,000 RCUs (strongly consistent)/36,000 RCUs (eventually consistent)

1,000 WCUs

Table US-EAST-1

Partition key	Sort key	Payload
id#0	timestamp#attribute	Blob
id#1	timestamp#attribute	Blob
id#2	timestamp#attribute	Blob

Table US-EAST-2

Partition key	Sort key	Payload
id#0	timestamp#attribute	Blob
id#1	timestamp#attribute	Blob
id#2	timestamp#attribute	Blob

Static + GT + GSI throughput

Table US-EAST-1

Partition key	Sort key	Payload
id#0	timestamp#attribute	Blob
id#1	timestamp#attribute	Blob
id#2	timestamp#attribute	Blob

Table US-EAST-2

Partition key	Sort key	Payload
id#0	timestamp#attribute	Blob
id#1	timestamp#attribute	Blob
id#2	timestamp#attribute	Blob

GSI US-EAST-1

Partition key	Sort key	Payload
id#0	timestamp#attribute	Blob
id#1	timestamp#attribute	Blob
id#2	timestamp#attribute	Blob

GSI US-EAST-2

Partition key	Sort key	Payload
id#0	timestamp#attribute	Blob
id#1	timestamp#attribute	Blob
id#2	timestamp#attribute	Blob

Our current throughput setting

Table		
Partition key	Sort key	Payload
id#0	timestamp#attribute	Blob
id#1	timestamp#attribute	Blob

GSI		
Partition key	Sort key	Payload
id#0	timestamp#attribute	Blob
id#1	timestamp#attribute	Blob

GSI		
Partition key	Sort key	Payload
id#0	timestamp#attribute	Blob
id#1	timestamp#attribute	Blob

How did we decide?

Table		
Partition key	Sort key	Payload
id#0	timestamp#attribute	Blob
id#1	timestamp#attribute	Blob

GSI		
Partition key	Sort key	Payload
id#0	timestamp#attribute	Blob
id#1	timestamp#attribute	Blob

GSI		
Primary key	Sort key	Payload
id#0	timestamp#attribute	Blob
id#1	timestamp#attribute	Blob

Our DynamoDB design considerations

Table
structure

Throughput

Table
management



Is this static?

Table management

How do we grow throughput?

How do we shrink throughput?

Table management

Table		
Partition key	Sort key	Payload
id#0	timestamp#attribute	Blob
id#1	timestamp#attribute	Blob

GSI		
Partition key	Sort key	Payload
id#0	timestamp#attribute	Blob
id#1	timestamp#attribute	Blob

GSI		
Partition key	Sort key	Payload
id#0	timestamp#attribute	Blob
id#1	timestamp#attribute	Blob

Table management

Table		
Partition key	Sort key	Payload
id#0	timestamp#attribute	Blob
id#1	timestamp#attribute	Blob
id#2	timestamp#attribute	Blob

GSI		
Partition key	Sort key	Payload
id#0	timestamp#attribute	Blob
id#1	timestamp#attribute	Blob
id#2	timestamp#attribute	Blob

GSI		
Partition key	Sort key	Payload
id#0	timestamp#attribute	Blob
id#1	timestamp#attribute	Blob
id#2	timestamp#attribute	Blob


Table management – Expansion



Table management – Expansion



Table management – Slow expansion



Our table management (pilot light)



Table set manager

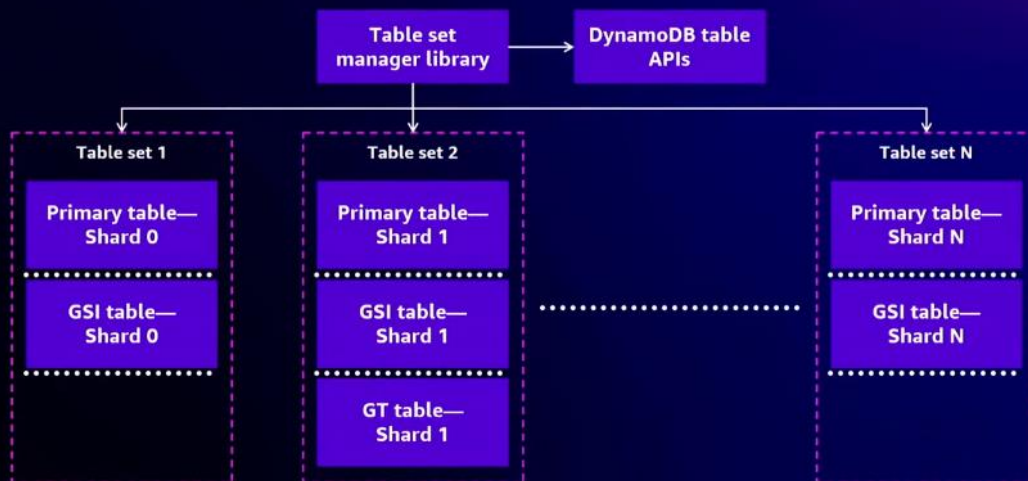



Table set manager

DynamoDB enables our flexibility

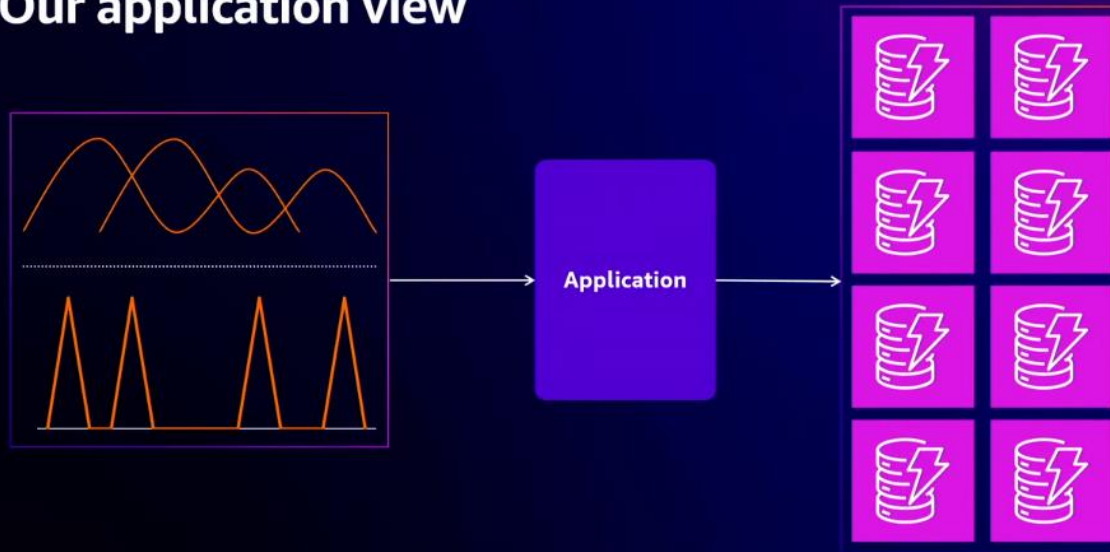
Pilot
light

Warm standby/
active-active

Our application view



Our application view



Results

80% capacity utilization

Max 0.0008% throttles per 60MM RCUs

Better operational recovery

Conclusions



High availability and high throughput

Pilot light operating model

Dynamic throughput

What's next



Evaluate pilot light to warm standby/active-active

Double down on our design pattern

If you take away just three things:

Resilient applications require a resilient database

DynamoDB offers a rich set of resilience features

AWS helps you build for resilience

Thank you!



Please complete the session survey in the mobile app

Jeff Duffy

Product Manager,
Amazon DynamoDB
Amazon Web Services

Tom Skinner

Director,
Measurement Infrastructure
Amazon Ads

Richard Edwards III

Principal Software Engineer
Amazon Ads