

Common Single-Table design modeling mistakes with DynamoDB



Tyler Walch
58 subscribers

Subscribe

53



Share

Download

Clip



1,292 views Aug 13, 2023

Advanced DynamoDB Techniques: Designing your access patterns is hard, designing your keys for scale is even harder! Learn these advanced techniques to modeling your DynamoDB keys before you write records to your database. I have been working with DynamoDB for years and have seen a lot of pitfalls from naively designed keys, in this video I go over every trick I have to model modeling DynamoDB keys with a focus on single-table design.

If you use TypeScript/JavaScript, checkout the library ElectroDB at <https://electrodb.dev> for a type-safe way to accomplish all of these techniques out of the box!

00:00:00 Introduction
00:00:56 Agenda/Overview
00:01:27 Partition and Sort keys
00:05:48 Single-Table design pros/cons
00:07:32 Our challenge
00:07:55 Presentation conventions
00:08:40 The "Course" entity
00:11:47 Identify natural keys
00:13:06 Separate your attributes from key fields
00:15:28 Standize casing and spacing
00:18:18 Use delimiters after each composite attribute
00:18:58 Why delimiters matter
00:24:29 Zero pad composite attribute numbers
00:26:16 Order composite attribures by hierarchy
00:29:18 Revisit "Course" entity access patterns
00:31:41 The "Certificate" entity
00:31:52 "Certificate" and "Course" entity comparison
00:34:06 Cross entity access patterns
00:39:35 Entity Isolation
00:42:03 Include static prefixes
00:46:23 High Volume of records per partition
00:47:38 High relationship density per partition
00:50:01 Isolating entity versions with keys
00:53:03 Partition key distribution with sharding
01:01:49 Final access patterns
01:05:13 ElectroDB for TypeScript users

Success with DynamoDB through sensible key design

Tyler W. Walch
Staff Engineer



@TINKERTAMPER

ElectroDB

- Open-Source JavaScript/TypeScript Project
- Define type-safe models and schemas for your Entities
- Abstracts away composing complex key patterns for you
- Puts Single Table Design within reach for existing projects

`npm install electrodb | yarn add electrodb`

Agenda

Start with a quick overview on DynamoDB Keys

Introduce our modeling challenge

Iterate on a key design

Intro to DynamoDB Keys

Partition Key

- Required on a DynamoDB table
- Impacts how your data is distributed within DynamoDB's storage nodes
- Must be provided in full when performing queries

Sort Key

- Optional on a DynamoDB table
- Impacts the order your query results are returned
- Can be partially provided when performing queries

Primary Key

The combination of the **partitionKey** and the **sortKey** can give you a unique Primary Key

“Single Table Design is just a fancy way of saying thoughtful string concatenation”

- Tyler W. Walch

Single Table Design

Pros

- Gain the ability to query across entities
- Less resources, less expensive, less queries

Cons

- Shared throughput limits across all entities
- More considerations, more complexity, more pitfalls

Our Challenge

- Model a new technical certification platform
- Design a key schema for **"Courses"** and **"Certificates"**
- Use Single Table Designs principles



Intro to DynamoDB Keys

PRESENTATION CONVENTIONS

- Partition Keys will be expressed with the color **YELLOW**
- Sort Keys will be expressed with the color **BLUE**
- All other attributes will be expressed with the color **GREY**

```
{  
  "pk": "your_partition_key",  
  "sk": "your_sort_key",  
  "field": "your_field",  
}
```

The "Course" Entity

The "Course" Entity

ACCESS PATTERNS

- Each record represents a single instance of a Course event that you can attend
- Table design only uses a **Partition Key**
- The Partition Key is named **"id"**

```
{  
  "id": "b921bd1f-af64...",  
  "courseName": "Intro to DynamoDB",  
  "location": "Building 1",  
  "startDate": "03/15/2022",  
  "courseType": "DevChat"  
}
```

The "Course" Entity

ACCESS PATTERNS

- Get Course by "id"

```
{
  "id": "b921bd1f-af64...",
  "courseName": "Intro to DynamoDB",
  "location": "Building 1",
  "startDate": "03/15/2022",
  "courseType": "DevChat"
}
```

Identify natural keys

Original Key

```
{
  "id": "b921bd1f-af64...",
  "courseName": "Intro to DynamoDB",
  "location": "Building 1",
  "startDate": "03/15/2022",
  "courseType": "DevChat"
}
```

Natural Keys

```
{
  "courseName": "Intro to DynamoDB",
  "location": "Building 1",
  "startDate": "03/15/2022",
  "courseType": "DevChat"
}
```

Separate your attributes from key fields

Before

```
{
  "courseName": "Intro to DynamoDB",
  "location": "Building 1",
  "startDate": "03/15/2022",
  "courseType": "DevChat"
}
```

After

```
{
  "pk": "Intro to DynamoDB",
  "sk": "03/15/2022 Building 1",
  "courseName": "Intro to DynamoDB",
  "startDate": "03/15/2022",
  "location": "Building 1",
  "courseType": "DevChat"
}
```


Standardize casing and spacing

Before

```
{  
  "pk": "Intro to DynamoDB",  
  "sk": "03/15/2022 Building 1",  
  "courseName": "Intro to DynamoDB",  
  "startDate": "03/15/2022",  
  "location": "Building 1",  
  "courseType": "DevChat"  
}
```

After

```
{  
  "pk": "introtodynamodb",  
  "sk": "03/15/2022building1",  
  "courseName": "Intro to DynamoDB",  
  "startDate": "03/15/2022",  
  "location": "Building 1",  
  "courseType": "DevChat"  
}
```

Use delimiters after each composite attribute

Before

```
{  
  "pk": "introtodynamodb",  
  "sk": "03/15/2022building1",  
  "courseName": "Intro to DynamoDB",  
  "startDate": "03/15/2022",  
  "location": "Building 1",  
  "courseType": "DevChat"  
}
```

After

```
{  
  "pk": "introtodynamodb",  
  "sk": "03/15/2022#building1#",  
  "courseName": "Intro to DynamoDB",  
  "startDate": "03/15/2022",  
  "location": "Building 1",  
  "courseType": "DevChat"  
}
```

Why delimiters matter

```
{city}, {state} {zip}  
Jackson, Wyoming 83002  
jacksonwyoming83002  
wyomingjackson83002
```

Rearrange to start broad and then make things more specific

```
begins_with("wyoming")  
  
wyomingcheyenne82001  
wyomingjackson83002  
wyominglaramie82073
```

```
begins_with("wyomingjackson")
```

```
wyomingjackson83002  
wyomingjacksonville82001
```

```
begins_with("wyoming#jackson")
```

```
wyoming#jackson#83002#  
Wyoming#jacksonville#82001#
```

```
begins_with("wyoming#jackson#")
```

```
wyoming#jackson#83002#
```

Using delimiters will get you all the zip codes in Wyoming and Jackson city and prevent over-querying DynamoDB tables

Use delimiters after each composite attribute

Before

```
{  
  "pk": "introtodynamodb",  
  "sk": "03/15/2022building1",  
  "courseName": "Intro to DynamoDB",  
  "startDate": "03/15/2022",  
  "location": "Building 1",  
  "courseType": "DevChat"  
}
```

After

```
{  
  "pk": "introtodynamodb",  
  "sk": "03/15/2022#building1#",  
  "courseName": "Intro to DynamoDB",  
  "startDate": "03/15/2022",  
  "location": "Building 1",  
  "courseType": "DevChat"  
}
```

Zero pad composite attribute numbers

Before

```
{
  "pk": "introtodynamodb",
  "sk": "03/15/2022#building1#",
  "courseName": "Intro to DynamoDB",
  "startDate": "03/15/2022",
  "location": "Building 1",
  "courseType": "DevChat"
}
```

After

```
{
  "pk": "introtodynamodb",
  "sk": "03/15/2022#building01#",
  "courseName": "Intro to DynamoDB",
  "startDate": "03/15/2022",
  "location": "Building 1",
  "courseType": "DevChat"
}
```

Order composite attributes by hierarchy

Before

```
{
  "pk": "introtodynamodb",
  "sk": "03/15/2022#building01#",
  "courseName": "Intro to DynamoDB",
  "startDate": "03/15/2022",
  "location": "Building 1",
  "courseType": "DevChat"
}
```

After

```
{
  "pk": "introtodynamodb",
  "sk": "2022/03/15#building01#",
  "courseName": "Intro to DynamoDB",
  "startDate": "03/15/2022",
  "location": "Building 1",
  "courseType": "DevChat"
}
```

How do you access your courses data, by date or by building?

The "Course" Entity

ACCESS PATTERNS

- Get Course by **"id"**

```
{
  "id": "b921bd1f-af64...",
  "courseName": "Intro to DynamoDB",
  "location": "Building 1",
  "startDate": "03/15/2022",
  "courseType": "DevChat"
}
```

The "Course" Entity

ACCESS PATTERNS

- Get Course occurrence by **"name"**, **"date"**, and **"location"**

```
{
  "pk": "introtodynamodb",
  "sk": "2022/03/15#building01#",
  "courseName": "Intro to DynamoDB",
  "startDate": "03/15/2022",
  "location": "Building 1",
  "courseType": "DevChat"
}
```

The "Course" Entity

ACCESS PATTERNS

- Get Course occurrence by **"name"**, **"date"**, and **"location"**
- Get Course occurrences by **"name"**

```
{
  "pk": "introtodynamodb",
  "sk": "2022/03/15#building01#",
  "courseName": "Intro to DynamoDB",
  "startDate": "03/15/2022",
  "location": "Building 1",
  "courseType": "DevChat"
}
```

The "Course" Entity

ACCESS PATTERNS

- Get Course occurrence by **"name"**, **"date"**, and **"location"**
- Get Course occurrences by **"name"**
- Get Course occurrences by **"name"** and **"year"**

```
{
  "pk": "introtodynamodb",
  "sk": "2022/03/15#building01#",
  "courseName": "Intro to DynamoDB",
  "startDate": "03/15/2022",
  "location": "Building 1",
  "courseType": "DevChat"
}
```


The "Course" Entity

ACCESS PATTERNS

- Get Course occurrence by **"name"**, **"date"**, and **"location"**
- Get Course occurrences by **"name"**
- Get Course occurrences by **"name"** and **"year"**
- Get Course occurrences by **"name"** and **"month"**

```
{
  "pk": "introtodynamodb",
  "sk": "2022/03/15#building01#",
  "courseName": "Intro to DynamoDB",
  "startDate": "03/15/2022",
  "location": "Building 1",
  "courseType": "DevChat"
}
```

The "Course" Entity

ACCESS PATTERNS

- Get Course occurrence by **"name"**, **"date"**, and **"location"**
- Get Course occurrences by **"name"**
- Get Course occurrences by **"name"** and **"year"**
- Get Course occurrences by **"name"** and **"month"**
- Get Course occurrences by **"name"** and **"date"**

```
{
  "pk": "introtodynamodb",
  "sk": "2022/03/15#building01#",
  "courseName": "Intro to DynamoDB",
  "startDate": "03/15/2022",
  "location": "Building 1",
  "courseType": "DevChat"
}
```

The "Course" Entity

ACCESS PATTERNS

- Get Course occurrence by **"name"**, **"date"**, and **"location"**
- Get Course occurrences by **"name"**
- Get Course occurrences by **"name"** and **"year"**
- Get Course occurrences by **"name"** and **"month"**
- Get Course occurrences by **"name"** and **"date"**
- Get Course occurrences by **"name"** and **"date"** by partial **"location"**

```
{
  "pk": "introtodynamodb",
  "sk": "2022/03/15#building01#",
  "courseName": "Intro to DynamoDB",
  "startDate": "03/15/2022",
  "location": "Building 1",
  "courseType": "DevChat"
}
```

The "Certificate" Entity


The "Certificate" Entity

"Course" Entity

```
{
  "pk": "introtodynamodb",
  "sk": "2022/03/15#building01#",
  "courseName": "Intro to DynamoDB",
  "startDate": "03/15/2022",
  "location": "Building 1",
  "courseType": "DevChat"
}
```

"Certificate" Entity

```
{
  "pk": "introtodynamodb",
  "sk": "2022/03/15#davidspurdy#",
  "certName": "Intro to DynamoDB",
  "issuedDate": "03/15/2022",
  "student": "David Spurdy",
  "certType": "Completion"
}
```




All Courses and associated Certificates

"Course" Entity

```
{
  "pk": "introtodynamodb",
  "sk": "2022/03/15#building01#",
  "courseName": "Intro to DynamoDB",
  "startDate": "03/15/2022",
  "location": "Building 1",
  "courseType": "DevChat"
}
```

"Certificate" Entity

```
{
  "pk": "introtodynamodb",
  "sk": "2022/03/15#davidspurdy#",
  "certName": "Intro to DynamoDB",
  "issuedDate": "03/15/2022",
  "student": "David Spurdy",
  "certType": "Completion"
}
```




All Courses and Certificates by "date"

"Course" Entity

```
{
  "pk": "introtodynamodb",
  "sk": "2022/03/15#building01#",
  "courseName": "Intro to DynamoDB",
  "startDate": "03/15/2022",
  "location": "Building 1",
  "courseType": "DevChat"
}
```

"Certificate" Entity

```
{
  "pk": "introtodynamodb",
  "sk": "2022/03/15#davidspurdy#",
  "certName": "Intro to DynamoDB",
  "issuedDate": "03/15/2022",
  "student": "David Spurdy",
  "certType": "Completion"
}
```



Single Table Design

"Course" Entity

```
{
  "pk": "introtodynamodb",
  "sk": "2022/03/15#building01#",
  "courseName": "Intro to DynamoDB",
  "startDate": "03/15/2022",
  "location": "Building 1",
  "courseType": "DevChat"
}
```

"Certificate" Entity

```
{
  "pk": "introtodynamodb",
  "sk": "2022/03/15#davidspurdy#",
  "certName": "Intro to DynamoDB",
  "issuedDate": "03/15/2022",
  "student": "David Spurdy",
  "certType": "Completion"
}
```

Entity Isolation

- Prevent record leaking from over-querying
- Optimize for either precision or cross entity joins

```
{
  "pk": "introtodynamodb",
  "sk": "2022/03/15#davidspurdy#",
  "certName": "Intro to DynamoDB",
  "issuedDate": "03/15/2022",
  "student": "David Spurdy",
  "certType": "Completion"
}
```

Entity Isolation: Static Prefixes

- Add static strings in your keys to namespace your entities
- Where you place your prefix can impact your access patterns

```
{
  "pk": "introtodynamodb",
  "sk": "2022/03/15#davidspurdy#"
}
```

```
{
  "pk": "introtodynamodb",
  "sk": "cert#2022/03/15#davidspurdy#"
}
```

Entity Isolation: Static Prefixes

- Add static strings in your keys to namespace your entities
- Where you place your prefix can impact your access patterns
- Location is unique to your individual use case

```
{  
  "pk": "introtodynamodb",  
  "sk": "cert#2022/03/15#davidspurdy#"  
}
```

```
{  
  "pk": "introtodynamodb",  
  "sk": "2022/03/15#cert#davidspurdy#"  
}
```

Entity Isolation: Static Prefixes



High Volume of Records per Partition

```
{  
  "pk": "introtodynamodb",  
  "sk": "cert#2022/03/15#davidspurdy#"  
}
```



High Relationship Density per Partition

```
{  
  "pk": "introtodynamodb",  
  "sk": "2022/03/15#cert#davidspurdy#"  
}
```

Entity Isolation: Versioning

- Adding a version number to the right of our static prefix
- Ensure the correct entity version is queried
- Leverage indexes in your migrations

```
{  
  "pk": "introtodynamodb",  
  "sk": "2022/03/15#cert#davidspurdy#"  
}
```

```
{  
  "pk": "introtodynamodb",  
  "sk": "2022/03/15#cert#01#davidspurdy#"  
}
```


Partition Key Distribution: Sharding

```
{
  "pk": "introtodynamodb",
  "sk": "2022/03/15#cert#01#davidspurdy#",
  "gsilpk": "completion",
  "gsilsk": "cert#01#tylerwalch#",
  "instructor": "Tyler Walch",
  "certType": "Completion"
}
```

Partition Key Distribution: Sharding

- Put/Delete/Update throughput can be throttled by your weakest Partition Key
- Access Patterns still need to consider Partition Key distribution

```
{
  "gsilpk": "completion",
  "gsilsk": "cert#01#tylerwalch#",
}
```

```
fun calcShardID(sortKey: String): Int {
    return sumUTFCodePoints(sortKey) % 20
}
```

```
{
  "gsilpk": "completion#18#",
  "gsilsk": "cert#01#tylerwalch#",
}
```

Final Access Patterns

```
{
  "pk": "introtodynamodb",
  "sk": "2022/03/15#cert#01#davidspurdy#",
  "gsilpk": "completion#19#",
  "gsilsk": "cert#01#tylerwalch#",
}
```

Final Access Patterns

- Get Certificate

```
{
  "pk": "introtodynamodb",
  "sk": "2022/03/15#cert#01#davidspurdy#",
  "gsilpk": "completion#19#",
  "gsilsk": "cert#01#tylerwalch#",
}
```

Final Access Patterns

- Get Certificate
- Get Certs by Date

```
{
  "pk": "introtodynamodb",
  "sk": "2022/03/15#cert#01#davidspurdy#",
  "gsilpk": "completion#19#",
  "gsilsk": "cert#01#tylerwalch#"
}
```

Final Access Patterns

- Get Certificate
- Get Certs (and Courses) by Date

```
{
  "pk": "introtodynamodb",
  "sk": "2022/03/15#cert#01#davidspurdy#",
  "gsilpk": "completion#19#",
  "gsilsk": "cert#01#tylerwalch#"
}
```

```
{
  "pk": "introtodynamodb",
  "sk": "2022/03/15#course#01#building01#"
}
```



Final Access Patterns

- Get Certificate
- Get Certs (and Courses) by Date
- Get Certs by Date

```
{
  "pk": "introtodynamodb",
  "sk": "2022/03/15#cert#01#davidspurdy#",
  "gsilpk": "completion#19#",
  "gsilsk": "cert#01#tylerwalch#"
}
```

Final Access Patterns

- Get Certificate
- Get Certs (and Courses) by Date
- Get Certs by Date
- Get Certs by Entity Version

```
{
  "pk": "introtodynamodb",
  "sk": "2022/03/15#cert#01#davidspurdy#",
  "gsilpk": "completion#19#",
  "gsilsk": "cert#01#tylerwalch#"
}
```

Final Access Patterns

- Get Certificate
- Get Certs (and Courses) by Date
- Get Certs by Date
- Get Certs by Entity Version
- Get All Completions

```
{
  "pk": "introtodynamodb",
  "sk": "2022/03/15#cert#01#davidspurdy#",
  "gsilpk": "completion#19#",
  "gsilsk": "cert#01#tylerwalch#"
}
```

Final Access Patterns

- Get Certificate
- Get Certs (and Courses) by Date
- Get Certs by Date
- Get Certs by Entity Version
- Get All Completion Certificates

```
{
  "pk": "introtodynamodb",
  "sk": "2022/03/15#cert#01#davidspurdy#",
  "gsilpk": "completion#19#",
  "gsilsk": "cert#01#tylerwalch#"
}
```

Final Access Patterns

- Get Certificate
- Get Certs (and Courses) by Date
- Get Certs by Date
- Get Certs by Entity Version
- Get All Completion Certificates
- Get All "v1" Completion Certificates

```
{
  "pk": "introtodynamodb",
  "sk": "2022/03/15#cert#01#davidspurdy#",
  "gsilpk": "completion#19#",
  "gsilsk": "cert#01#tylerwalch#"
}
```

Final Access Patterns

- Get Certificate
- Get Certs (and Courses) by Date
- Get Certs by Date
- Get Certs by Entity Version
- Get All Completion Certificates
- Get All "v1" Completion Certificates
- Get all Completion Certificates taught by Tyler Walch

```
{  
  "pk": "introtodynamodb",  
  "sk": "2022/03/15#cert#01#davidspurdy#",  
  "gs1pk": "completion#19#",  
  "gs1sk": "cert#01#tylerwalch#"  
}
```



ElectroDB

- Open-Source JavaScript/TypeScript Project
- Define type-safe models and schemas for your Entities
- Abstracts away composing complex key patterns for you
- Puts Single Table Design within reach for existing projects

`npm install electrodb | yarn add electrodb`

