

12-Factor Agents: Patterns of reliable LLM applications — Dex Horthy, HumanLayer

AI Engineer

147K subscribers

Subscribe

15

Share

Clip

Save

164 views Jul 3, 2025 [AIEWF 2025 Complete Playlist](#)

Hi, I'm Dex. I've been hacking on AI agents for a while.

I've tried every agent framework out there, from the plug-and-play crew/langchains to the "minimalist" smolagents of the world to the "production grade" langraph, griptape, etc.

I've talked to a lot of really strong founders who are all building really impressive things with AI. Most of them are rolling the stack themselves. I don't see a lot of frameworks in production customer-facing agents.

I've been surprised to find that most of the products out there billing themselves as "AI Agents" are not all that agentic. A lot of them are mostly deterministic code, with LLM steps sprinkled in at just the right points to make the experience truly magical.

Agents, at least the good ones, don't follow the "here's your prompt, here's a bag of tools, loop until you hit the goal" pattern. Rather, they are comprised of mostly just software.

So, I set out to answer:

What are the principles we can use to build LLM-powered software that is actually good enough to put in the hands of production customers?

The Short Version: The 12 Factors

Even if LLMs continue to get exponentially more powerful, there will be core engineering techniques that make LLM-powered software more reliable, more scalable, and easier to maintain.

How We Got Here: A Brief History of Software

Factor 1: Natural Language to Tool Calls

Factor 2: Own your prompts

Factor 3: Own your context window

Factor 4: Tools are just structured outputs

Factor 5: Unify execution state and business state

Factor 6: Launch/Pause/Resume with simple APIs

Factor 7: Contact humans with tool calls

Factor 8: Own your control flow

Factor 9: Compact Errors into Context Window

Factor 10: Small, Focused Agents

Factor 11: Trigger from anywhere, meet users where they are

Factor 12: Make your agent a stateless reducer

https://github.com/humanlayer/12-factor-agents

humanlayer / 12-factor-agents

Type to search

CodeIssues7Pull requests1ActionsProjectsSecurityInsights

12-factor-agentsPublic

Watch67Fork317Star5.1k

main14 Branches0 Tags

Go to fileAdd fileCode

dexhorthyUpdate brief-history-of-software.md1070b6d · yesterday241 Commits

content	Update brief-history-of-software.md	yesterday
hack/contributors_markdown	add contributors markdown stuff	2 months ago
img	update factor 9 image	3 months ago
packages/walkthroughgen	Fix typos	2 months ago
workshops	Fix typos	2 months ago
LICENSE	license stuff	2 months ago
README.md	Update README.md	last week

About

What are the principles we can use to build LLM-powered software that is actually good enough to put in the hands of production customers?

frameworkaimemoryorchestrationagents12-factorragprompt-engineeringllmscontext-window12-factor-agents

ReadmeView licenseActivityCustom properties

# 12-Factor Agents - Principles for building reliable LLM applications

Code [Apache 2.0](#) Content [CC BY-SA 4.0](#) chat [discord](#) youtube [deep dive](#)

In the spirit of [12 Factor Apps](#). The source for this project is public at <https://github.com/humanlayer/12-factor-agents>, and I welcome your feedback and contributions. Let's figure this out together!



## The Short Version: The 12 Factors

Even if LLMs [continue to get exponentially more powerful](#), there will be core engineering techniques that make LLM-powered software more reliable, more scalable, and easier to maintain.

- [How We Got Here: A Brief History of Software](#)
- [Factor 1: Natural Language to Tool Calls](#)
- [Factor 2: Own your prompts](#)
- [Factor 3: Own your context window](#)
- [Factor 4: Tools are just structured outputs](#)
- [Factor 5: Unify execution state and business state](#)
- [Factor 6: Launch/Pause/Resume with simple APIs](#)
- [Factor 7: Contact humans with tool calls](#)
- [Factor 8: Own your control flow](#)
- [Factor 9: Compact Errors into Context Window](#)
- [Factor 10: Small, Focused Agents](#)
- [Factor 11: Trigger from anywhere, meet users where they are](#)
- [Factor 12: Make your agent a stateless reducer](#)

Building Agents in 2025

# The Agent Journey

1. Decide you want to build an agent
2. Product design, UX mapping, what problems to solve
3. Want to move fast, so grab \$FRAMEWORK and *get to building*
4. Get to 70-80% quality bar
5. Realize that 80% isn't good enough
6. Realize that getting past 80% requires reverse-engineering the framework, prompts, flow, etc.
7. Start over from scratch

8. Realize that this isn't a good problem for agents!

## Principles of Reliable LLM Applications

1. There are some core things that make agents great
2. Greenfield rewrite w/ a framework may be counter-productive
3. Apply small, modular concepts to existing code
4. Don't need an AI Background

## THE TWELVE-FACTOR APP

### INTRODUCTION

In the modern era, software is commonly delivered as a service: called *web apps*, or *software-as-a-service*. The twelve-factor app is a methodology for building software-as-a-service apps that:

- Use **declarative** formats for setup automation, to minimize time and cost for new developers joining the project;
- Have a **clean contract** with the underlying operating system, offering **maximum portability** between execution environments;
- Are suitable for **deployment** on modern **cloud platforms**, obviating the need for servers and systems administration;
- **Minimize divergence** between development and production, enabling **continuous deployment** for maximum agility;
- And can **scale up** without significant changes to tooling, architecture, or development practices.

The twelve-factor methodology can be applied to apps written in any programming language, and which use any combination of backing services (database, queue, memory cache, etc).

### BACKGROUND

The contributors to this document have been directly involved in the development and deployment of hundreds of apps, and

## THE TWELVE FACTORS

### I. Codebase

One codebase tracked in revision control, many deploys

### II. Dependencies

Explicitly declare and isolate dependencies

### III. Config

Store config in the environment

### IV. Backing services

Treat backing services as attached resources

### V. Build, release, run

Strictly separate build and run stages

### VI. Processes

Execute the app as one or more stateless processes

### VII. Port binding

Export services via port binding

### VIII. Concurrency

Scale out via the process model

### IX. Disposability

Maximize robustness with fast startup and graceful shutdown

### X. Dev/prod parity

Keep development, staging, and production as similar as possible

### XI. Logs

Treat logs as event streams

### XII. Admin processes

Run admin/management tasks as one-off processes

# 12 Factor Agents

## Principles of Reliable LLM Applications



<https://hlyr.dev/12fa>



README License

## 12-Factor Agents - Principles for building reliable LLM applications

Code Apache 2.0 Content CC BY-SA 4.0 chat discord youtube deep

In the spirit of [12 Factor Apps](#). The source for this project is public at <https://github.com/humanlayer/12-factor-agents>, and I welcome your feedback figure this out together!



**Hacker News** new | threads | past | comments | ask | show | jobs | submit

\* 12-factor Agents: Patterns of reliable LLM applications (github.com/humanlayer)

475 points by dorthy 43 days ago | hide | past | favorite | 78 comments

I've been building AI agents for a while. After trying every framework out there and talking to ma that make it to production aren't actually that agentic. The best ones are mostly just well-enginee

So I set out to document what I've learned about building production-grade AI systems: <https://github.com/humanlayer/12-factor-agents> powered software that's reliable enough to put in the hands of production customers.

In the spirit of Heroku's 12 Factor Apps (<https://12factor.net/>), these principles focus on the engi maintainable. Even as models get exponentially more powerful, these core techniques will remain

I've seen many SaaS builders try to pivot towards AI by building greenfield new projects on agent bar with out-of-the-box tools. The ones that did succeed tended to take small, modular concepts t starting from scratch.

The full guide goes into detail on each principle with examples and patterns to follow. I've seen th

I'm sharing this as a starting point—the field is moving quickly so these principles will evolve. I w AI systems!



## Factor 1

### Natural Language to tool calls

Can you create a payment link to Terri for \$750 for sponsoring the February meetup?

```
{
  "function": {
    "name": "create_payment_link",
    "parameters": {
      "amount": 750,
      "customer": "cust_128934ddasf9",
      "product": "prod_8675309",
      "price": "prc_09874329fds",
      "quantity": 1,
      "memo": "Hey Jeff - see below for the payment link for the february ai tinkers meetup"
    }
  }
}
```

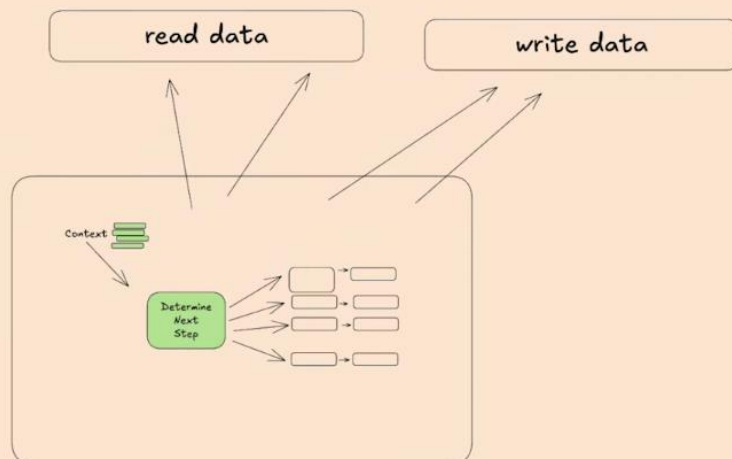
### Factor 1 - Natural Language -> Tool (API) Calls

Add a ticket to restock  
the file cabinet

Read this email thread  
and update the CRM for  
AcmeCorp

What's the status of my package

That's wrong, use the XYZ project



## Factor 4

### Tools are structured outputs

#### Edgar Dijkstra: Go To Statement Considered Harmful

##### Go To Statement Considered Harmful

**Key Words and Phrases:** go to statement, jump instruction, branch instruction, conditional clause, alternative clause, repetitive clause, program intelligibility, program sequencing  
**CR Categories:** 4.22, 5.23, 5.24

##### EDITOR:

For a number of years I have been familiar with the observation that the quality of programmers is a decreasing function of the density of **go to** statements in the programs they produce. More recently I discovered why the use of the **go to** statement has such disastrous effects, and I became convinced that the **go to** state-

dynamic progress is only characterized when we also give to which call of the procedure we refer. With the inclusion of procedures we can characterize the progress of the process via a sequence of textual indices, the length of this sequence being equal to the dynamic depth of procedure calling.

Let us now consider repetition clauses (like, **while B repeat A or repeat A until B**). Logically speaking, such clauses are now superfluous, because we can express repetition with the aid of recursive procedures. For reasons of realism I don't wish to exclude them: on the one hand, repetition clauses can be implemented quite comfortably with present day finite equipment; on the other hand, the reasoning pattern known as "induction"

## “Tool Use” Considered Harmful

### Tool Use

LLM outputs JSON

Deterministic Code Does Something

(maybe) Results Fed Back to LLM

```
class Issue:
    title: str
    description: str
    team_id: str
    assignee_id: str

class CreateIssue:
    intent: "create_issue"
    issue: Issue

class SearchIssues:
    intent: "search_issues"
    query: str
    what_youre_looking_for: str
```

```
if nextStep.intent == 'create_payment_link':
    stripe.paymentlinks.create(nextStep.parameters)
    return # or whatever you want, see below
elif nextStep.intent == 'wait_for_a_while':
    # do something monadic idk
else: #... the model didn't call a tool we know about
    # do something else
```

There's nothing special about tools

it's just JSON + Code

## Factor 4 - Tools are Structured Output

```
const openai_tools: ChatCompletionTool[] = [
  {
    type: "function",
    function: {
      name: "multiply",
      description: "multiply two numbers",
      parameters: {
        type: "object",
        properties: {
          a: { type: "number" },
          b: { type: "number" },
        },
        required: ["a", "b"],
      },
    },
  },
  {
    type: "function",
    function: {
      name: "add",
      description: "add two numbers",
      parameters: {
        type: "object",
        properties: {
          a: { type: "number" },
          b: { type: "number" },
        },
        required: ["a", "b"],
      },
    },
  },
];
```

JSON Schema

Prompt-Only

```
What should the next step be?

Answer in JSON using any of these schemas:
{
  // you can request more information from me
  intent: "request_more_information",
  message: string,
} or {
  intent: "create_issue",
  issue: {
    title: string,
    description: string,
    team_id: string,
    // name of the team to create the issue in
    team_name: string or null,
    project_id: string or null,
    // name of the project to create the issue in
    project_name: string or null,
    assignee_id: string or null,
    // name of the user to assign the issue to
    assignee_name: string or null,
    labels_ids: string[],
    labels_names: string[],
    // The priority of the issue.
    // 0 = No priority, 1 = Urgent, 2 = High, 3 = Normal, 4 = Low.
    priority: int or null,
  },
} or {
  intent: "list_teams",
  // what is your goal or desired outcome with this operation
  query_intent: string or null,
} or {
```


... better, but I know you want to  
RYTHING

Benefits: Flexibility, Prompt Control

## Factor 8 Own your control flow


### 60 years ago

software is a Directed Graph




### 20 years ago

DAG Orchestrators



# the promise of agents

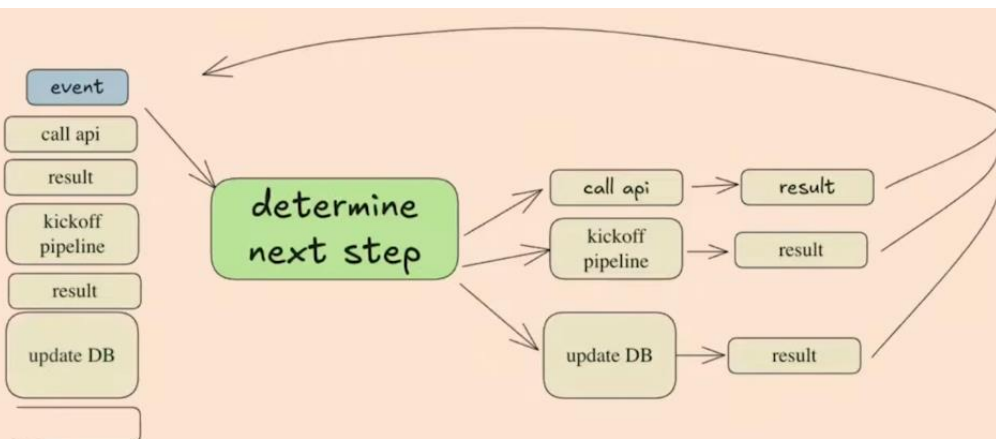
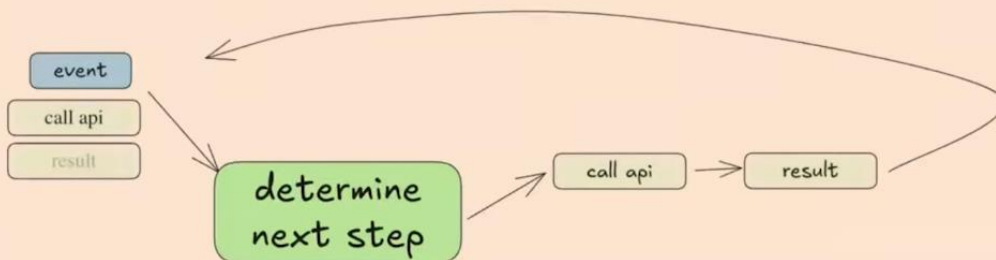
Stop doing DAGs




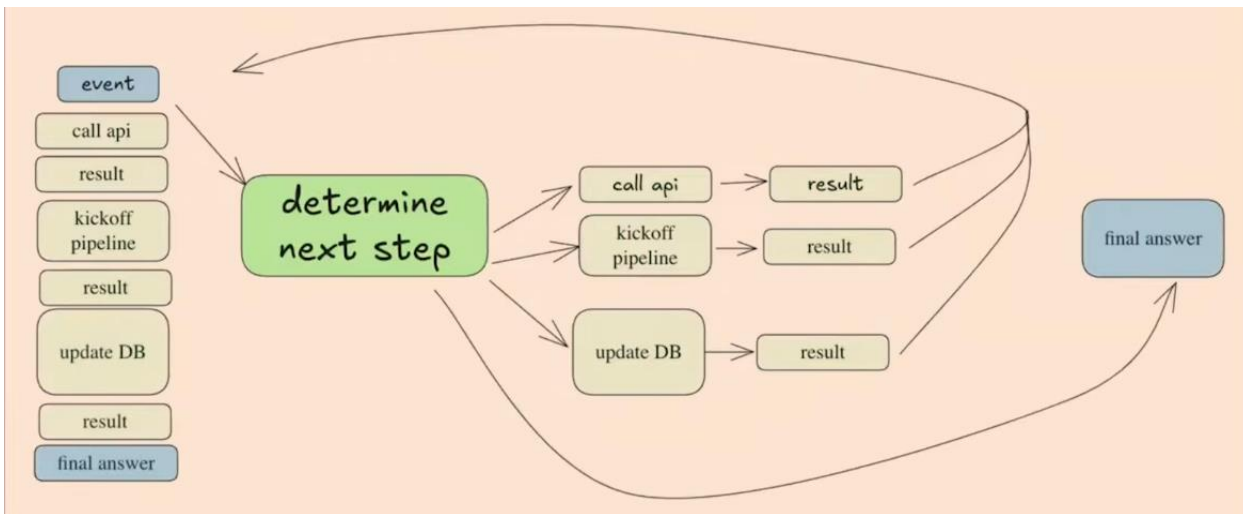
```
initial_event = {"message": "..."}
context = [initial_event]
while True:
    next_step = await llm.determine_next_step(context)
    context.append(next_step)

    if (next_step.intent == "done"):
        return next_step.final_answer

    result = await execute_step(next_step)
    context.append(result)
```





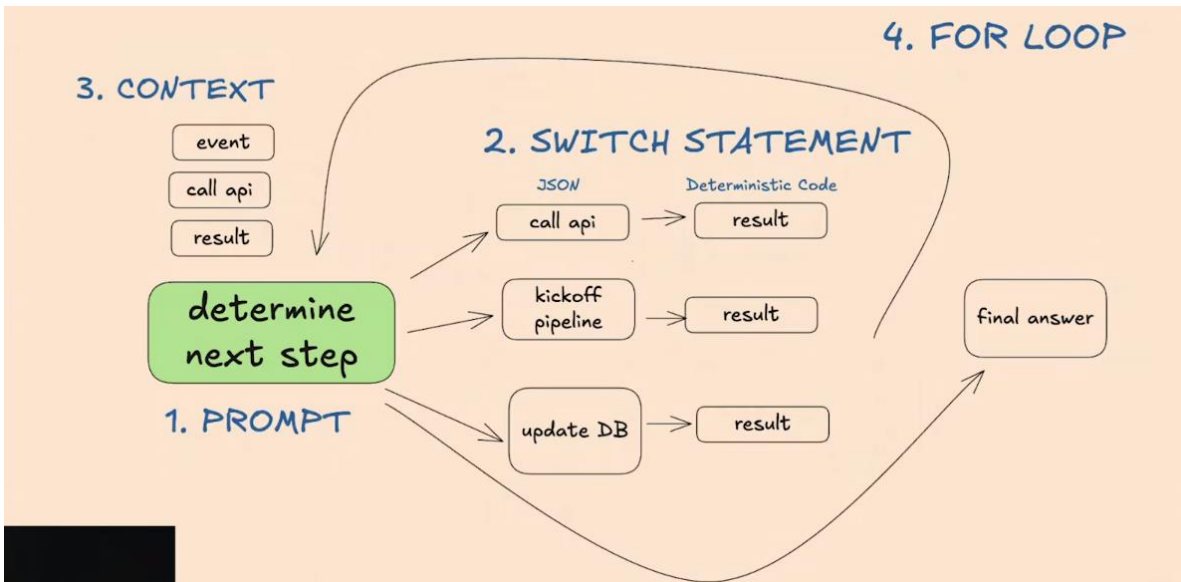


Turns out this doesn't really work

Mostly: Long Context windows

Even as models support longer and longer context windows...  
 ...you'll ALWAYS get better results with a small, focused prompt and context

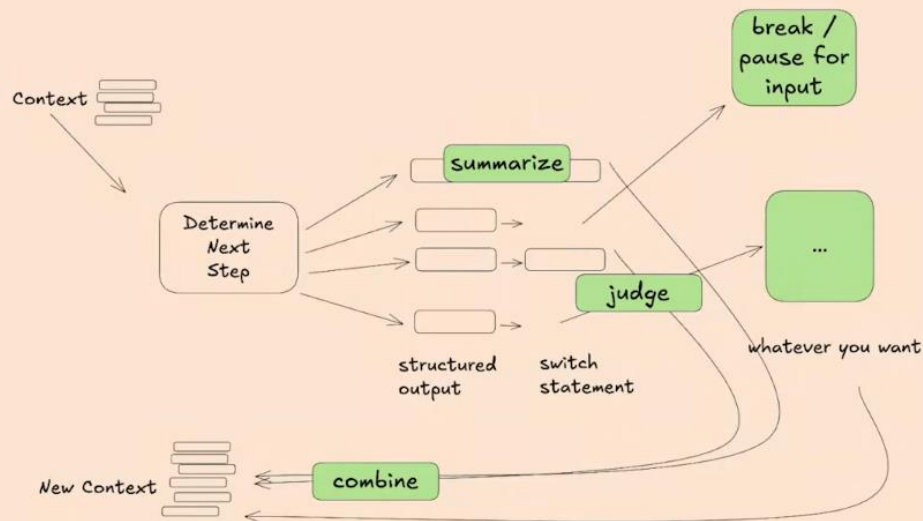
what's an agent really?!



If you own your control flow, you can

Break  
Switch  
Summarize  
Judge

### Factor 8 - Own Your Control Flow - BYO Switch Statement



Benefits: Interrupts, Process Control

### Factor 5 / 6

Unify Execution State and Business State

Launch / Pause / Resume with simple APIs

Execution state:

current step  
next step  
waiting status  
retry config  
etc

Business state: What's happened in the agent workflow so far

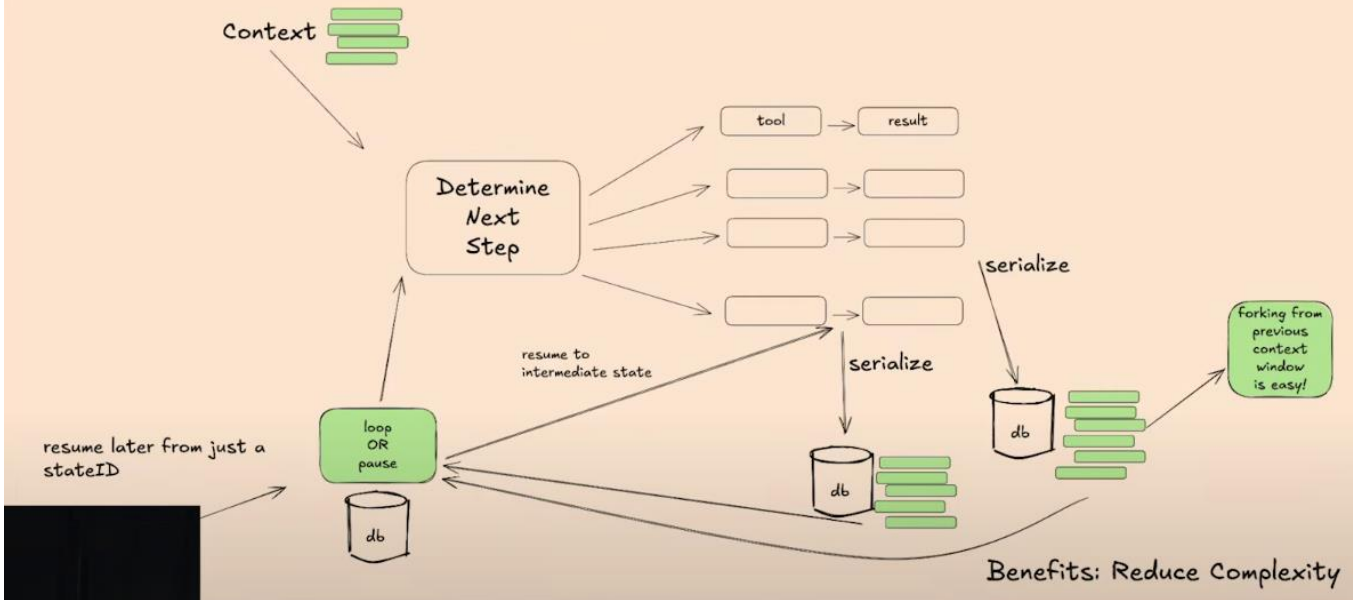
list of OpenAI messages  
list of tool calls and results  
etc

Launch

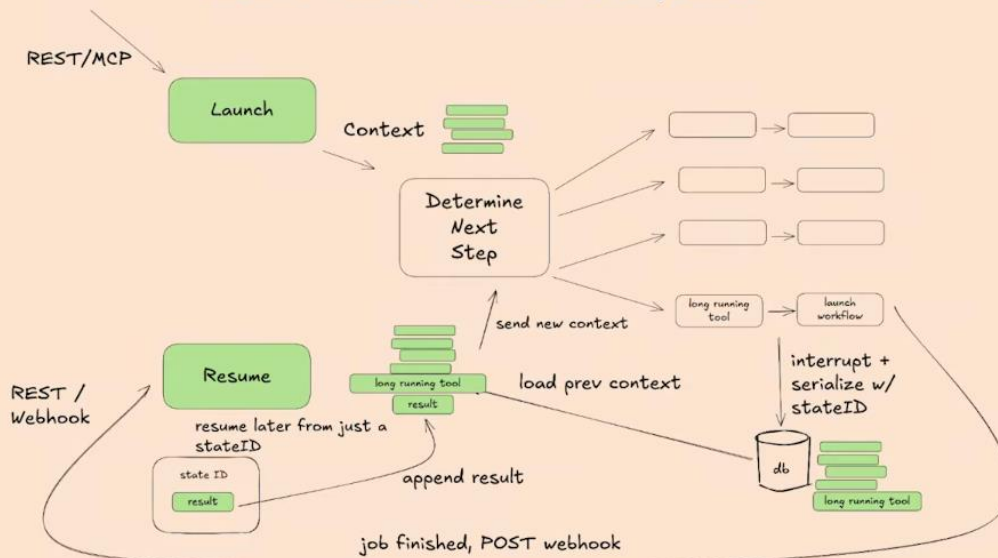
Pause

Resume

## Factor 5 - Unify Execution State and Business State



## Factor 6- Pause / Resume with simple APIs



# Building great agents requires flexibility and creativity

Don't box yourself in

## Factor 2

Own your prompts

```

agent = Agent(
    role="...",
    goal="...",
    personality="...",
    tools=[tool1, tool2, tool3]
)

task = Task(
    instructions="...",
    expected_output=OutputModel
)

result = agent.run(task)

```

## Black Box

```

function DetermineNextStep(thread: string) -> DoneForNow | ListGitTags | DeployBackend | DeployFrontend | RequestMo
prompt #"
  {{ _.role("system") }}

  You are a helpful assistant that manages deployments for frontend and backend systems.
  You work diligently to ensure safe and successful deployments by following best practices
  and proper deployment procedures.

  Before deploying any system, you should check:
  - The deployment environment (staging vs production)
  - The correct tag/version to deploy
  - The current system status

  You can use tools like deploy_backend, deploy_frontend, and check_deployment_status
  to manage deployments. For sensitive deployments, use request_approval to get
  human verification.

  Always think about what to do first, like:
  - Check current deployment status
  - Verify the deployment tag exists
  - Request approval if needed
  - Deploy to staging before production
  - Monitor deployment progress

  {{ _.role("user") }}

  {{ thread }}

  What should the next step be?
"#
}

```

## Full Control

## LLMs are Pure Functions

Tokens in → Tokens Out

### Factor 2 - Own your prompts

```

Agent:
  role: str
  goal: str
  personality: str

Task:
  instructions: str
  expected_output: Type[BaseModel]
  tools: list[tool]

agent.run(task)

```

black box

I don't know what's better, but I know you want to be able to try EVERYTHING

```

> Open Playground >
function DetermineNextStep(
  // To keep this clean, make the client turn the thread into a prompt-ready string.
  // didn't wanna solve that in Jinja (although long term that's probably the best solution)
  thread: string
) -> ClarificationRequest | CreateIssue | ListTeams | ListIssues | ListUsers | DoneForNow | AddComment |
  client CustomGPT4o

prompt #"
  {{ _.role("system") }}

  You are a helpful assistant that helps the user with their linear issue management.
  You work hard for whoever sent the inbound initial email, and want to do your best
  to help them do their job by carrying out tasks against the linear api.

  Before creating an issue, you should ensure you have accurate team/user/project ids.
  You can list_teams and list_users and list_projects functions to get ids.

  If you are BCC'd on a thread, assume that the user is asking you to look up the related issue and

  Always think about what to do first, like:

  - ...
  - ...
  - ...

  {{ _.role("user") }}

```

full control

Benefits: Flexibility, Optimization



## Factor 3

### Own your Context Building

```
[
  {
    "role": "system",
    "content": "You are a helpful assistant..."
  },
  {
    "role": "user",
    "content": "Can you deploy the backend?"
  },
  {
    "role": "assistant",
    "content": null,
    "tool_calls": [
      {
        "id": "1",
        "name": "list_git_tags",
        "arguments": "{}"
      }
    ]
  },
  {
    "role": "tool",
    "name": "list_git_tags",
    "content": "{\n  \"tags\": [\n    {\n      \"name\": \"v1.2.3\", \"commit\": \"abc123\", \"date\": \"2024-03-15T10:00:00Z\"\n    },\n    {\n      \"name\": \"v1.2.2\", \"commit\": \"def456\", \"date\": \"2024-03-14T15:30:00Z\"\n    },\n    {\n      \"name\": \"v1.2.1\", \"commit\": \"ghi789\", \"date\": \"2024-03-13T09:15:00Z\"\n    }\n  ]\n}",
    "tool_call_id": "1"
  }
]
```

## Standard Messages Format

```
{
  "role": "system",
  "content": "You are a helpful assistant..."
},
{
  "role": "user",
  "content": |
    Here's everything that happened so far:

    <slack_message>
      From: @alex
      Channel: #deployments
      Text: Can you deploy the backend?
    </slack_message>

    <list_git_tags>
      intent: "list_git_tags"
    </list_git_tags>

    <list_git_tags_result>
      tags:
        - name: "v1.2.3"
          commit: "abc123"
          date: "2024-03-15T10:00:00Z"
        - name: "v1.2.2"
          commit: "def456"
          date: "2024-03-14T15:30:00Z"
        - name: "v1.2.1"
          commit: "ghi789"
          date: "2024-03-13T09:15:00Z"
    </list_git_tags_result>

    what's the next step?
  |
}
}
```

## Agentic Trace in Single User Message

```
class Thread:
    events: List[Event]

class Event:
    # could just use string, or could be explicit - up to you
    type: Literal["list_git_tags", "deploy_backend", "deploy_frontend", "request_more_information", "done_for_now",
    data: ListGitTags | DeployBackend | DeployFrontend | RequestMoreInformation |
        ListGitTagsResult | DeployBackendResult | DeployFrontendResult | RequestMoreInformationResult | string

def event_to_prompt(event: Event) -> str:
    data = event.data if isinstance(event.data, str) \
        else stringifyToYaml(event.data)

    return f"<{event.type}>\n{data}\n</{event.type}>"

def thread_to_prompt(thread: Thread) -> str:
    return '\n\n'.join(event_to_prompt(event) for event in thread.events)
```

```

<slack_message>
  From: @alex
  Channel: #deployments
  Text: Can you deploy the latest backend to production?
  Thread: []
</slack_message>

<deploy_backend>
  intent: "deploy_backend"
  tag: "v1.2.3"
  environment: "production"
</deploy_backend>

<error>
  error running deploy_backend: Failed to connect to deployment service
</error>

<request_more_information>
  intent: "request_more_information_from_human"
  question: "I had trouble connecting to the deployment service, can you provide more details and/or check on the :
</request_more_information>

<human_response>
  data:
    response: "I'm not sure what's going on, can you check on the status of the latest workflow?"
</human_response>

```

Tokens	Characters	Tokens	Characters
87	336	57	182

```

[
  {
    "type": "user_input",
    "data": "can you multiply 3 and 4, then divide
    then divide the result by 2
    and then add 12 to that result?"
  },
  {
    "type": "tool_call",
    "data": {
      "intent": "multiply",
      "a": 3,
      "b": 4
    }
  }
]

```

```

<user_input>
  can you multiply 3 and 4, then divide
  the result by 2 and then add 12 to that result?
</user_input>


<multiply>
  a: 3
  b: 4
</multiply>

```

## Everything is Context Engineering

Prompt  
 Memory  
 RAG  
 Agentic History  
 Structured Output

# Everything is Context Engineering



## Factor 3 - Own your context building

```
{
  "role": "system",
  "content": "you are a helpful assistant"
},
{
  "role": "user",
  "content": "whats the weather in tokyo"
},
{
  "role": "assistant",
  "tool_calls": [...]
},
{
  "role": "tool",
  ...
},
{
  "role": "assistant",
  "tool_calls": [...]
},
}
```

```
{
  "role": "system",
  "content": "you are a helpful assistant"
},
{
  "role": "user",
  "content": "Here's what happened so far:

<initial_message>
  whats the weather in tokyo
</initial_message>

<check_weather>
  city: tokyo
</check_weather>

<check_weather_response>
  {city: "tokyo", "temp": "73", "skies": "cloudy"}
</check_weather_response>

whats the next step?"
}
```

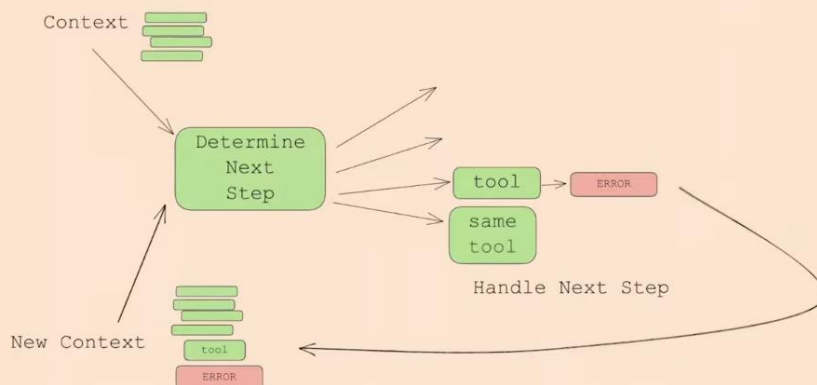
I don't know what's better, but I know you want to be able to try EVERYTHING

Benefits: Flexibility, Prompt Control

## Factor 9

### Compact Errors into context window

## Factor 9 - Compact Errors into Context Window



```

thread = {"events": [initial_message]}

while True:
    next_step = await determine_next_step(thread_to_prompt(thread))
    thread["events"].append({
        "type": next_step.intent,
        "data": next_step,
    })
    try:
        result = await handle_next_step(thread, next_step) # our switch statement
    except Exception as e:
        # if we get an error, we can add it to the context window and try again
        thread["events"].append({
            "type": 'error',
            "data": format_error(e),
        })
    # loop, or do whatever else here to try to recover

```

## What about Spin-Outs

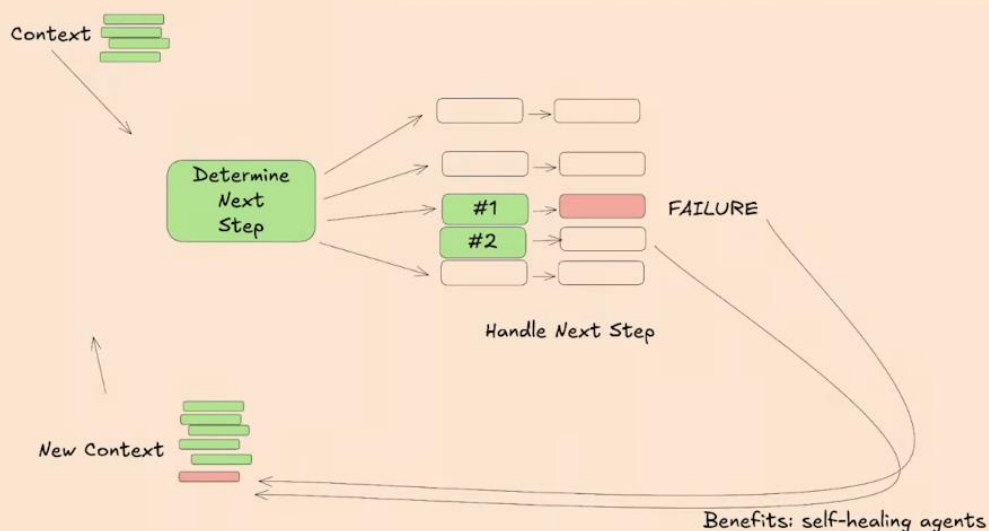
### Errors: Own your context window

Clear errors once you get a valid tool call

Summarize and redact

Limit Loop Count

### Factor 9 - Compact Errors into Context Window

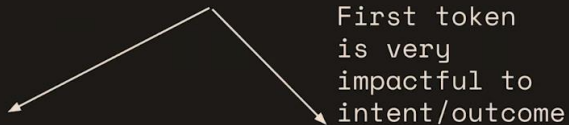


## Factor 7

Contact Humans with Tools



User: What's the weather in Tokyo?



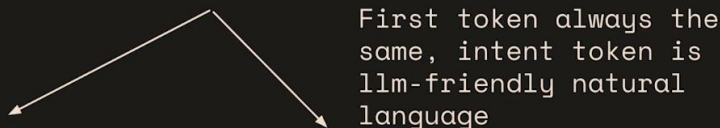
First token  
is very  
impactful to  
intent/outcome

```
|JSON> {  
  "function":  
    "check_weather",  
  "parameters":  
    [{"city_name": "tokyo"}]  
}
```

I'm a large language  
model, and I don't  
have access to  
up-to-date  
information about...

Tools push the emphasis to natural  
language **intent** token(s)

User: What's the weather in Tokyo?



First token always the  
same, intent token is  
llm-friendly natural  
language

```
|JSON> {  
  "function":  
    "check_weather",  
  "parameters":  
    [{"city_name": "tokyo"}]  
}
```

```
|JSON> {  
  "function":  
    "request_clarification",  
  "parameters":  
    [{"message": "is that  
the tokyo in japan?"}]  
}
```

```
class Options:  
    urgency: Literal["low", "medium", "high"]  
    format: Literal["free_text", "yes_no", "multiple_choice"]  
    choices: List[str]  
  
# Tool definition for human interaction  
class RequestHumanInput:  
    intent: "request_human_input"  
    question: str  
    context: str  
    options: Options  
  
# Example usage in the agent loop  
if nextStep.intent == 'request_human_input':  
    thread.events.append({  
        type: 'human_input_requested',  
        data: nextStep  
    })  
    thread_id = await save_state(thread)  
    await notify_human(nextStep, thread_id)  
    return # Break loop and wait for response to come back with thread ID  
else:  
    # ... other cases
```

```

<slack_message>
  From: @alex
  Channel: #deployments
  Text: Can you deploy backend v1.2.3 to production?
  Thread: []
</slack_message>

<request_human_input>
  intent: "request_human_input"
  question: "Would you like to proceed with deploying v1.2.3 to production?"
  context: "This is a production deployment that will affect live users."
  options: {
    urgency: "high"
    format: "yes_no"
  }
</request_human_input>

<human_response>
  response: "yes please proceed"
  approved: true
  timestamp: "2024-03-15T10:30:00Z"
  user: "alex@company.com"
</human_response>


<deploy_backend>
  intent: "deploy_backend"
  tag: "v1.2.3"
  environment: "production"
</deploy_backend>

```

## Gen 2 vs. Gen 3 Agents

### Gen 2: Reasoning / Collaboration / Tools

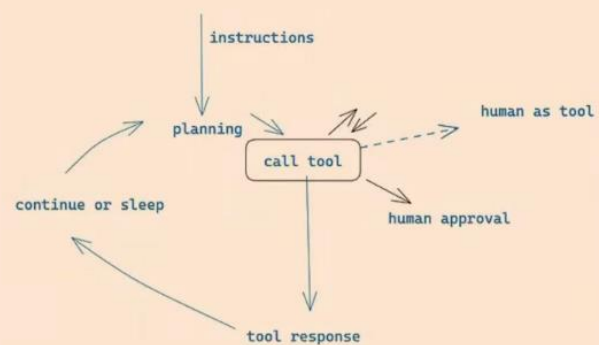
#### Planning and Reasoning and Routing



#### Tool Calling



### Gen 3: Outer Loop



## Factor 7 - Contact Human with tools

content: "the weather in tokyo is 57 and rainy"

OR

content: "for what postal code?"

OR

tool\_calls: [...]

model switches between  
tools and plaintext

```

tool_calls: [{
  function: "final_answer",
  arguments: {
    message: "the weather in tokyo is 57 and rainy",
  }
}]

```

OR

```

tool_calls: [{
  function: "request_clarification",
  arguments: {
    message: "for what postal code?",
  }
}]

```

OR

```

tool_calls: {
  function: "check_weather_in_city",
  arguments: {
    "city": "tokyo",
    "postal_code": "12345"
  }
}

```

everything  
is tools

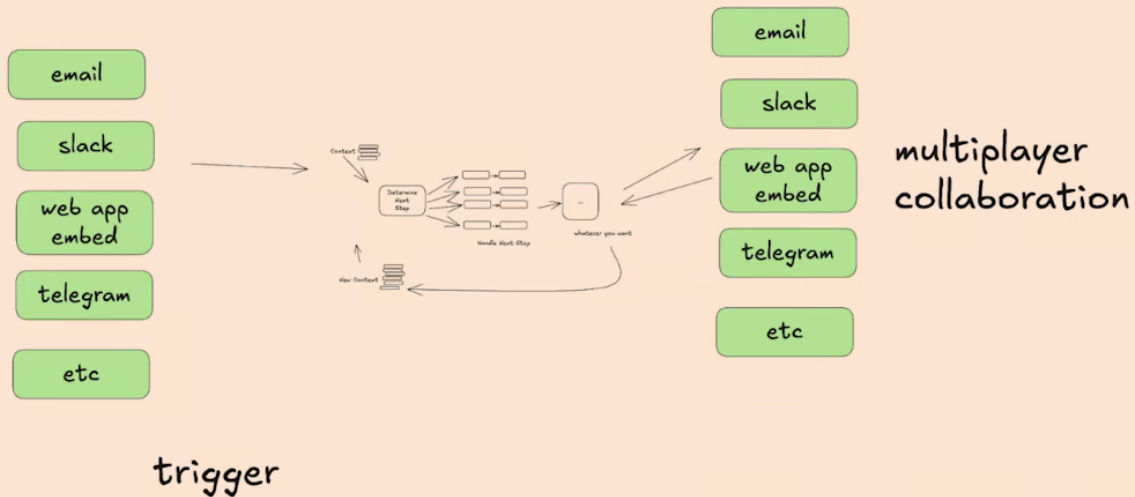
Final answer vs. I need input vs. call tool is fiddly

Benefits: Clear instructions for the model, inner vs. outer loop,  
access to multiple humans

## Factor 11

Trigger from anywhere, meet users where they are

### Factor 11 - Trigger From Anywhere, Meet Users where they are




Benefits: Agents that feel like coworkers

## Factor 10

Small Focused Agents

### 3. CONTEXT



### 2. SWITCH STATEMENT

JSON      Deterministic Code

call api → result

kickoff pipeline → result

update DB → result

determine next step

### 1. PROMPT

final answer


### 4. FOR LOOP

Turns out this doesn't really work

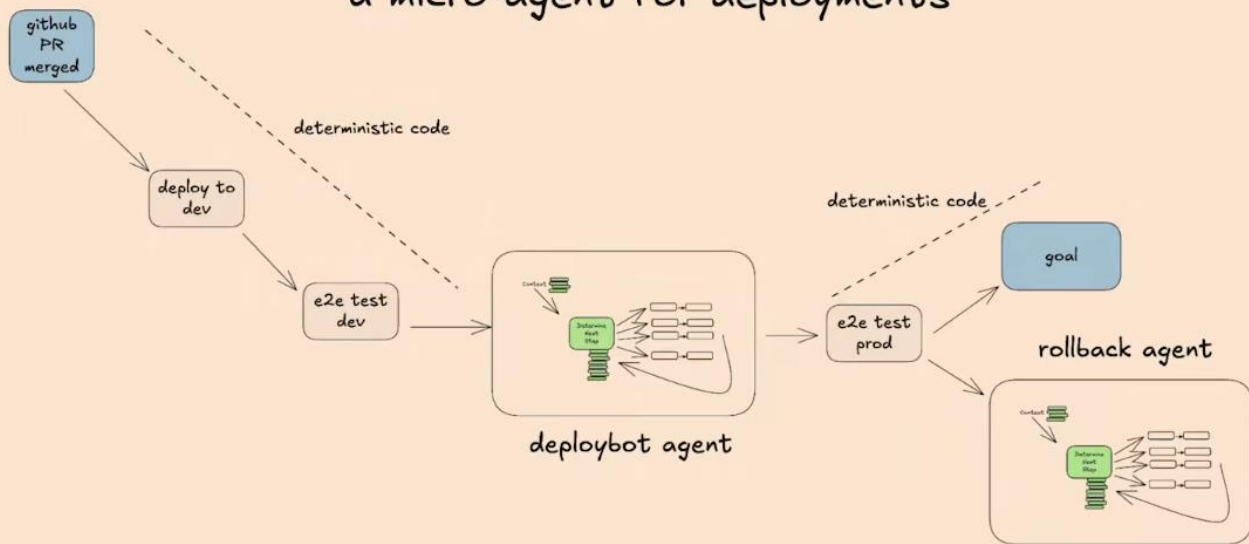
So what does?

## what about micro agents?

10-100 tools, 3-10 steps as part of larger pipeline



## a micro-agent for deployments

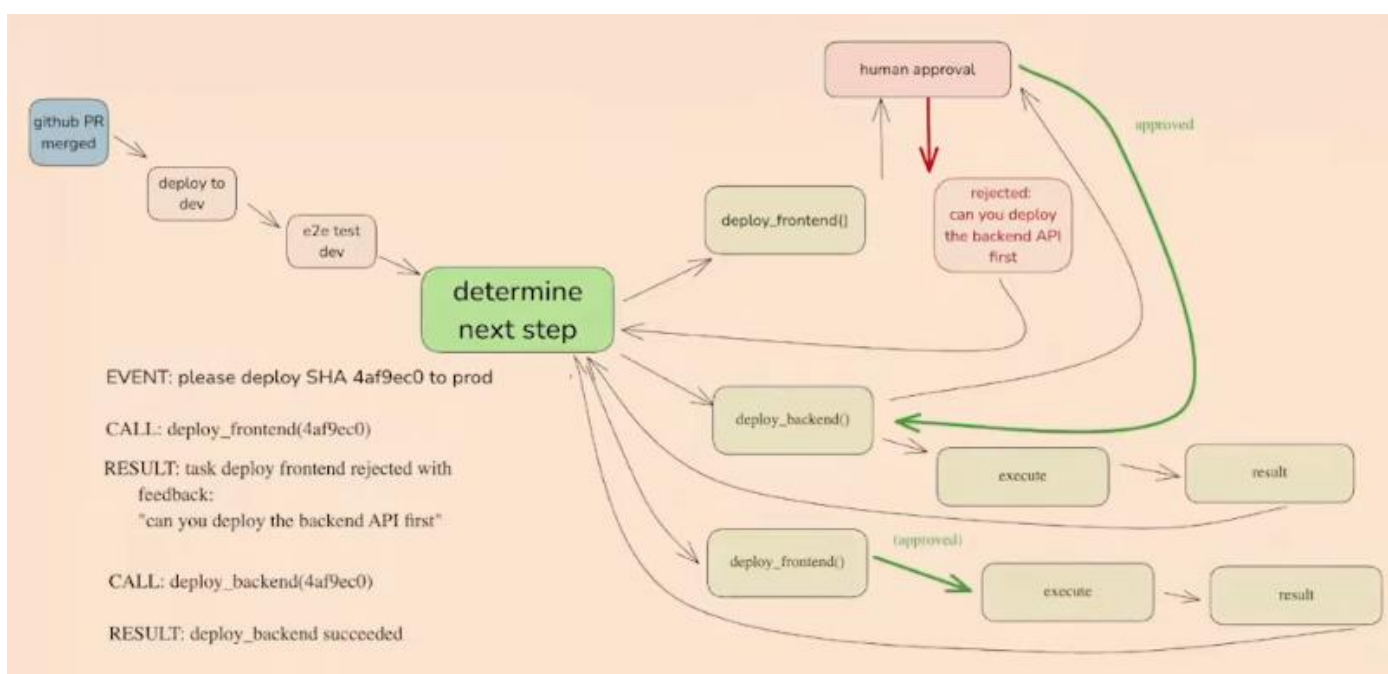
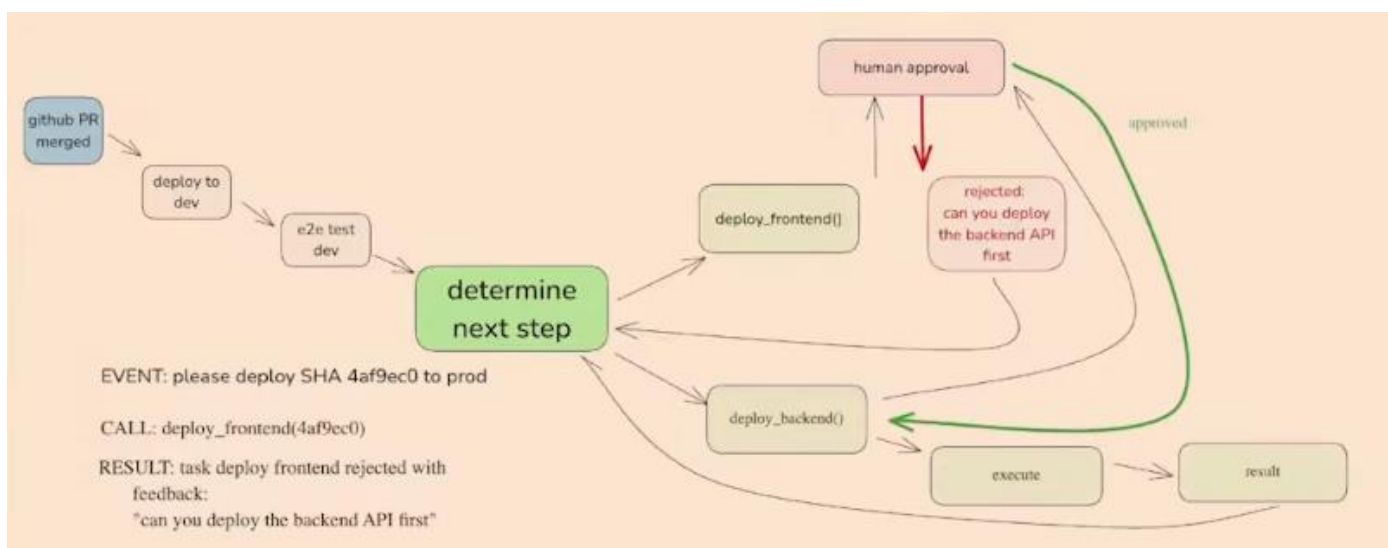
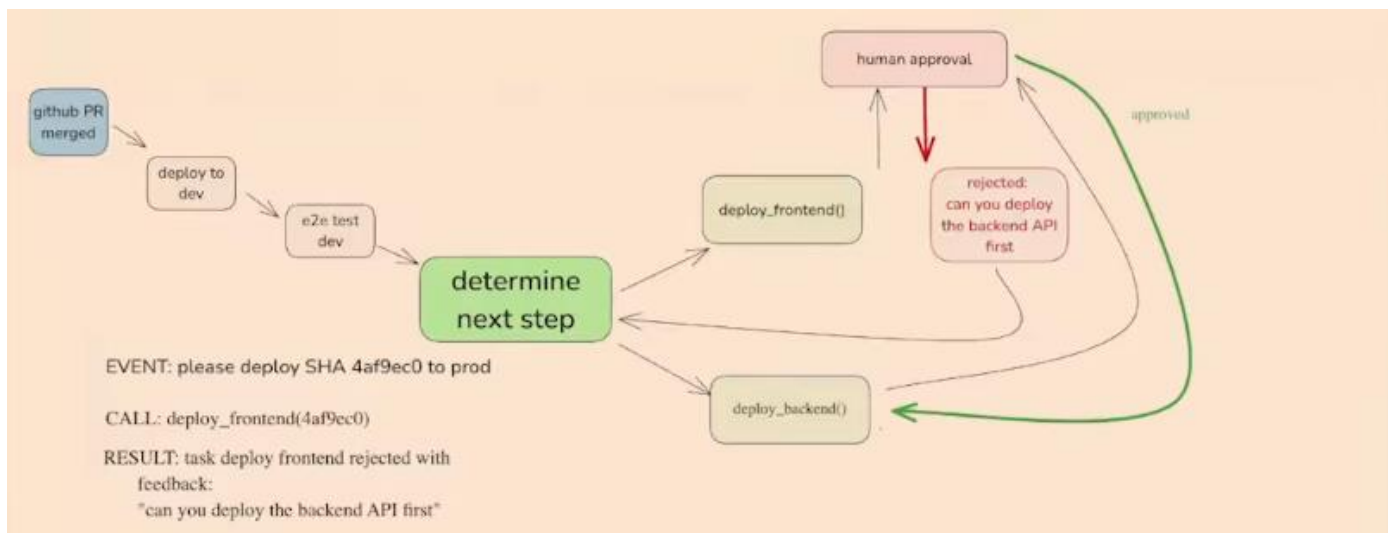


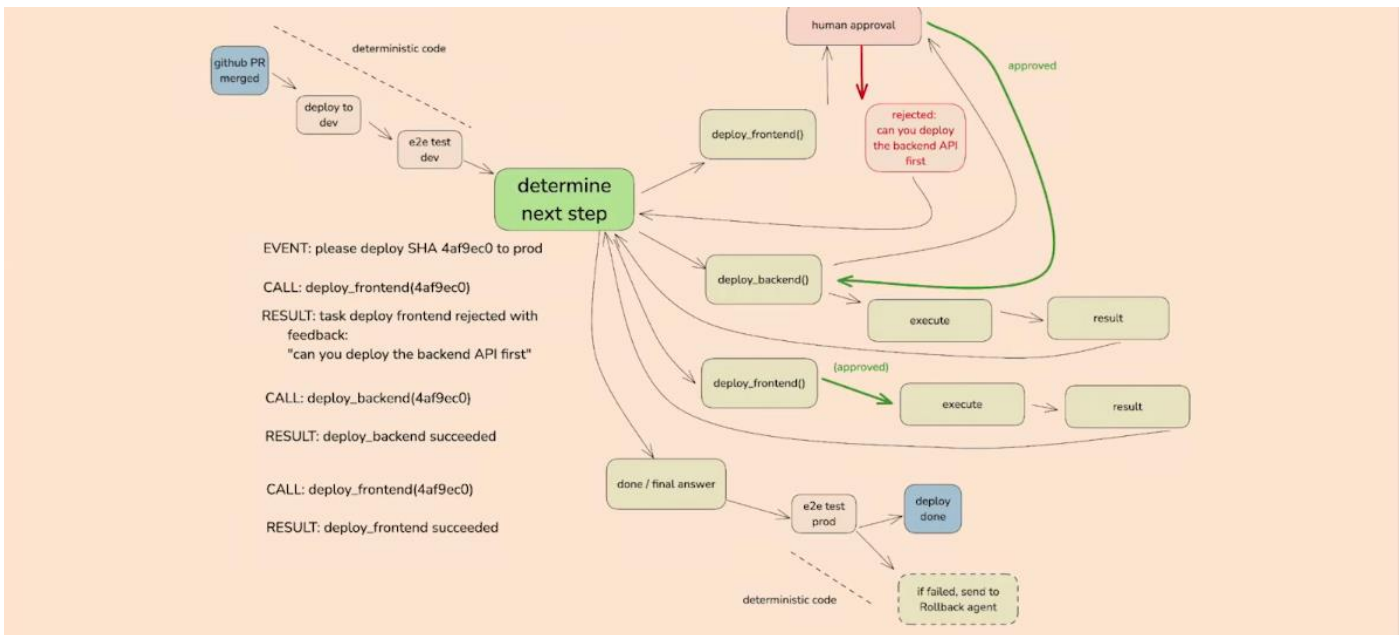
EVENT: please deploy SHA 4af9ec0 to prod

CALL: deploy\_frontend(4af9ec0)

RESULT: task deploy frontend rejected with







**HL-1 Deploybot** APR 9:33 PM  
 Agent `agent-6a5853e5` wants to call `promote_vercel_deployment`

function: `promote_vercel_deployment`  
 new\_deployment\_sha: ddec3f8  
 new\_deployment:  
 Commit ddec3f8 by dexhorthy: 'Merge pull request #1004 from balanceiskey/sundeeep/eng-1074-' - View Deployment  
 previous\_deployment:  
 Commit 3fc12af by dexhorthy: 'Merge pull request #1005 from balanceiskey/sundeeep/eng-1032-' - View Deployment

✗ Rejected by: @dex at 04:36 AM UTC  
 Reason: can you deploy the API first

Agent `agent-325ea819` wants to call `tag_push_prod`

function: `tag_push_prod`

---

sha\_to\_deploy:  
 ddec3f83f4620338ab7461007330c8cc9308e4ed  
 new\_commit:  
 Commit ddec3f83 by dexhorthy: 'Merge pull request #1004 from balanceiskey/sundeeep/eng-1074-form-updates-for-byo-slack'  
 previous\_commit:  
 Commit d590591 by dexhorthy: 'Update tag-and-push-prod.yaml'

✓ This request was approved by @dex at 04:46 AM UTC

**HL-1 Deploybot** APR 9:58 PM  
 Agent `agent-b2f3d659` wants to call `promote_vercel_deployment`

## Small Focused Agents

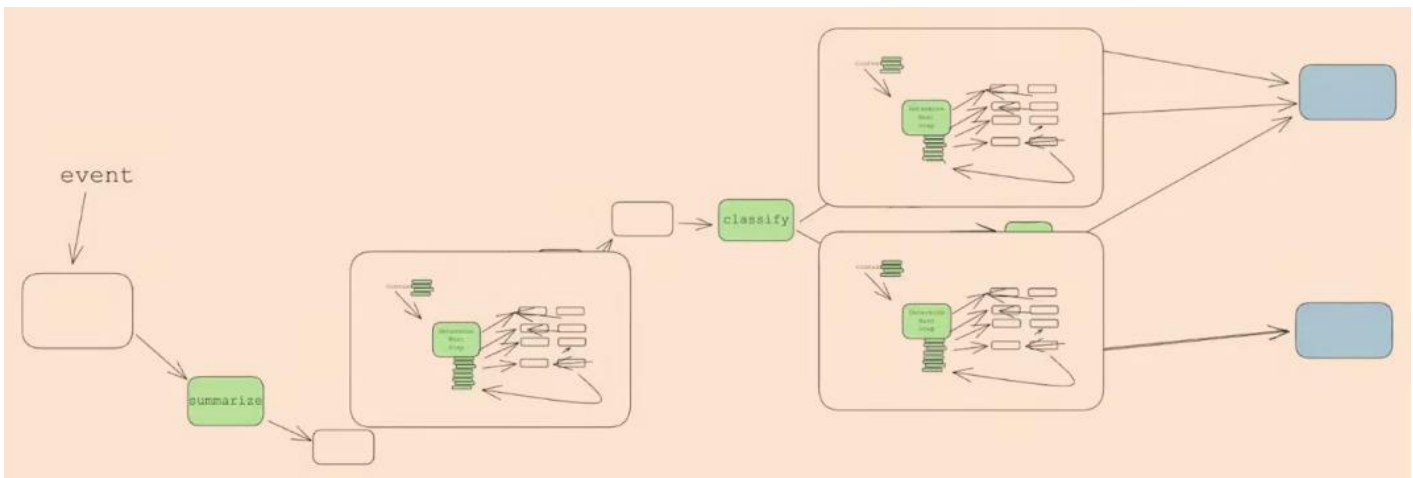
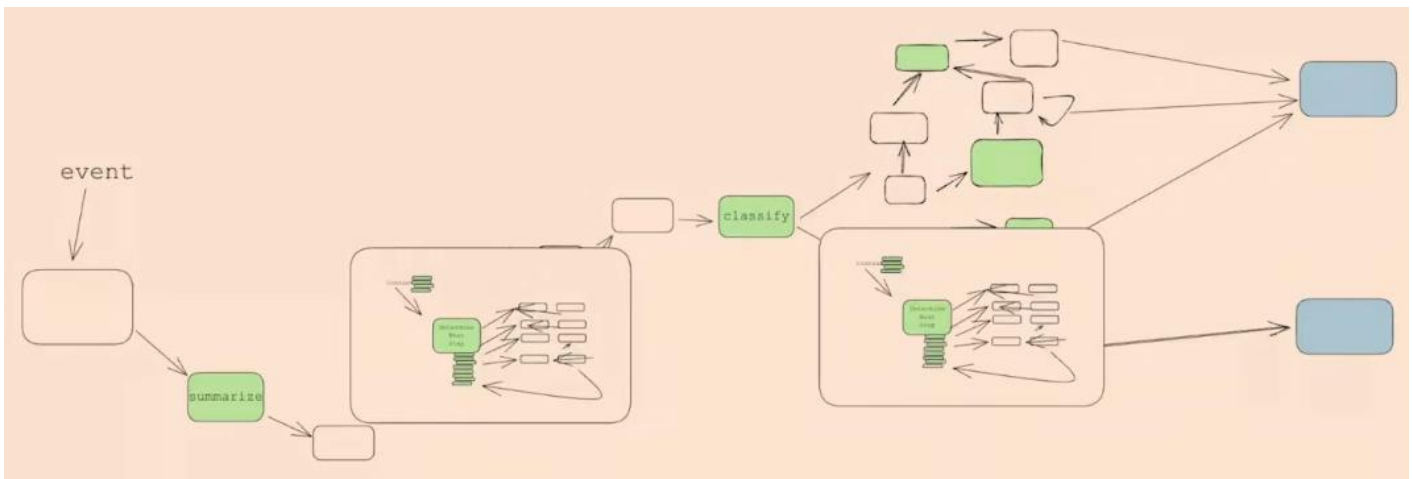
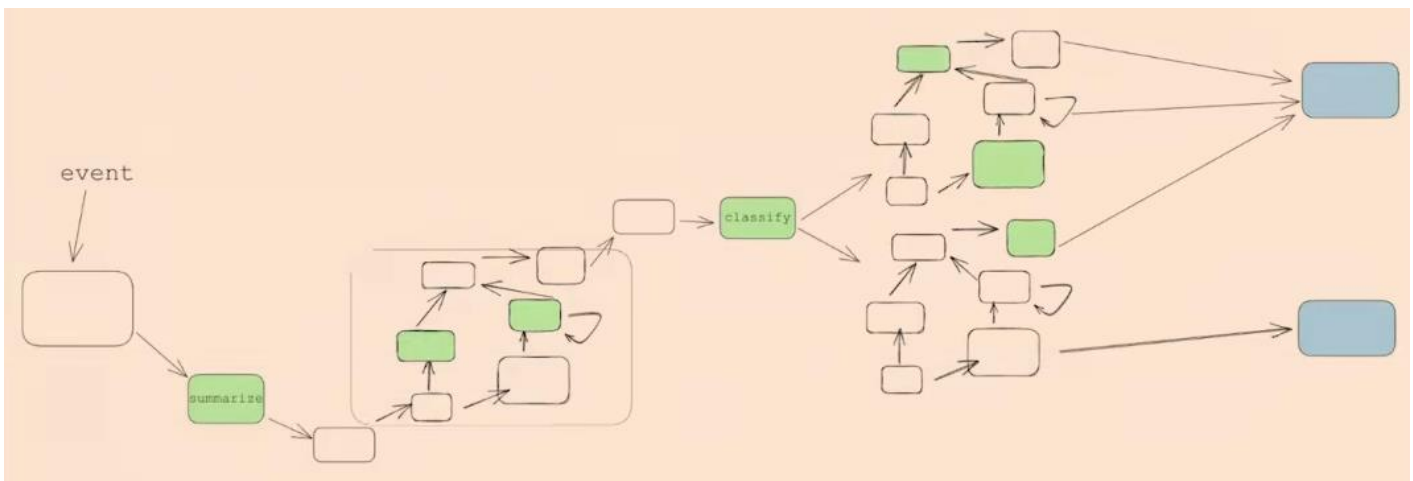
~100 tools | ~20 steps

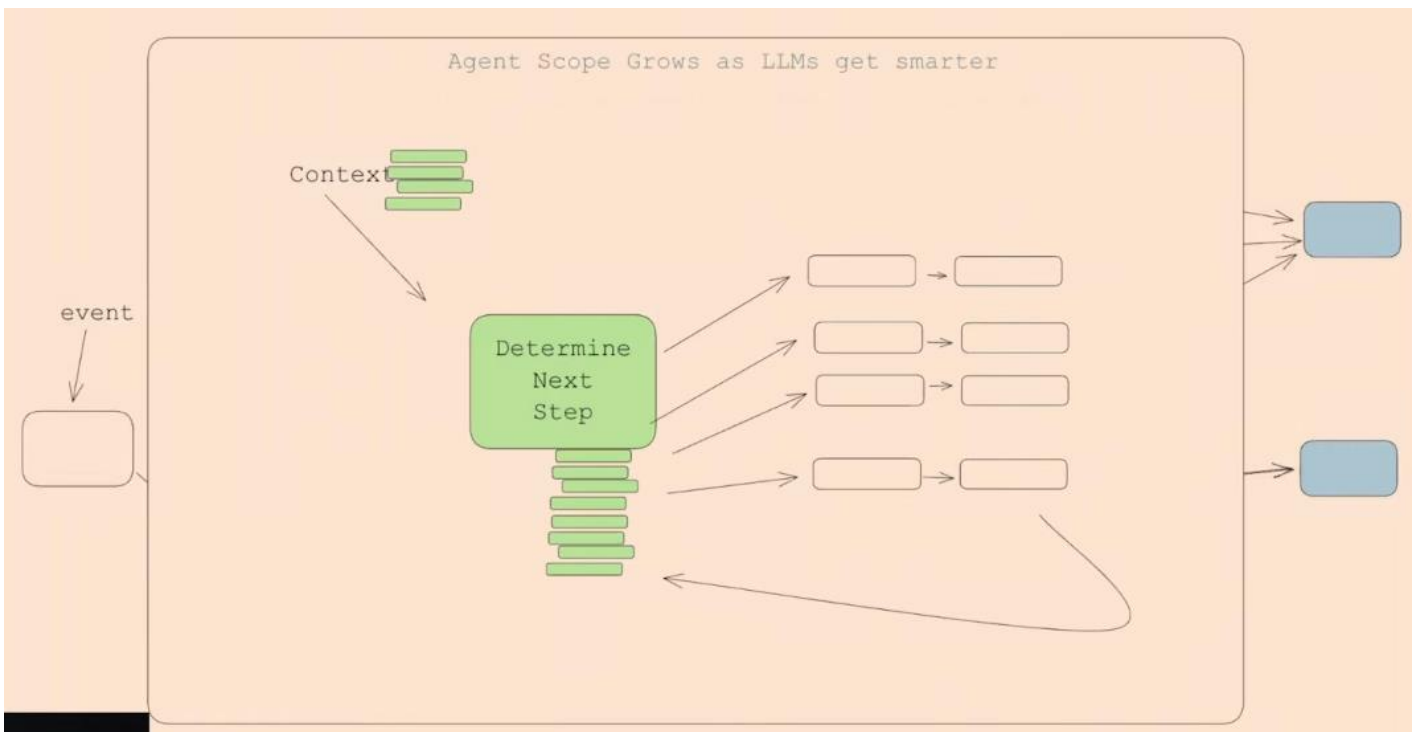
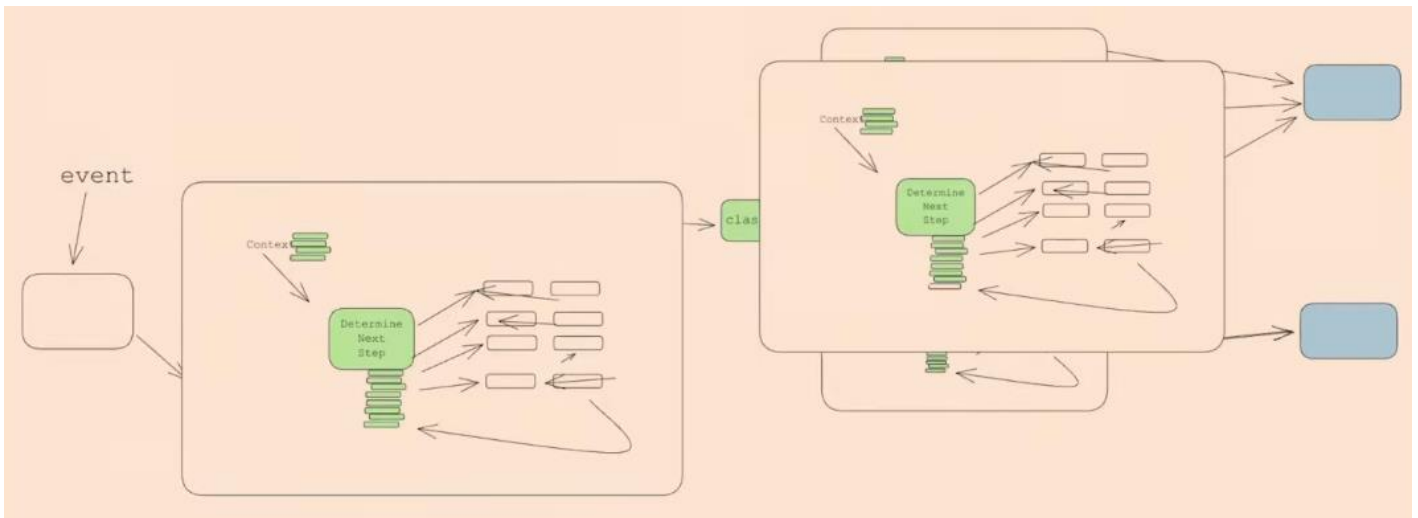

Manageable Context  
 Clear Responsibilities  
 More Reliability  
 Easier Testing / Debugging

What if LLMs keep getting smarter?

*"I feel like consistently, the most magical moments out of AI building come about for me when I'm really, really, really just close to the edge of the model capability"*


– Usama Bin Shafqat  
NotebookLM Team  
[latent space](#)





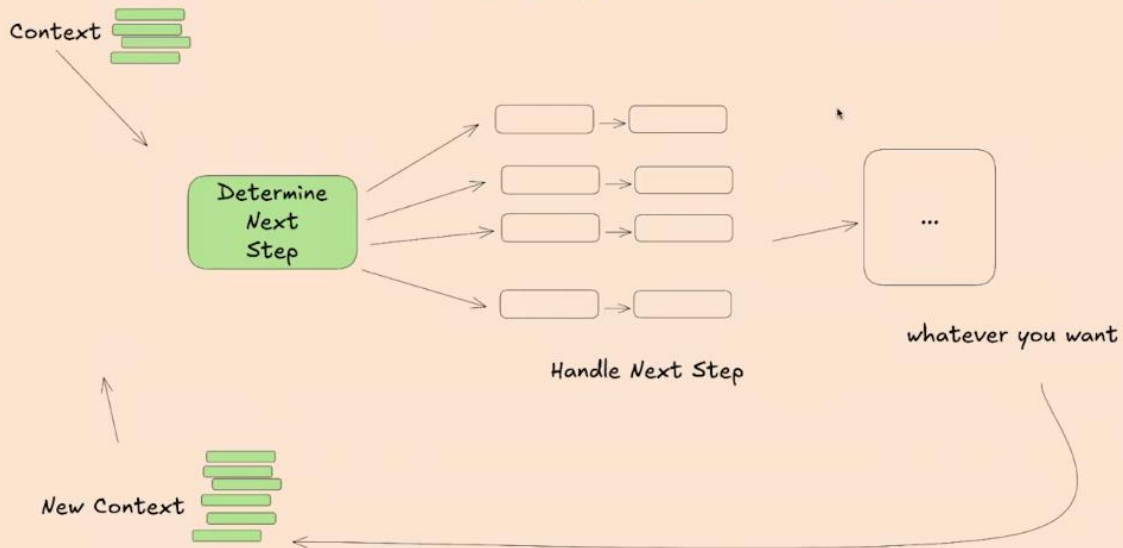


## Factor 10 - Small, Focused Agents



## Factor 12 On stateless reducers

### Factor 12 - Make your agent a stateless reducer!



We're still finding the right abstractions

```
$ npx create-12-factor-agent
```

Agents don't need bootstrap...

...they need shadcn



Agent-to-Human (A2H) protocol v0

