

## DeepSeek-R1 with LangGraph - OpenSource Fallback Mechanism



Coding Crash Courses  
16.3K subscribers

Subscribe

68



Share

Clip

Save



2,245 views Jan 25, 2025

In this Video you learn the following:

- How to use DeepSeek with Ollama
- How to integrate it into LangGraph
- How to switch from openai to open source model dynamically
- and more :)

Code: <https://github.com/Coding-Crashcourse/...>

Timestamps:

0:00 Introduction

0:35 Ollama Installation

2:00 DeepSeek in LangGraph

4:40 Fallback (Modelswitcher)

https://github.com/Coding-Crashcourse/LangGraph-Tutorial

Coding-Crashcourse / LangGraph-Tutorial

Code Issues 4 Pull requests Actions Projects Security Insights

LangGraph-Tutorial

Public

Watch 4

Fork 65

Star 142

main 1 Branch 0 Tags

Go to file

Add file

Code

About

No description, website, or topics provided.

Readme

Activity

142 stars

4 watching

65 forks

Report repository

Releases

No releases published

Packages

No packages published

Languages

Jupyter Notebook 95.3% Python 4.6% Dockerfile 0.1%

## LangGraph + DeepSeek R1

```
ChatOllama(model="deepseek-r1:7b", temperature=0)
```

## LangGraph + DeepSeek R1

```
ollama_config = {  
    "configurable": {  
        "model_type": "ollama"  
    }  
}  
graph.invoke({"question": "What's the highest mountain in the world?"}, ollama_config)
```

# LangGraph + DeepSeek R1

```
import re
import time
from langgraph.graph import StateGraph

class ModelSwitcher:
    def __init__(self, graph: StateGraph):
        self.graph = graph
        self.last_openai_fail_time = None
        ...

model_switcher.invoke("Which city is the capital of France?")

-----
OpenAI is still in cooldown. Time until OpenAI is active again: 280.01 seconds.
Using fallback (Ollama).
Using Ollama (deepseek-r1:7b).
```

## LangGraph + DeepSeek R1

**Ollama is the easiest way  
to setup DeepSeek R1**

```
File Edit Selection View Go Run ... LangGraph-Tutorial
fallback_opensource.ipynb M testscript.py langgraph.ipynb
from dotenv import load_dotenv

load_dotenv()

[1] ✓ 0.0s Python
... True

from typing_extensions import TypedDict

from langgraph.graph import StateGraph, START, END
from langchain_core.runnables.config import RunnableConfig
from langchain_core.messages import HumanMessage
from langchain_openai import ChatOpenAI
from langchain_ollama import ChatOllama

class AgentState(TypedDict):
    question: str
    answer: str

def call_model(state: AgentState, config: RunnableConfig) -> AgentState:
    model_type = config["configurable"].get("model_type", "openai")
    if model_type == "ollama":
        print("Using Ollama (deepseek-r1:7b).")
        llm = ChatOllama(model="deepseek-r1:7b", temperature=0)
    else:
        print("Using OpenAI (gpt-4o-mini).")
        llm = ChatOpenAI(model_name="gpt-4o-mini", temperature=0)

    messages = [HumanMessage(content=state["question"])]
```

```
File Edit Selection View Go Run ... ← → LangGraph-Tutorial
EXPLORER
LANGGRAPH-TUTORIAL
> .venv
> app
> fastapi_app
> sql
> env
.gitignore
agent_team.ipynb
code.ipynb
crag_dynamic_models.ipynb
crag.ipynb
custom_persistence.ipynb
customer_support.ipynb
docker-compose.yaml
fallback_opensource.ipynb M
human_in_loop.ipynb
langgraph.ipynb
lcel_vs_langgraph.ipynb
mapreduce.ipynb
README.md
requirements.txt
testscript.py

+ Code + Markdown | Run All | Restart | Clear All Outputs | Jupyter Variables | Outline ...
Python 3.11.0

from typing_extensions import TypedDict
from langgraph.graph import StateGraph, START, END
from langchain_core.runnables.config import RunnableConfig
from langchain_core.messages import HumanMessage
from langchain_openai import ChatOpenAI
from langchain_ollama import ChatOllama

class AgentState(TypedDict):
    question: str
    answer: str

def call_model(state: AgentState, config: RunnableConfig) -> AgentState:
    model_type = config["configurable"].get("model_type", "openai")
    if model_type == "ollama":
        print("Using Ollama (deepseek-r1:7b).")
        llm = ChatOllama(model="deepseek-r1:7b", temperature=0)
    else:
        print("Using OpenAI (gpt-4o-mini).")
        llm = ChatOpenAI(model_name="gpt-4o-mini", temperature=0)

    messages = [HumanMessage(content=state["question"])]
    response = llm.invoke(messages)
    state["answer"] = response.content
    return state

workflow = StateGraph(AgentState)
workflow.add_edge(START, "agent")
workflow.add_node("agent", call_model)
workflow.add_edge("agent", END)
graph = workflow.compile()

[2] Python
```

```
File Edit Selection View Go Run ... ← → LangGraph-Tutorial
EXPLORER
LANGGRAPH-TUTORIAL
> .venv
> app
> fastapi_app
> sql
> env
.gitignore
agent_team.ipynb
code.ipynb
crag_dynamic_models.ipynb
crag.ipynb
custom_persistence.ipynb
customer_support.ipynb
docker-compose.yaml
fallback_opensource.ipynb M
human_in_loop.ipynb
langgraph.ipynb
lcel_vs_langgraph.ipynb
mapreduce.ipynb
README.md
requirements.txt
testscript.py

+ Code + Markdown | Interrupt | Restart | Clear All Outputs | Go To | Jupyter Variables | Outline ...
Python 3.11.0

ollama_config = {
    "configurable": {
        "model_type": "openai"
    }
}

llm = ChatOpenAI(model_name="gpt-4o-mini", temperature=0)

messages = [HumanMessage(content=state["question"])]
response = llm.invoke(messages)
state["answer"] = response.content
return state

workflow = StateGraph(AgentState)
workflow.add_edge(START, "agent")
workflow.add_node("agent", call_model)
workflow.add_edge("agent", END)
graph = workflow.compile()

graph.invoke({"question": "What's the highest mountain in the world?"}, ollama_config)

... Using OpenAI (gpt-4o-mini).

import re

[2] ✓ 2.8s Python
[4] 1.1s Python
```

The screenshot shows a VS Code editor with a Jupyter notebook open. The notebook contains the following Python code:

```
ollama_config = {  
    llm = ChatOpenAI(model_name="gpt-4o-mini", temperature=0)  
  
    messages = [HumanMessage(content=state["question"])]  
    response = llm.invoke(messages)  
    state["answer"] = response.content  
    return state  
  
    workflow = StateGraph(AgentState)  
    workflow.add_edge(START, "agent")  
    workflow.add_node("agent", call_model)  
    workflow.add_edge("agent", END)  
    graph = workflow.compile()  
    graph.invoke({"question": "What's the highest mountain in the world?"}, ollama_config)
```

The output of the notebook shows the following text:

```
Using Ollama (deepseek-r1:7b).  
  
{'question': 'What's the highest mountain in the world?',  
 'answer': '<think>\n\n</think>\n\nThe highest mountain in the world is Mount Everest.'}
```

We can now use OpenAI and DeepSeek as the fallback as below

The screenshot shows a VS Code editor with a Jupyter notebook open. The notebook contains the following Python code:

```
import re  
import time  
from langgraph.graph import StateGraph  
  
class ModelSwitcher:  
    def __init__(self, graph: StateGraph):  
        self.graph = graph  
        self.last_openai_fail_time = None  
        self.openai_config = {  
            "configurable": {  
                "model_type": "openai",  
            }  
        }  
        self.fallback_config = {  
            "configurable": {  
                "model_type": "ollama",  
            }  
        }  
  
    def invoke(self, question: str, remove_think: bool = True) -> str:  
        if self._should_skip_openai():  
            return self._invoke_fallback(question, remove_think)  
  
        if question.lower() == "force error":  
            print("Forcing error with ChatOpenAI.")
```



```
import re

model_type": "openai",
}

self.fallback_config = {
    "configurable": {
        "model_type": "ollama",
    }
}

def invoke(self, question: str, remove_think: bool = True) -> str:
    if self._should_skip_openai():
        return self._invoke_fallback(question, remove_think)

    if question.lower() == "force error":
        print("Forcing error with ChatOpenAI.")
        self.last_openai_fail_time = time.time()
        print("OpenAI is disabled for 5 minutes. Invoking fallback.")
        return self._invoke_fallback(question, remove_think)

    try:
        print("Attempting invocation with OpenAI...")
        result = self.graph.invoke({"question": question}, self.openai_config)
        answer = result["answer"]
        return self._clean_if_needed(answer, remove_think)
    except Exception as e:
        print("Error with ChatOpenAI occurred:", e)
        self.last_openai_fail_time = time.time()
        print("OpenAI is disabled for 5 minutes. Invoking fallback.")
        return self._invoke_fallback(question, remove_think)

def _invoke_fallback(self, question: str, remove_think: bool) -> str:
    print("Using fallback (Ollama).")
```

```
except Exception as e:
    print("Error with ChatOpenAI occurred:", e)
    self.last_openai_fail_time = time.time()
    print("OpenAI is disabled for 5 minutes. Invoking fallback.")
    return self._invoke_fallback(question, remove_think)

def _invoke_fallback(self, question: str, remove_think: bool) -> str:
    print("Using fallback (Ollama).")
    result = self.graph.invoke({"question": question}, self.fallback_config)
    answer = result["answer"]
    return self._clean_if_needed(answer, remove_think)

def _should_skip_openai(self) -> bool:
    if self.last_openai_fail_time is None:
        return False
    elapsed = time.time() - self.last_openai_fail_time
    if elapsed < 300:
        remaining = 300 - elapsed
        print(f"OpenAI is still in cooldown. Time until OpenAI is active again: {remaining:.2f} seconds.")
        return True
    return False

def _clean_if_needed(self, text: str, remove_think: bool) -> str:
    if not remove_think:
        return text
    return self._remove_thinking_tokens(text)

def _remove_thinking_tokens(self, text: str) -> str:
    pattern = r"<think>.*?</think>"
    text_no_think = re.sub(pattern, "", text, flags=re.DOTALL)
    return text_no_think.lstrip("\n")
```

```
def _clean_if_needed(self, text: str, remove_think: bool) -> str:
    if not remove_think:
        return text
    return self._remove_thinking_tokens(text)

def _remove_thinking_tokens(self, text: str) -> str:
    pattern = r"<think>.*</think>"
    text_no_think = re.sub(pattern, "", text, flags=re.DOTALL)
    return text_no_think.lstrip("\n")
```

```
model_switcher = ModelSwitcher(graph)
```

```
model_switcher.invoke("What's the highest mountain in the world?")
```

Attempting invocation with OpenAI...  
Using OpenAI (gpt-4o-mini).

'The highest mountain in the world is Mount Everest, which stands at an elevation of 8,848.86 meters (29,031.7 feet) above sea lev

```
model_switcher.invoke("force error")
```

```
text_no_think = re.sub(pattern, "", text, flags=re.DOTALL)
return text_no_think.lstrip("\n")
```

```
model_switcher = ModelSwitcher(graph)
```

```
model_switcher.invoke("What's the highest mountain in the world?")
```

Attempting invocation with OpenAI...  
Using OpenAI (gpt-4o-mini).

'The highest mountain in the world is Mount Everest, which stands at an elevation of 8,848.86 meters (29,031.7 feet) above sea lev

```
model_switcher.invoke("force error")
```

Forcing error with ChatOpenAI.  
OpenAI is disabled for 5 minutes. Invoking fallback.  
Using fallback (Ollama).  
Using Ollama (deepseek-r1:7b).

'It seems like you\'re asking about how to fix an error, but you didn\'t specify the type of error. Could you please provide more

```
File Edit Selection View Go Run ... LangGraph-Tutorial
EXPLORER
LANGGRAPH-TUTORIAL
  .env
  app
  fastapi_app
  sql
  agent_team.ipynb
  code.ipynb
  crag_dynamic_models.ipynb
  crag.ipynb
  custom_persistence.ipynb
  customer_support.ipynb
  docker-compose.yaml
  fallback_opensource.ipynb M
  human_in_loop.ipynb
  langgraph.ipynb
  lcel_vs_langgraph.ipynb
  mapreduce.ipynb
  README.md
  requirements.txt
  testscript.py
  OUTLINE
  TIMELINE
main* 0 0 16 0

[8] ✓ 1.9s Python
...
model_switcher.invoke("Which city is the capital of France?")
...
Attempting invocation with OpenAI...
Using OpenAI (gpt-4o-mini).
...
'The highest mountain in the world is Mount Everest, which stands at an elevation of 8,848.86 meters (29,031.7 feet) above sea level'

[9] ✓ 5.8s Python
...
model_switcher.invoke("force error")
...
Forcing error with ChatOpenAI.
OpenAI is disabled for 5 minutes. Invoking fallback.
Using fallback (Ollama).
Using Ollama (deepseek-r1:7b).
...
'It seems like you\'re asking about how to fix an error, but you didn\'t specify the type of error. Could you please provide more

[10] ✓ 0.5s Python
...
model_switcher.invoke("Which city is the capital of France?")
...
OpenAI is still in cooldown. Time until OpenAI is active again: 280.01 seconds.
Using fallback (Ollama).
Using Ollama (deepseek-r1:7b).
...
'The capital of France is Paris.'
```

```
File Edit Selection View Go Run ... LangGraph-Tutorial
EXPLORER
LANGGRAPH-TUTORIAL
  .env
  app
  fastapi_app
  sql
  agent_team.ipynb
  code.ipynb
  crag_dynamic_models.ipynb
  crag.ipynb
  custom_persistence.ipynb
  customer_support.ipynb
  docker-compose.yaml
  fallback_opensource.ipynb M
  human_in_loop.ipynb
  langgraph.ipynb
  lcel_vs_langgraph.ipynb
  mapreduce.ipynb
  README.md
  requirements.txt
  testscript.py
  OUTLINE
  TIMELINE
main* 0 0 16 0

[10] ✓ 0.6s Python
...
model_switcher.invoke("Which city is the capital of France?")
...
OpenAI is still in cooldown. Time until OpenAI is active again: 280.01 seconds.
Using fallback (Ollama).
Using Ollama (deepseek-r1:7b).
...
'The capital of France is Paris.'
```

```
[11] ✓ 0.3s Python
...
model_switcher.invoke("Which city is the capital of France?")
...
OpenAI is still in cooldown. Time until OpenAI is active again: 267.49 seconds.
Using fallback (Ollama).
Using Ollama (deepseek-r1:7b).
...
'The capital of France is Paris.'
```

```
[12] ✓ 0.3s Python
...
model_switcher.invoke("Which city is the capital of France?", remove_think=False)
...
OpenAI is still in cooldown. Time until OpenAI is active again: 249.42 seconds.
Using fallback (Ollama).
Using Ollama (deepseek-r1:7b).
...
'<think>\n\n</think>\n\nThe capital of France is Paris.'
```