

LangGraph vs LangChain vs LangFlow vs LangSmith : Which One To Use & Why?



FuturMinds
11.4K subscribers

Subscribe

2.5K



Share

Save

Download



70,256 views Aug 21, 2024

Are you curious about the differences between LangChain, Langflow, Langgraph, and Langsmith? In this video, I'll explain what is LangChain, what is Langflow, what is Langgraph, and what is Langsmith, breaking down each tool's features, use cases, and why they were created.

We'll explore what Langsmith is used for and how it integrates with LangChain for monitoring and debugging workflows. Learn how Langgraph simplifies agent management and discover the differences in Langchain vs Langgraph and Langgraph vs Langchain for building AI applications. I also cover the Langchain vs Langflow comparison, showing how Langflow's visual interface can speed up prototyping for non-coders.

If you're wondering about Langchain vs Langsmith, this video also breaks down how Langsmith helps in tracking and optimizing your applications. Whether you're a developer or just starting with AI workflows, this video will clear up the confusion around these tools, helping you understand which one to choose for your next project.

★ Timestamps:

00:00 Intro

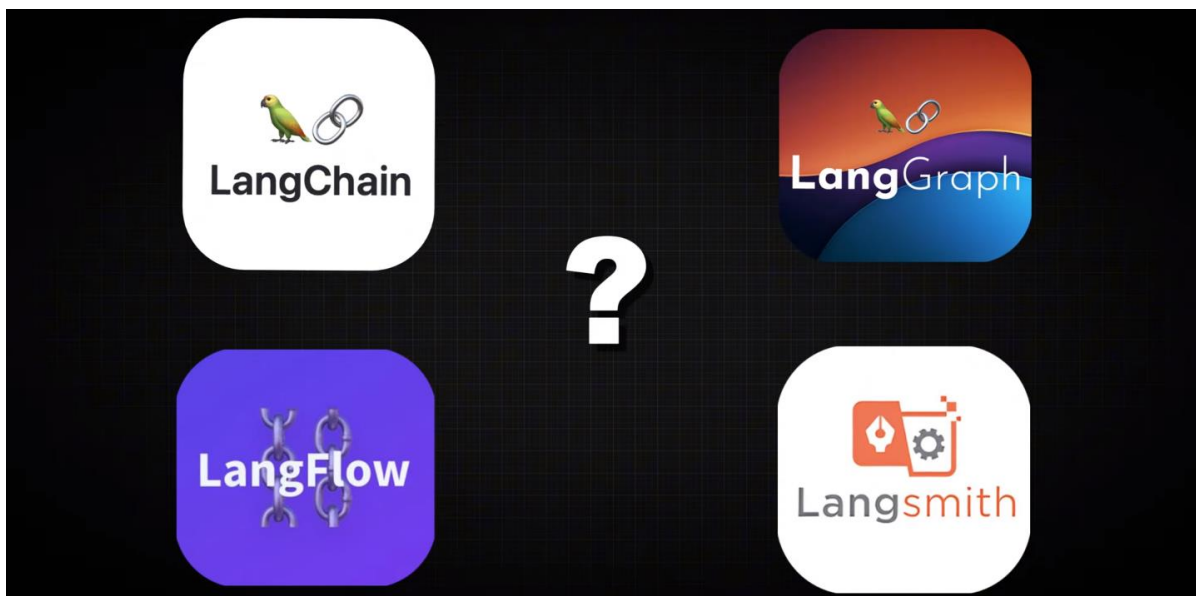
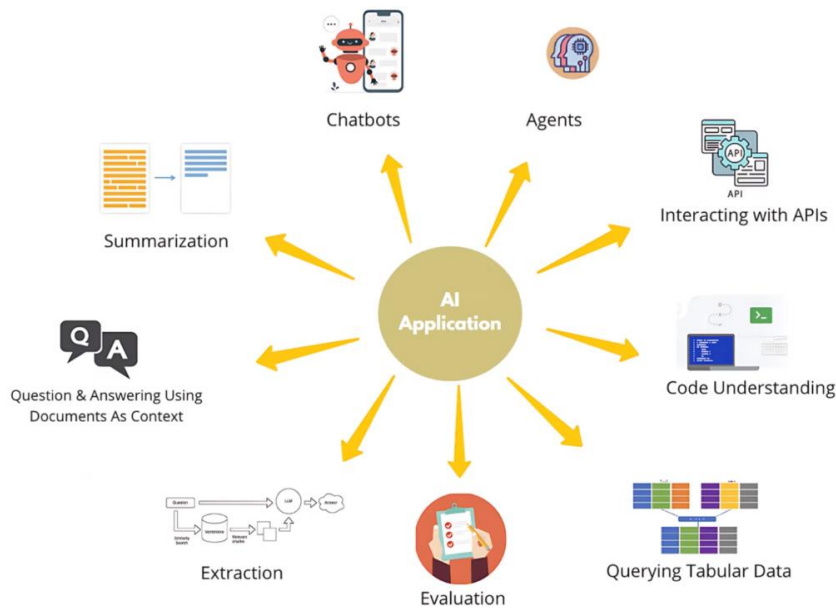
00:40 LangChain (What, Why, How)

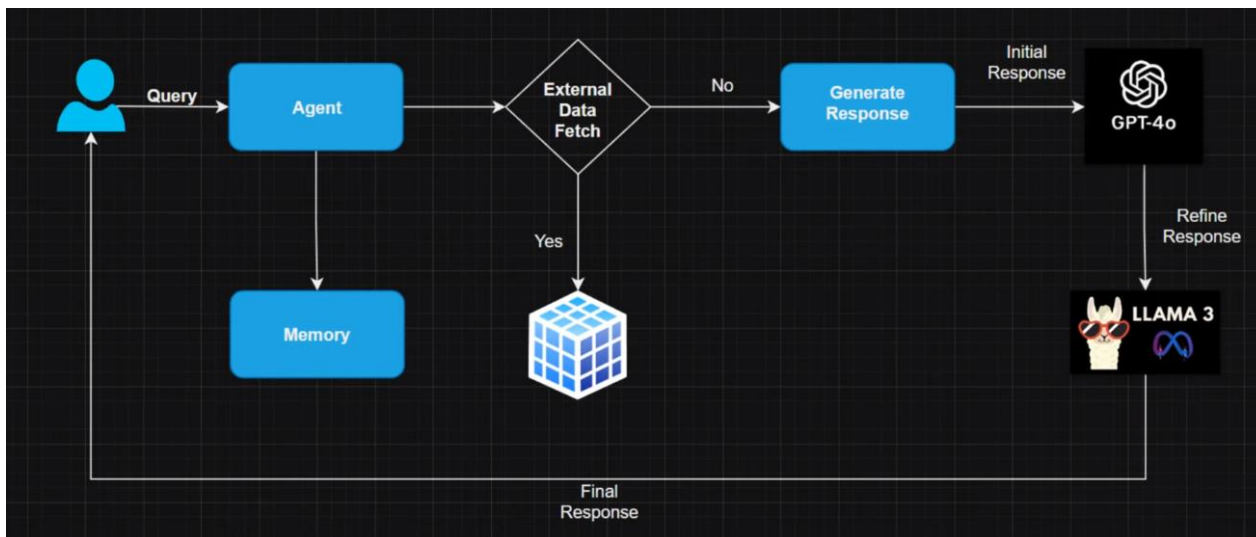
03:40 LangGraph (What, Why, How)

05:37 LangFlow (What, Why, How)

07:11 LangSmith (What, Why, How)

14:10 What's Next?





```
import openai
import requests

# GPT-4 API function
def query_gpt4(prompt):
    openai.api_key = "your-gpt4-api-key"
    response = openai.Completion.create(
        engine="gpt-4",
        prompt=prompt,
        max_tokens=100
    )
    return response.choices[0].text

# Llama 3 API function
def query_llama3(prompt):
    url = "https://llama3-api-endpoint"
    headers = {
        "Authorization": "Bearer your-llama3-api-key",
        "Content-Type": "application/json"
    }
    data = {
        "prompt": prompt,
        "max_tokens": 100
    }
    response = requests.post(url, headers=headers, json=data)
    return response.json()["response"]
```

```

def query_llama3(prompt):
    url = "https://llama3-api-endpoint"
    headers = {
        "Authorization": "Bearer your-llama3-api-key",
        "Content-Type": "application/json"
    }
    data = {
        "prompt": prompt,
        "max_tokens": 100
    }
    response = requests.post(url, headers=headers, json=data)
    return response.json()["response"]

# Simulating Memory (you might use a file or database in reality)
memory = {}

# Manually managing memory
def update_memory(user_input, response):
    memory["last_interaction"] = {"input": user_input, "response": response}

def get_memory():
    return memory.get("last_interaction")

# Agent function: Decides which model to use
def agent(query):
    # Retrieve memory
    ..

```

```

        "prompt": prompt,
        "max_tokens": 100
    }
    response = requests.post(url, headers=headers, json=data)
    return response.json()["response"]

# Simulating Memory (you might use a file or database in reality)
memory = {}

# Manually managing memory
def update_memory(user_input, response):
    memory["last_interaction"] = {"input": user_input, "response": response}

def get_memory():
    return memory.get("last_interaction")

# Agent function: Decides which model to use
def agent(query):
    # Retrieve memory
    last_interaction = get_memory()

    # Logic for using GPT-4 or Llama 3 based on the query
    if "complex" in query:
        response = query_gpt4(query)
    else:
        response = query_llama3(query)

```

You can write the code to manually do all the above things as above or you can use a framework like LangGraph

Build

LangChain is a framework to build with LLMs by chaining interoperable components. LangGraph is the framework for building controllable agentic workflows.

Run

Deploy your LLM applications at scale with LangGraph Cloud, our infrastructure purpose-built for agents.

Manage

Debug, collaborate, test, and monitor your LLM app in LangSmith - whether it's built with a LangChain framework or not.



Build your app with LangChain

Build context-aware, reasoning applications with LangChain's flexible framework that leverages your company's data and APIs. Future-proof your application by making vendor optionality part of your LLM infrastructure design.

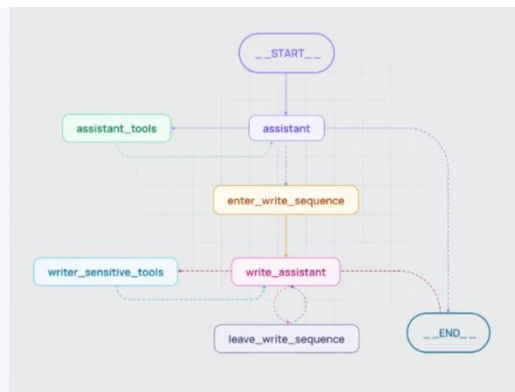
[Learn more about LangChain ↗](#)



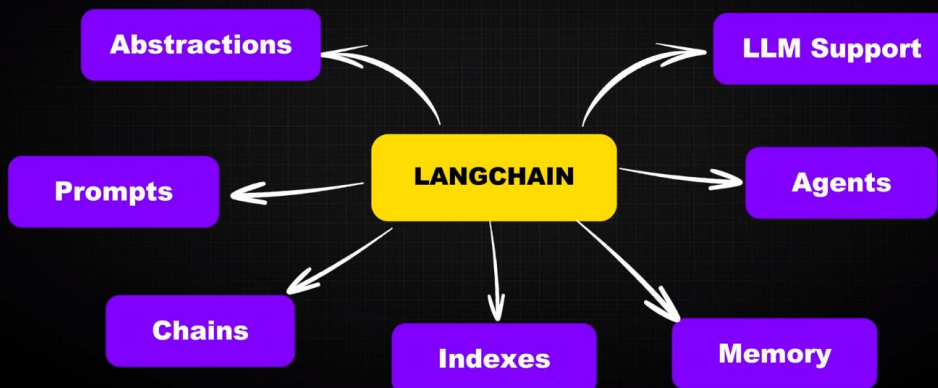
Run at scale with LangGraph Cloud

Deploy your LangGraph app with LangGraph Cloud for fault-tolerant scalability - including support for async background jobs, built-in persistence, and distributed task queues.

[Learn more about LangGraph ↗](#)



Langchain's Key Features



```

from langchain import OpenAI, LlamaCpp, Memory, Agent
from langchain.chains import SequentialChain

# Step 1: Define the models
gpt4_llm = OpenAI(api_key="your-gpt4-api-key", model="gpt-4")
llama3_llm = LlamaCpp(model_path="path-to-your-llama3-model")

# Step 2: Set up memory
memory = Memory()

# Step 3: Create an agent that decides whether to use GPT-4 or Llama 3
def agent_fn(query, memory):
    last_interaction = memory.load()

    # Simple agent decision: Use GPT-4 for complex queries
    if "complex" in query:
        return gpt4_llm(query)
    else:
        return llama3_llm(query)

agent = Agent(agent_fn, memory=memory)

# Step 4: Set up LangChain to manage the sequence of tasks
chain = SequentialChain([agent], memory=memory)

# Example usage

```



Products ▾

Methods ▾

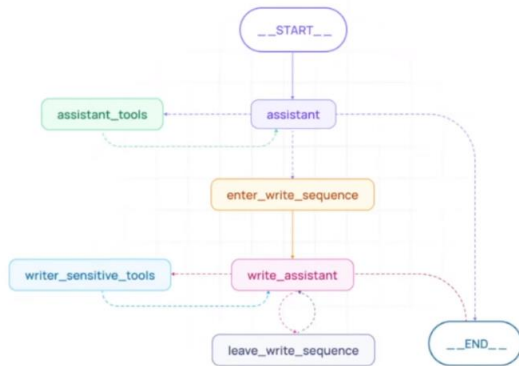
Resources ▾

Docs ▾

Company ▾

Pricing

Get a demo




Controllable cognitive architecture for any ta

LangGraph's flexible API supports diverse control flows - multi-agent, hierarchical, sequential - and robustly handles complex scenarios.

Ensure reliability with easy-to-add moderation and quality checks to prevent agents from veering off course.

[See the docs](#)



Products ▾

Methods ▾

Resources ▾

Docs ▾

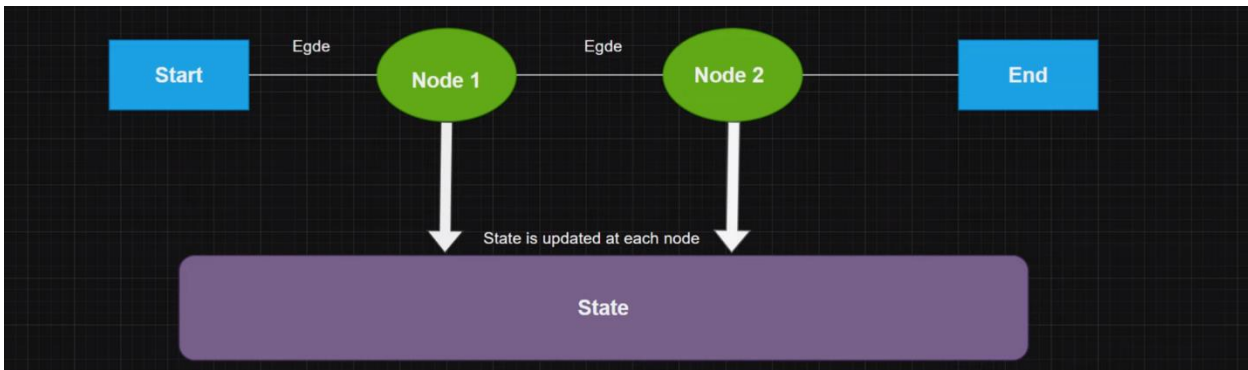
Company ▾

Pricing

Get a demo

Deploy agents at scale, monitor carefully, iterate boldly

With LangGraph Cloud, quickly deploy and scale your application, with infrastructure purpose-built for agents. [Now in beta](#) for all users on LangSmith Plus and Enterprise plans.



Installation

```
pip install -U langgraph
```

Example

One of the central concepts of LangGraph is state. Each graph execution creates a state that is passed between nodes in the graph as they execute, and each node updates this internal state with its return value after it executes. The way that the graph updates its internal state is defined by either the type of graph chosen or a custom function.

Let's take a look at a simple example of an agent that can use a search tool.

```
pip install langchain-anthropic
```

```
export ANTHROPIC_API_KEY=sk-...
```

Optionally, we can set up [LangSmith](#) for best-in-class observability.

```
export LANGSMITH_TRACING=true
export LANGSMITH_API_KEY=lsv2_sk...
```

```
from typing import Annotated, Literal, TypedDict

from langchain_core.messages import HumanMessage
from langchain_anthropic import ChatAnthropic
from langchain_core.tools import tool
from langgraph.checkpoint.memory import MemorySaver
from langgraph.graph import END, StateGraph, MessagesState
from langgraph.prebuilt import ToolNode
```

```
# Define the tools for the agent to use
@tool
def search(query: str):
    """Call to surf the web."""
    # This is a placeholder, but don't tell the LLM that...
    if "sf" in query.lower() or "san francisco" in query.lower():
        return "It's 60 degrees and foggy."
    return "It's 90 degrees and sunny."

tools = [search]

tool_node = ToolNode(tools)

model = ChatAnthropic(model="claude-3-5-sonnet-20240620", temperature=0).bind_tools(tools)

# Define the function that determines whether to continue or not
def should_continue(state: MessagesState) -> Literal["tools", END]:
    messages = state['messages']
    last_message = messages[-1]
    # If the LLM makes a tool call, then we route to the "tools" node
    if last_message.tool_calls:
```



Install Langflow Locally

⚠ CAUTION

Langflow **requires** Python version 3.10 or greater and [pip](#) or [pipx](#) to be installed on your system.

Install Langflow with pip:

```
1 python -m pip install langflow -U
```

Install Langflow with pipx:

```
1 pipx install langflow --python python3.10 --fetch-missing-python
```

Pipx can fetch the missing Python version for you with `--fetch-missing-python`, but you can also install the Python version manually. Use `--force-reinstall` to ensure you have the latest version of Langflow and its dependencies.

Having a problem?

If you encounter a problem, see [Common Installation Issues](#).

To get help in the Langflow CLI:

```
1 python -m langflow --help
```

🔧 Run Langflow

1. To run Langflow, enter the following command.

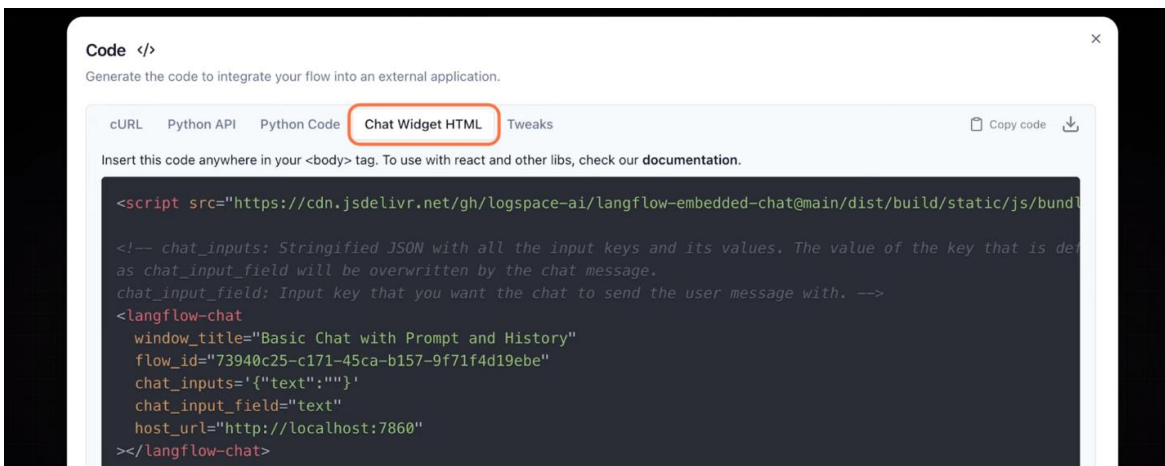
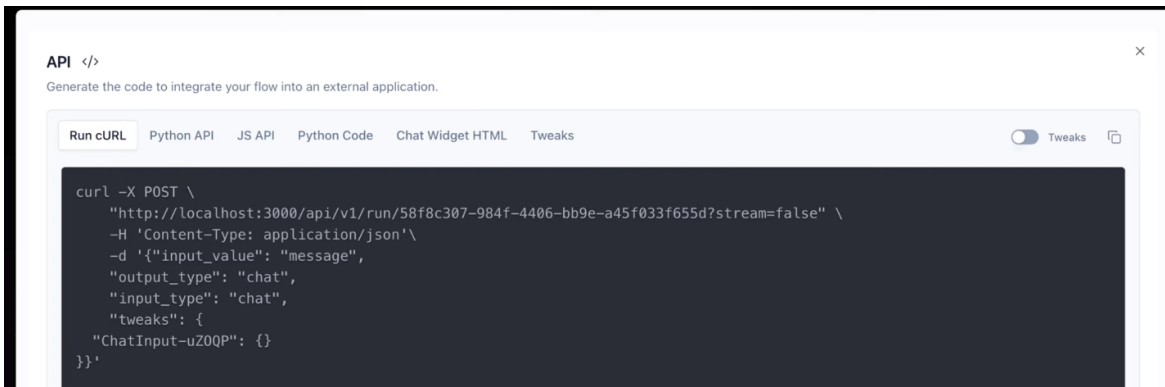
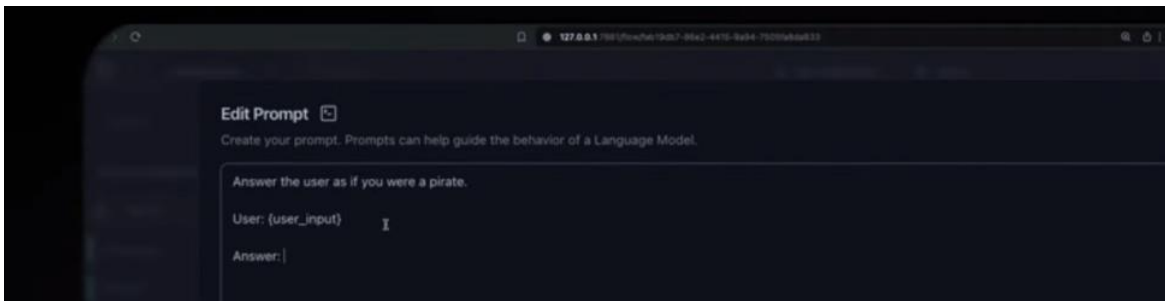
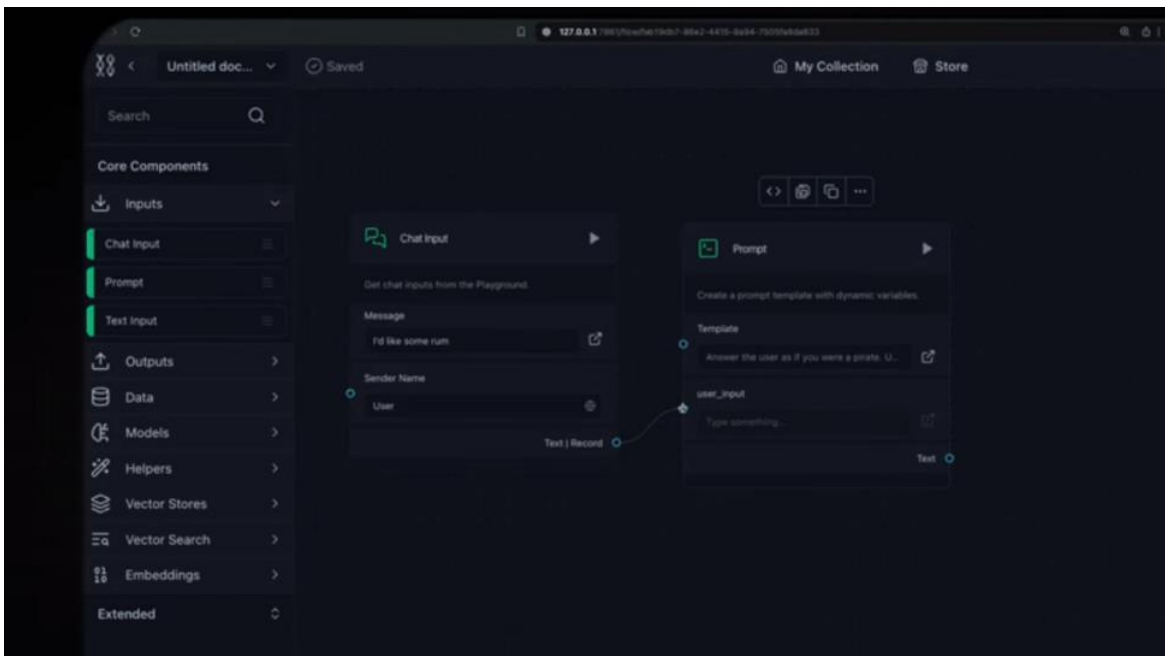
```
1 python -m langflow run
```

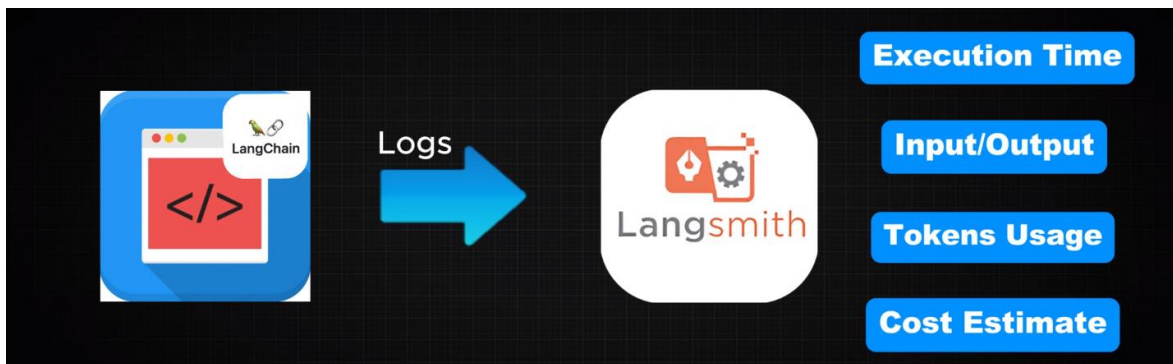
2. Confirm that a local Langflow instance starts by visiting `http://127.0.0.1:7860` in a Chromium-based browser.

Welcome to 🧩 Langflow

Collaborate, and contribute at our [GitHub Repo](#) 🚀

Access `http://127.0.0.1:7860`

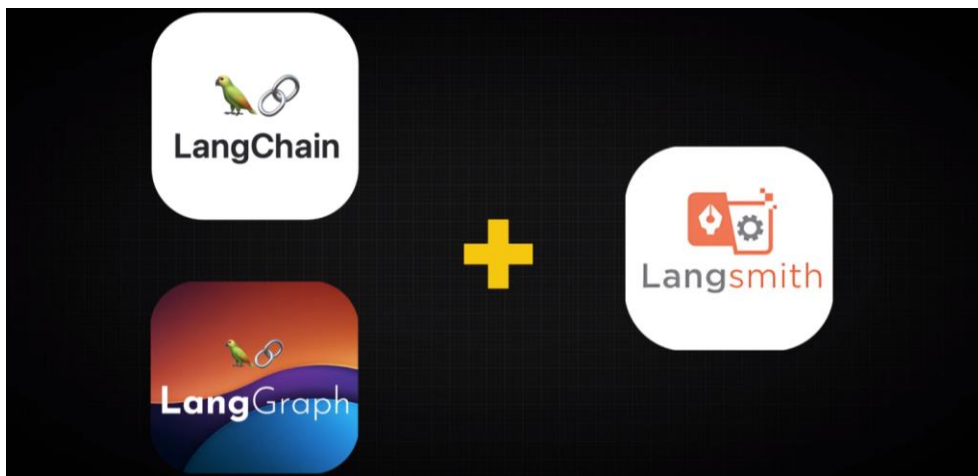




The platform for your LLM development lifecycle

LLM-apps are powerful, but have peculiar characteristics. The non-determinism, coupled with unpredictable, natural language inputs, make for countless ways the system can fall short. Traditional engineering best practices need to be re-imagined for working with LLMs, and LangSmith supports all phases of the development lifecycle.

[Sign Up](#)



Comparing Experiments

Input	rag-qa-gpt4-1106-d7a86cae ↗	rag-qa-gpt4o-2cc85e63 ↗
I'm passing {'a':1} and want to create an output map of {'a':1, 'b':2, 'c':3}. H...	In LCEL, you can achieve this by using the 'RunnablePassthrough.assign()' method to add additional ke... ANSWER_ACCURACY 0.5 → 10.68% SUCCESS 61.356 \$0.617	You can achieve this by using the 'RunnablePassthrough' an... ANSWER_ACCURACY 0.5 → 9.18% SUCCESS 66.259
With a RAG chain in LCEL, why are documents retrieved automatically f...	In the RAG chain example, when we construct 'chain = ({'context': retriever, 'question': RunnablePassth... ANSWER_ACCURACY 1 → 16.52% SUCCESS 61.512 \$0.621	When you construct the chain using {'context': retriever, 'qu... ANSWER_ACCURACY 1 → 16.59% SUCCESS 66.416
How can I configure the temperature of an LLM when invoking the LCEL...	In LCEL, you can configure the temperature of a large language model (LLM) such as GPT-3 using the '... ANSWER_ACCURACY 1 → 23.09% SUCCESS 61.552 \$0.623	You can configure the temperature of a language model (LL... ANSWER_ACCURACY 0.8 → 13.96% SUCCESS 66.276
How can I dynamically route logic based on input using a RunnableLam...	In LangChain, you can dynamically route logic based on input by creating a custom function that uses a... ANSWER_ACCURACY 1 → 21.12% SUCCESS 61.459 \$0.621	To dynamically route logic based on input using a 'Runnable... ANSWER_ACCURACY 1 → 23.05% SUCCESS 66.737
I am passing a map with {'num': 1} to a LCEL chain. How can I add an extr...	To add an extra key 'num2' that adds 1 to the value of 'num', and then assign the new map to a key nam... ANSWER_ACCURACY 0.8 → 17.64% SUCCESS 61.466 \$0.62	To achieve this, you can use the 'RunnablePassthrough.assi... ANSWER_ACCURACY 0.8 → 16.17% SUCCESS 66.201
I am passing text key 'foo' to my prompt and want to process it with a fu...	To process text with a custom function in LCEL before feeding it into a prompt, you can use a 'Runnable... ANSWER_ACCURACY 0.9 → 18.76% SUCCESS 61.441 \$0.619	You can achieve this by introducing a custom function using... ANSWER_ACCURACY 1 → 13.44% SUCCESS 66.267

Get started with LangSmith

LangSmith is a platform for building production-grade LLM applications. It allows you to closely monitor and evaluate your application, so you can ship quickly and with confidence. Use of LangChain is not necessary - LangSmith works on its own!

1. Install LangSmith

Python TypeScript

```
pip install -U langsmith
```

2. Create an API key

To create an API key head to the Settings page. Then click **Create API Key**.

3. Set up your environment

Shell

```
export LANGCHAIN_TRACING_V2=true
export LANGCHAIN_API_KEY=<your-api-key>

# The below examples use the OpenAI API, though it's not necessary in general
export OPENAI_API_KEY=<your-openai-api-key>
```

4. Log your first trace

💡 TRACING TO LANGSMITH FOR LANGCHAIN USERS

There is no need to use the LangSmith SDK directly if your application is built entirely on LangChain (either Python and JS).

We've outlined a tracing guide specifically for LangChain users [here](#).

We provide multiple ways to log traces to LangSmith. Below, we'll highlight how to use `traceable`. See more on the [Annotate code for tracing](#) page.

Python TypeScript

```
import openai
from langsmith.wrappers import wrap_openai
from langsmith import traceable

# Auto-trace LLM calls in-context
client = wrap_openai(openai.Client())

@traceable # Auto-trace this function
def pipeline(user_input: str):
    result = client.chat.completions.create(
        messages=[{"role": "user", "content": user_input}],
        model="gpt-3.5-turbo"
    )
    return result.choices[0].message.content

pipeline("Hello, world!")
# Out: Hello there! How can I assist you today?
```

2

0

langsmith-demo

P50 Latency: 13.43sP99 Latency: 18.56sTotal Tokens: 4,328

TracesAll RunsSetup

and(eq(name, "ChatOpenAI"), gte(latency, "10s"))

Columns

Run Type	Name	Input	Start Time	Latency	Tokens	Tags
LLM	ChatOpenAI	human: How di...	8/6/2023, 12:37:57 ...	16.11s	422	
LLM	ChatOpenAI	human: How d...	8/6/2023, 12:37:32 ...	17.51s	480	
LLM	ChatOpenAI	human: How d...	8/6/2023, 12:33:18 ...	18.70s	459	
LLM	ChatOpenAI	human: How di...	8/6/2023, 12:12:51 ...	11.23s	454	
LLM	ChatOpenAI	human: Tell me...	8/6/2023, 12:08:55 ...	15.00s	382	
LLM	ChatOpenAI	human: How di...	8/5/2023, 9:31:02 PM	13.87s	471	
LLM	ChatOpenAI	human: Tell me...	8/2/2023, 5:29:12 PM	13.43s	397	
LLM	ChatOpenAI	human: Tell me...	8/2/2023, 5:28:37 PM	13.77s	400	
LLM	ChatOpenAI	human: Tell me...	8/2/2023, 5:24:47 PM	12.87s	392	

Filters

Full-Text Search

Search...

Name

☒ ChatOpenAI

Run Type

☐ Llm

Status

☐ Success

Other

☒ Latency >= 10s

☐ Tokens >= 1,000

☐ Today

☐ Canceled

LangChain Inc. > Projects > chat-langchain

chat-langchain

See: <https://chat.langchain.com/>

TracesLLM CallsMonitorSetup

LatencyP50P95

Trace Latency (s)

LLM Latency (s)

LLM Calls per Trace

Tokens / sec