

## Memory Masterclass: Make Your AI Agents Remember What They Do! – Mark Bain, AIUS

 AI Engineer  
141K subscribers

[Subscribe](#)

 22   
 Share  Clip  Save 

547 views Jun 27, 2025 [AIEWF 2025 Complete Playlist](#)

Are you ready to give your AI agents a memory upgrade?

Join us for a fast-paced workshop exploring how memory can transform your agents.

### What You'll Do:

Learn Leading Memory Solutions: Gain practical experience with open-source tools like Neo4j, Cognee, Graphiti, and MemO.

Explore Memory Types: Understand the theory behind long-term, short-term, episodic, semantic, and other memory types.

Discover Memory Benefits: Learn how memory improves recall, contextual awareness, and reasoning in autonomous agents.

Compare Implementations: Get a snapshot of how different solutions implement memory—what's built-in, flexible, and experimental. We'll also demonstrate GraphRAG memory solutions and a GraphRAG chat implemented with Google ADK.

Whether you're working on AI copilots, agentic workflows, or research prototypes, this workshop will help you embed real memory into your AI stack.

### About Mark Bain

I'm a deep tech founder building a private research lab: AIUS Technologies. We're on a mission to develop artificial life through R&D of long-term memory.

Previously I led a cybersecurity lab and an edtech startup. Worked with defense and national security clients, banks and Fortune 500 clients.

## Make Your Agents Remember What They Do!

Host: Mark Bain

Guest Speakers:

Vasilije Markovic

Alex Gilmore

Daniel Chalef

### Agenda

## Timeline

1 

AI Memory

Power Talk

2 

Live Demos

by Four Guest Speakers

3 

GraphRAG Chat arena

Working with Multi-Memory  
Chat Agents

4 

Q&A

with Speakers

Join the slack channel:  
**#workshop-graphrag-chat**

Follow along the walkthrough





Passionate about deep physics  
of memory architectures  
and its applications

Mark Bain

AI Researcher & Founder  
@ AIUS Technologies

## AI memory

Any data, in any format – including code, algorithms, hardware, and any causal changes that affect them

Attention (Gravity) + Diffusion (Entropy)  
+ VAEs (Ricci Flow/Smoothing) = Asymmetries

Memory x Compute =  $i^2$

Our universe is a network database.  
Graph is its hidden causal structure  
and everything else is a fuzzy diffusion

Relationships are the universes causal links

LLMs -> correlations, weights & biases  
Graphs -> preserve causality

# Reasoning requires causal links hence GraphRAG is the next Big Thing



Vasilije Markovic

Founder @ Cognee



Alexander Gilmore

AI Solutions Architect,  
Neo4j



Daniel Chalet

Founder @ Zep AI  
& Graphiti



Mark Bain

AI Researcher & Founder  
@ AIUS Technologies



Vasilije Markovic

Founder @ Cognee

Let your Agents load,  
search and optimize memory  
based on the task at hand

The screenshot shows a web browser window for 'Cognee' at localhost:3000. The page features a dark header with the Cognee logo and navigation links. Below the header, there's a sidebar with a 'Mexican Standoff' section for comparing GitHub profiles, an 'Activity Log' showing 'No activity logged.', and a main area for graph visualization.

**Main Area:** The main area displays a graph with three nodes: 'Add', 'Cognify', and 'Search'. Arrows indicate a flow from 'Add' to 'Cognify' (labeled 'but don't forget to...'), and from 'Cognify' to 'Search' (labeled 'and after that you can...'). To the right of the graph is a 'Node Details' panel with settings for 'Graph Shape: Left-right' and a 'Fit Graph into View' button. A message says 'No node selected.'

**Powered By:** The page is powered by neo4j, LanceDB, Deepnote, and crewai.

**Sidebar:**

- Mexican Standoff:** Allows users to compare GitHub profiles by entering two GitHub usernames (e.g., hajdulBB and lxobr) and clicking 'Start Mexican Standoff'.
- Activity Log:** Displays the message 'No activity logged.'

**Bottom:** A red button with '1 Issue X' and a footer with icons for GitHub, LinkedIn, and other social media.

Create Next App × Cognee ×

localhost:3000

cognee

Powered by: neo4j LanceDB Deepnote crewai

**Cognee Dev Mexican Standoff**

Agents compare GitHub profiles, and make a decision who is a better developer

GitHub username GitHub username

hajdul88 lxobr

Start Mexican Standoff ↗

**Activity Log**

04/06/2025, 13:18:36: Dispatching hiring crew agents

Add → but don't forget to → Cognify → and after that you can → Search

Node Details Feedback

Graph Shape: Left-right Fit Graph into View

No node selected.

Issue X

cognee cognee\_adjustment\_crewai\_demo

Project README.md cognee/api/.env cognee-freontend/.env

Available Languages: Portuguese · [中文] · Python

**cognee builds AI Memory**  
Improve your AI infrastructure

**cognee** Ask cognee 89.4% answer relevancy

**Vector Store** Ask Vector Solution Potato 5% answer relevancy

**ChatGPT** Ask ChatGPT

Run client

```
/Users/vasilije/Projects/cognee_demo_example/cognee/.venv/bin/python /Users/vasilije/Projects/cognee_demo_example/cognee/cognee/api/client.py
INFO: Started server process [6445]INFO: Waiting for application startup.INFO: Application startup complete.INFO: Uvicorn running on http://0.0.0.0:58008
INFO: 127.0.0.1:58008 - "GET /api/v1/datasets HTTP/1.1" 307 Temporary RedirectINFO: 127.0.0.1:58008 - "GET /api/v1/datasets/ HTTP/1.1" 401 UnauthorizedINFO: 127.0.0.1:58008 - "GET /api/v1/datasets/ HTTP/1.1" 200 OKINFO: 127.0.0.1:60915 - "OPTIONS /api/v1/crewai/run HTTP/1.1" 200 OKINFO: ("127.0.0.1", 2025-06-04T13:18:36.566249) [error] PermissionDeniedError: Request owner does not have necessary permission: [write] for all datasets requested. (Status code: 401)
```

1:1 LF UTF-8 2 spaces\* Python 3.11 (cognee) (2)

Create Next App × Cognee × +

cognee Powered by neo4j LanceDB Deepnote crewai

Cognee Dev Mexican Standoff

Agents compare GitHub profiles, and make a decision who is a better developer

GitHub username GitHub username

hajdul88 lxbcr

Start Mexican Standoff ↗

Activity Log

04/06/2025, 13:18:36: GitHub applicant data ingestion started

04/06/2025, 13:18:36: Dispatching hiring crew agents

Add → Cognify → Search

No node selected.

Graph Shape: Left-right Fit Graph into View

N 1 Issue X

Create Next App × Cognee × +

cognee Powered by neo4j LanceDB Deepnote crewai

Cognee Dev Mexican Standoff

Agents compare GitHub profiles, and make a decision who is a better developer

GitHub username GitHub username

hajdul88 lxbcr

Start Mexican Standoff ↗

Activity Log

04/06/2025, 13:19:22: Applicant's (hajdul88) data ingestion finished

04/06/2025, 13:19:50: Graph updated: | GitHubUser: 1 | Repository: 1 | Issue: 7 | Comment: 7 | NodeSet: 2 | Commit: 3 | FileChange: 4

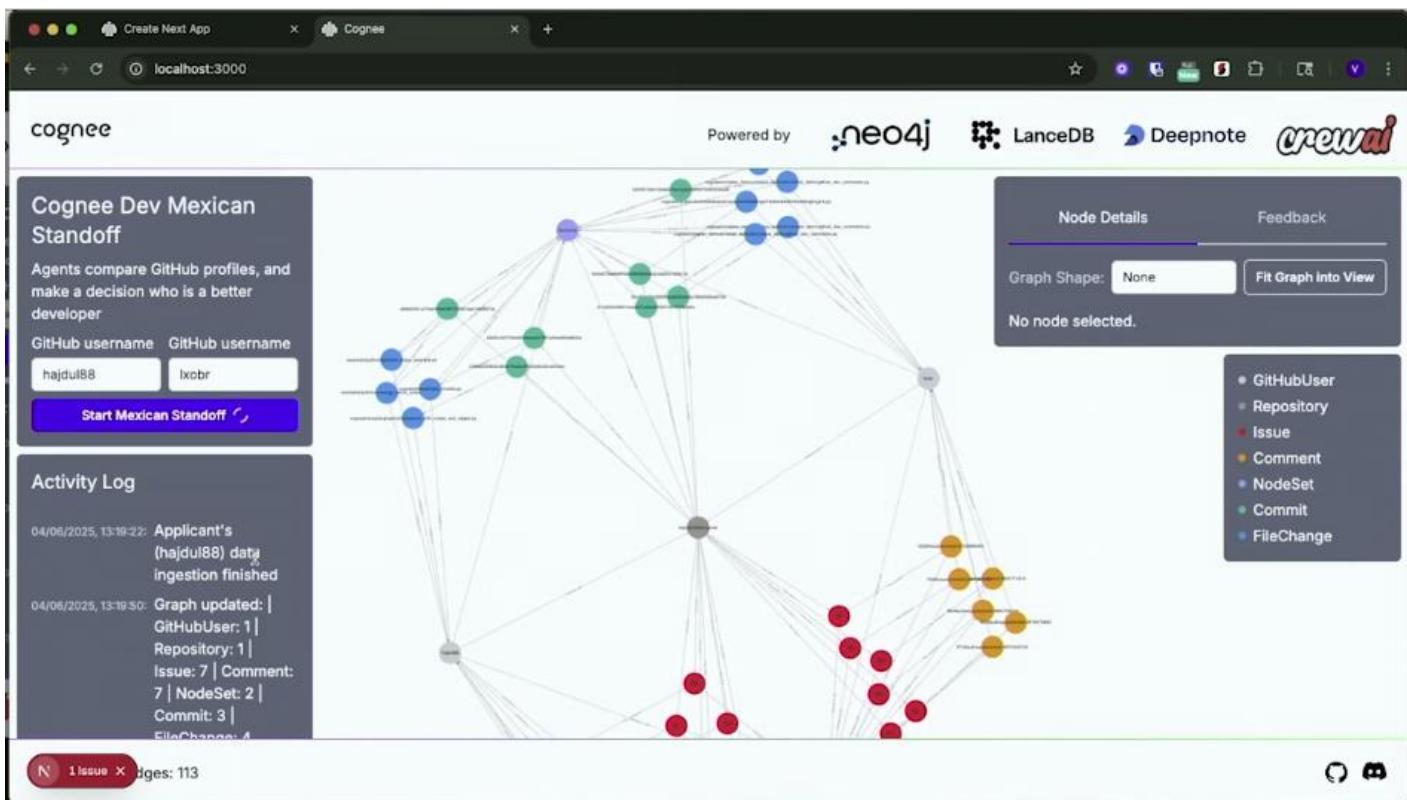
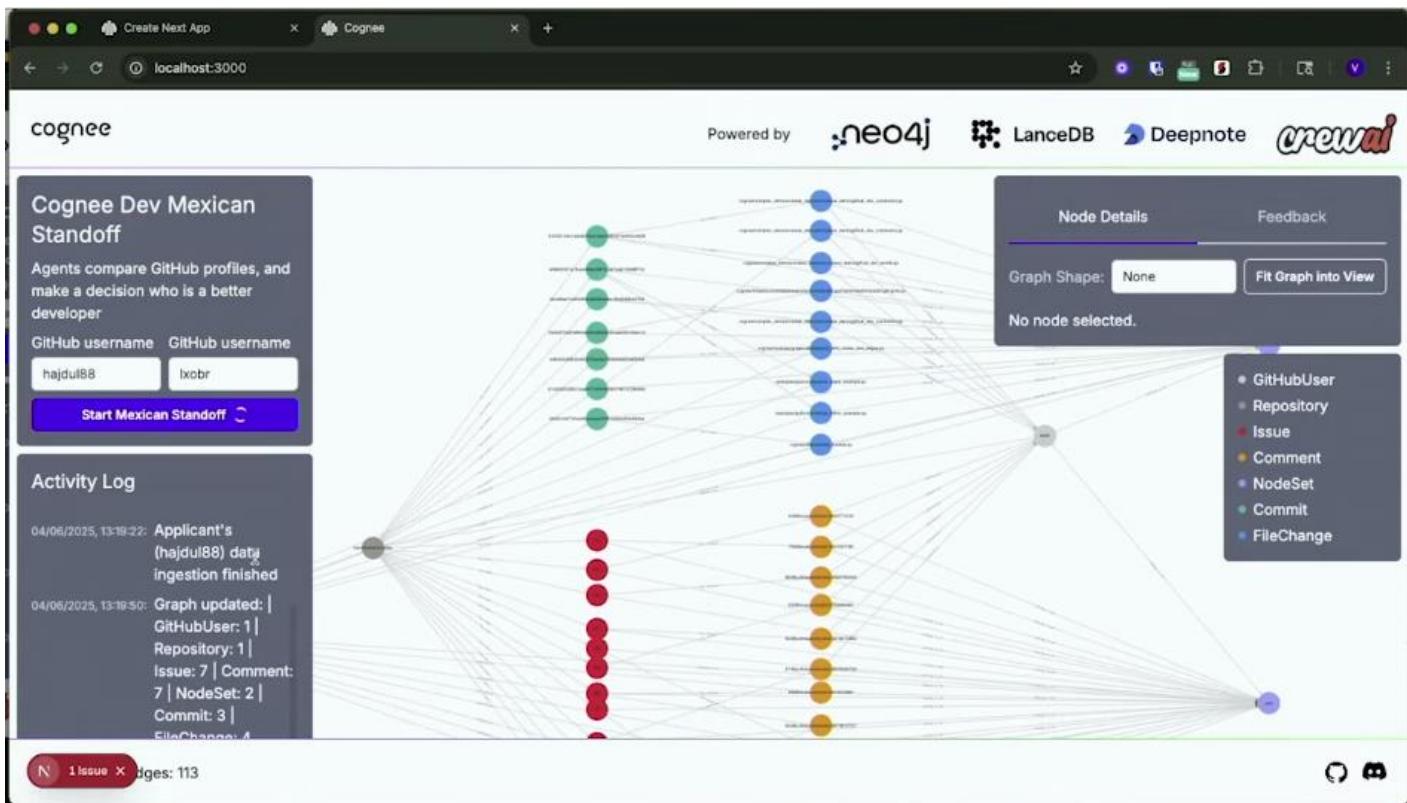
N 1 Issue X Edges: 57

Node Details Feedback

Graph Shape: Left-right Fit Graph into View

No node selected.

- GitHubUser
- Repository
- Issue
- Comment
- NodeSet
- Commit
- FileChange



```
Project Run client .env README.md cognee/api/.env cognee-frontend/.env
Run: client
Name: crew
ID: 76583f2a-9844-4190-89e7-87df5c9d56b7

Crew: crew
Task: 717a7621-528b-4c0e-8a94-eac0e184469a
Status: Executing Task...

Crew: crew
Task: 717a7621-528b-4c0e-8a94-eac0e184469a
Status: Executing Task...
Agent: Focused on communication, collaboration, and documentation excellence.

Status: In Progress

# Agent: Focused on communication, collaboration, and documentation excellence.
## Task: Search cognee for comments authored by 'hajdul88'. Use the "search_from_cognee" tool to collect information. Evaluate their communication clarity, communication clarity, community engagement, and kindness.
--- Example Output ---
Input:
applicant_1: Sarah Jennings

Output:
- Name: Sarah Jennings
- communication_clarity: 0.92
- community_engagement: 0.88
- Kindness: 0.95
```

```
Project Run client .env README.md cognee/api/.env cognee-frontend/.env
Run: client
applicant_1: Sarah Jennings

Output:
- Name: Sarah Jennings
- communication_clarity: 0.92
- community_engagement: 0.88
- Kindness: 0.95
- reasoning:
    Sarah consistently communicates with clarity and structure. In several threads, her responses broke down complex issues into actionable steps, showing strong explanatory skills. She uses inclusive language like "let's", "we should", and frequently thanks others for their input, which indicates a high degree of Kindness. Sarah also initiates or joins collaborative threads, offering feedback or connecting people with relevant documentation. Her tone is encouraging and non-defensive, even when correcting misunderstandings. These patterns were observed across over 8 threads involving different team members over a 3-week span.

Agent: Focused on communication, collaboration, and documentation excellence.

Status: In Progress
Thinking...

Agent: Focused on communication, collaboration, and documentation excellence.

Status: In Progress
Agent: Focused on communication, collaboration, and documentation excellence.

Status: In Progress
comments authored by 'hajdul88'
```

```
cognee cognify_adjustment_crewai_demo .env README.md cognee/api/.env cognee-frontend/.env
Project Run client
Agent: Focused on communication, collaboration, and documentation excellence.
Status: In Progress
# Agent: Focused on communication, collaboration, and documentation excellence.
## Task: Search cognee for comments authored by 'lxobr'. Use the "search_from_cognee" tool to collect information. Evaluate their communication clarity, community engagement, and kindness.
--- Example Output ---
Input:
applicant_1: Sarah Jennings

Output:
- Name: Sarah Jennings
- communication_clarity: 0.92
- community_engagement: 0.88
- kindness: 0.95
- reasoning: >
    Sarah consistently communicates with clarity and structure. In several threads, her responses broke down complex issues into actionable steps, showing strong explanatory skills. She uses inclusive language like "let's", "we should", and frequently thanks others for their input, which indicates a high degree of kindness. Sarah also initiates or joins collaborative threads, offering feedback or connecting people with relevant documentation. Her tone is encouraging and non-defensive, even when correcting misunderstandings. These patterns were observed across over 8 threads involving different team members over a 3-week span.

Agent: Focused on communication, collaboration, and documentation excellence.
Status: In Progress
Comments authored by lxobr
```

```
cognee cognify_adjustment_crewai_demo .env README.md cognee/api/.env cognee-frontend/.env
Project Run client
- community_engagement: 0.88
- kindness: 0.95
- reasoning: >
    Sarah consistently communicates with clarity and structure. In several threads, her responses broke down complex issues into actionable steps, showing strong explanatory skills. She uses inclusive language like "let's", "we should", and frequently thanks others for their input, which indicates a high degree of kindness. Sarah also initiates or joins collaborative threads, offering feedback or connecting people with relevant documentation. Her tone is encouraging and non-defensive, even when correcting misunderstandings. These patterns were observed across over 8 threads involving different team members over a 3-week span.

Agent: Focused on communication, collaboration, and documentation excellence.
Status: In Progress
Comments authored by lxobr

# Agent: Focused on communication, collaboration, and documentation excellence.
## Thought: I need to gather information about comments authored by 'lxobr' to assess their communication clarity, community engagement, and kindness within the organization.
## Using tool: search_from_cognee
## Tool Input:
"{'query': '\"Comments authored by lxobr\"'}"
## Tool Output:
[{'github': ['Comments authored by lxobr:\n\n1. I also fixed a little bug with the TaskGetters.\n2. Looks good, nice work!\n3. Nice work!\n4. Great work!\n5. Note\nthe spelling of \"TaskGetters\".\n\nGreat work!']}

Agent: Focused on communication, collaboration, and documentation excellence.

Status: In Progress
Thinking...
```

Create Next App Cognee localhost:3000

**cognee**

Powered by **neo4j** **LanceDB** **Deepnote** **crewai**

**Cognee Dev Mexican Standoff**  
Agents compare GitHub profiles, and make a decision who is a better developer

GitHub username GitHub username  
hajdul88 Ixobr

**Start Mexican Standoff**

**Activity Log**

- (search\_from\_cognee execution started)
- 04/06/2025, 13:19:56: Applicant's (Ixobr) data ingestion finished
- 04/06/2025, 13:19:22: Applicant's (hajdul88) data ingestion finished
- 04/06/2025, 13:19:50: Graph updated: |

N 1 Issue X Edges: 113

**Node Details** **Feedback**

Graph Shape: Bottom-up **Fit Graph Into View**

No node selected.

Legend:

- GitHubUser
- Repository
- Issue
- Comment
- NodeSet
- Commit
- FileChange

Create Next App Cognee localhost:3000

**cognee**

Powered by **neo4j** **LanceDB** **Deepnote** **crewai**

**Cognee Dev Mexican Standoff**  
Agents compare GitHub profiles, and make a decision who is a better developer

GitHub username GitHub username  
hajdul88 Ixobr

**Start Mexican Standoff**

**Activity Log**

- (search\_from\_cognee execution started)
- 04/06/2025, 13:19:56: Applicant's (Ixobr) data ingestion finished
- 04/06/2025, 13:19:22: Applicant's (hajdul88) data ingestion finished
- 04/06/2025, 13:19:50: Graph updated: |

N 1 Issue X Edges: 113

**Node Details** **Feedback**

Graph Shape: Radial-in **Fit Graph Into View**

No node selected.

Legend:

- GitHubUser
- Repository
- Issue
- Comment
- NodeSet
- Commit
- FileChange

Create Next App × Cognee × +

cognee

Powered by neo4j LanceDB Deepnote crewai

### Cognee Dev Mexican Standoff

Agents compare GitHub profiles, and make a decision who is a better developer

GitHub username GitHub username

hajdul88 lxobr

**Start Mexican Standoff**

**Activity Log**

- 04/06/2025, 13:21:14: Agent tool call (search\_from\_cognee) execution started
- 04/06/2025, 13:21:01: Agent tool call (search\_from\_cognee) execution completed
- 04/06/2025, 13:20:45: Agent tool call (search\_from\_cognee) execution started

N 1 Issue X Edges: 113

**Node Details**

Graph Shape: Top-down Fit Graph into View

ID: f7267f21-2a6c-588b-a4bb-a2d73699e468  
Label: cognee/complex\_demos/crewai\_demo/src/cree  
Ontology\_valid:  
Version: 1  
Topological\_rank: 0  
Metadata: {"index\_fields":["text"]}  
Filename: cognee/complex\_demos/crewai\_demo/src/  
Status: modified  
Additions: 4  
Deletions: 4  
Changes: 8  
Text: @@ -16,22 +16,22 @@ def \_\_init\_\_(

• GitHubUser Repository

Create Next App × Cognee × +

cognee

Powered by neo4j LanceDB Deepnote crewai

### Cognee Dev Mexican Standoff

Agents compare GitHub profiles, and make a decision who is a better developer

GitHub username GitHub username

hajdul88 lxobr

**Start Mexican Standoff**

**Activity Log**

- 04/06/2025, 13:21:14: Agent tool call (search\_from\_cognee) execution started
- 04/06/2025, 13:21:01: Agent tool call (search\_from\_cognee) execution completed
- 04/06/2025, 13:20:45: Agent tool call (search\_from\_cognee) execution started

N 1 Issue X Edges: 113

**Node Details**

Graph Shape: Top-down Fit Graph into View

```
self.include_issue_details =
    include_issue_details def
    get_issue_comments(self): - """Fetches
        comments made by the user on Issues across
        repositories within timeframe."""
    + """Fetches
        comments made by the user on issues and
        PRs across repositories within timeframe."""
    if
        not self.profile.user: return None date_filter =
        self._get_date_filter(self.days) - query =
        f"commenter:{self.profile.username}"
    isissue(date_filter)" + query = f"commenter:
        {self.profile.username}(date_filter)" return
    self._get_comments_from_search(query) def
    get_repo_issue_comments(self, repo_name): -
        """Fetches comments made by the user on
        issues in a specific repository within
```

• GitHubUser Repository

Create Next App × Cognee × +

localhost:3000

cognee

Powered by neo4j LanceDB Deepnote crewai

### Cognee Dev Mexican Standoff

Agents compare GitHub profiles, and make a decision who is a better developer

GitHub username GitHub username

hajdul88 lxcbr

**Start Mexican Standoff**

**Activity Log**

- 04/06/2025, 13:21:14: Agent tool call (search\_from\_cognee) execution started
- 04/06/2025, 13:21:01: Agent tool call (search\_from\_cognee) execution completed
- 04/06/2025, 13:20:45: Agent tool call (search\_from\_cognee) execution started

A complex network graph visualization showing connections between GitHub users, issues, and comments. The graph is highly interconnected, with many nodes (GitHubUser, Issue, Comment) and edges representing relationships.

**Node Details**

Graph Shape: Top-down Fit Graph into View

```
issues in a specific repository within timeframe." + """Fetches comments made by the user on issues and PRs in a specific repository within timeframe.""" if not self.profile.user: return None date_filter = self._get_date_filter(self.days) - query = f"repo:{repo_name} is:issue commenter: {self.profile.username}{date_filter}" + query = f"repo:{repo_name} commenter: {self.profile.username}{date_filter}" self.profile.github.get_repo(repo_name) return self._get_comments_from_search(query)
```

Commit sha: 2bcd8aa74d2f03f0ba8d3a1f6d5f41  
Repo: topoteretes/cognee  
Modifies: cognee/complex\_demos/crewai\_demo/src/

**Feedback**

GitHubUser Repository Issue Comment NodeSet

N 1 Issue X Edges: 113

Create Next App × Cognee × +

localhost:3000

cognee

Powered by neo4j LanceDB Deepnote crewai

### Cognee Dev Mexican Standoff

Agents compare GitHub profiles, and make a decision who is a better developer

GitHub username GitHub username

hajdul88 lxcbr

**Start Mexican Standoff**

**Activity Log**

- 04/06/2025, 13:21:14: Agent tool call (search\_from\_cognee) execution started
- 04/06/2025, 13:21:01: Agent tool call (search\_from\_cognee) execution completed
- 04/06/2025, 13:20:45: Agent tool call (search\_from\_cognee) execution started

A complex network graph visualization showing connections between GitHub users, issues, and comments. The graph is highly interconnected, with many nodes (GitHubUser, Issue, Comment) and edges representing relationships.

**Node Details**

Graph Shape: Top-down Fit Graph into View

ID: 03097dce-b497-5ed1-87da-e26f0aed1bf0  
Label: hajdul88  
Ontology\_valid:  
Version: 1  
Topological\_rank: 0  
Metadata: {"index\_fields":["name"]}  
Public\_repos: 3  
Followers: 4  
Following: 1

**Feedback**

GitHubUser Repository Issue Comment NodeSet

N 1 Issue X Edges: 113

cognee cognify\_adjustment\_crewai\_demo

Project Run client

```
Status: Completed
Task: e85a3f85-ffed-41c4-ac35-2f5deec956ce
Assigned to: Specialized in evaluating technical skills and code quality.

Status: Completed
Agent: Specialized in evaluating technical skills and code quality.

Status: Completed
Task: 77003c63-16f4-4f81-b3f3-89ecb66ad59b
Assigned to: Specialized in evaluating technical skills and code quality.

Status: Completed
Agent: Specialized in evaluating technical skills and code quality.

Status: Completed
Task: a1200225-f681-4933-9486-1429fa35624d
Status: Executing Task...
Agent: CTO/CEO-level decision maker who integrates expert feedback.

Status: In Progress

# Agent: CTO/CEO-level decision maker who integrates expert feedback.
## Task: Review the technical and soft skill assessment task outputs for candidates: -hajdul88 and -lxobr, then decide HIRE or NO_HIRE for each candidate with a defensible rationale.
-hajdul88
-lxobr
We have to hire one of them.
Prepare the final output for the ingest_hiring_decision_task.
```

cognee > cognee

1:1 LF UTF-8 2 spaces\* Python 3.11 (cognee) (2)

cognee cognify\_adjustment\_crewai\_demo

Project Run client

```
Agent: Focused on communication, collaboration, and documentation excellence.

Status: In Progress

# Agent: Focused on communication, collaboration, and documentation excellence.
## Task: Search cognee for comments authored by 'hajdul88'. Use the "search_from_cognee" tool to collect information. Evaluate their communication clarity, community engagement, and reasoning skills.
--- Example Output ---
Input:
applicant_1: Sarah Jennings

Output:
- Name: Sarah Jennings
- communication_clarity: 0.92
- community_engagement: 0.88
- kindness: 0.95
- reasoning: >
    Sarah consistently communicates with clarity and structure. In several threads, her responses broke down complex issues into actionable steps, showing strong explanatory skills. She uses inclusive language like "let's", "we should", and frequently thanks others for their input, which indicates a high degree of kindness. Sarah also initiates or joins collaborative threads, offering feedback or connecting people with relevant documentation. Her tone is encouraging and non-defensive, even when correcting misunderstandings. These patterns were observed across over 8 threads involving different team members over a 3-week span.

Agent: Focused on communication, collaboration, and documentation excellence.

Status: In Progress

Comments authored by 'hajdul88'
```

cognee > cognee

1:1 LF UTF-8 2 spaces\* Python 3.11 (cognee) (2)

Create Next App × Cognee ×

localhost:3000 cognee

Powered by neo4j LanceDB Deepnote crewai

### Cognee Dev Mexican Standoff

Agents compare GitHub profiles, and make a decision who is a better developer

GitHub username GitHub username

hajdul88 lxobr

**Start Mexican Standoff**

**Activity Log**

- 04/06/2025, 13:21:53: Agent tool call (search\_from\_cognee) execution started
- 04/06/2025, 13:21:44: Research round 2 started
- 04/06/2025, 13:21:30: Agent tool call (search\_from\_cognee) execution completed
- 04/06/2025, 13:21:14: Agent tool call

N 1 Issue Edges: 113

**Node Details**

Graph Shape: Top-down Fit Graph into View

ID: 03097dce-b497-5ed1-87da-e26f0aed1bf0  
Label: hajdul88  
Ontology.valid:  
Version: 1  
Topological\_rank: 0  
Metadata: {"index\_fields":["name"]}  
Public\_repos: 3  
Followers: 4  
Following: 1

GitHubUser Repository Issue Comment NodeSet

presentation\_demo\_recording ×

localhost:3000 cognee

Powered by neo4j LanceDB Deepnote crewai

### Cognee Dev Mexican Standoff

Agents compare GitHub profiles, and make a decision who is a better developer

GitHub username GitHub username

hajdul88 lxobr

**Start Mexican Standoff**

**Activity Log**

- 04/06/2025, 10:48:35: Agent tool call (search\_from\_cognee) execution started
- 04/06/2025, 10:48:35: Graph updated: | Entity: 4 | NodeSet: 3 | Repository: 1 | Issue: 13 | Comment: 13 | Commit: 7 | FileChange: 9 | DocumentChunk: 1 | EntityType: 2 | TextDocument: 1 | TextSummary: 1
- 04/06/2025, 10:48:15: Agent tool call (ingest\_report\_to\_cognee) execution started

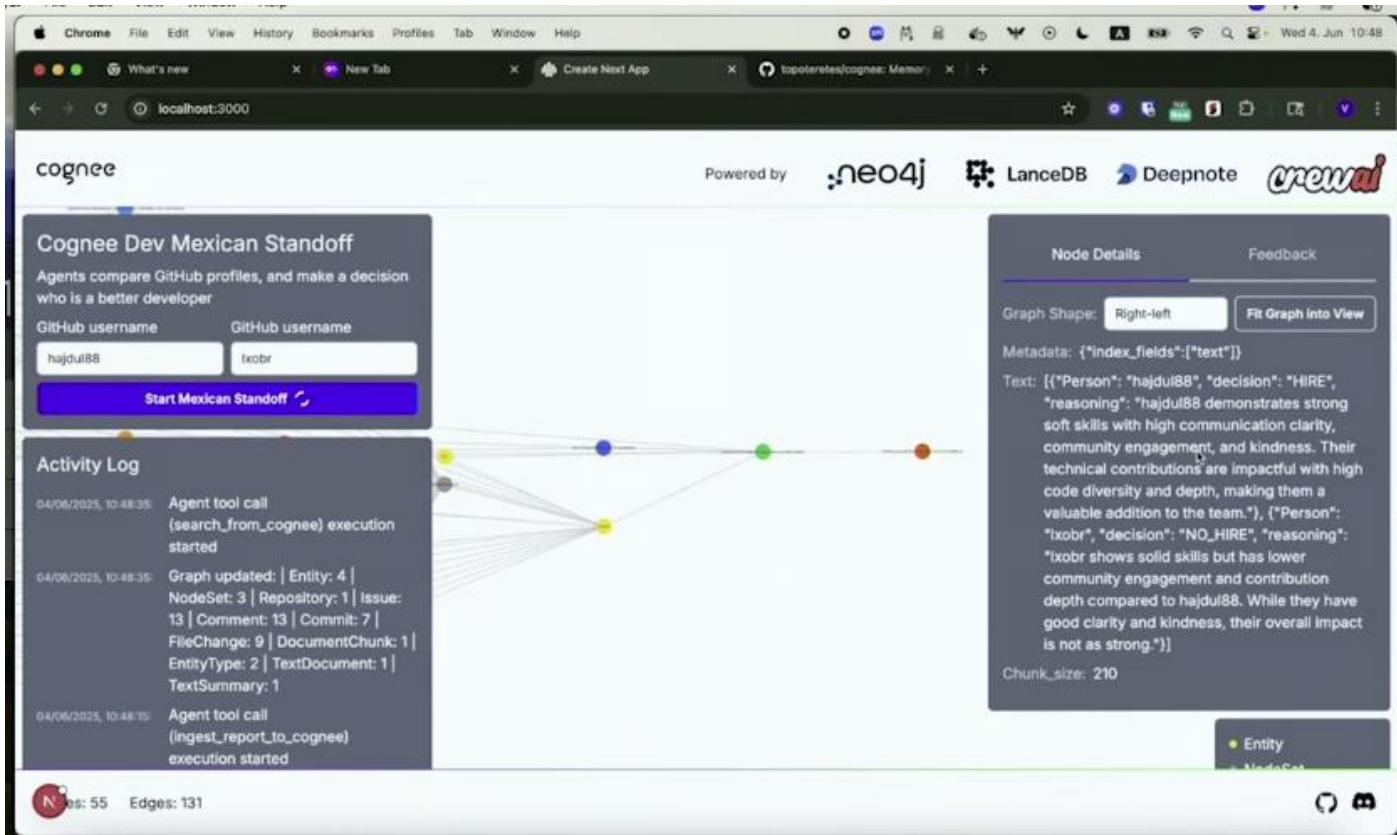
Notes: 55 Edges: 131

**Node Details**

Graph Shape: Right-left Fit Graph into View

Label: DocumentChunk\_9b9e35/8-14a9-5d84-ac40-29fb4cfab8ac  
Ontology.valid:  
Version: 1  
Topological\_rank: 0  
Metadata: {"index\_fields":["text"]}  
Text: [{"Person": "hajdul88", "decision": "HIRE", "reasoning": "hajdul88 demonstrates strong soft skills with high communication clarity, community engagement, and kindness. Their technical contributions are impactful with high code diversity and depth, making them a valuable addition to the team."}, {"Person": "lxobr", "decision": "NO\_HIRE", "reasoning": "lxobr shows solid skills but has lower communication clarity and less community engagement compared to hajdul88."}]

Entity NodeSet



+++

Preview README.md — ai-engineer-neo4j-memory-demo

claudie\_desktop\_config.json README.md Preview README.md

## AI Engineer Neo4j Memory MCP Demo

This is a short demo of the Neo4j Memory MCP server.

This file describes how to configure Claude Desktop with agentic memory courtesy of the Neo4j Memory MCP Server.

### Set Up

#### Install uv Python Package Manager

The Neo4j Memory MCP Server is written in Python and managed by uv. Follow the directions on this page to properly install.

#### Neo4j Database

Create a Neo4j database instance by either

- Navigating to the Neo4j Aura Console to create a cloud hosted Neo4j Aura instance
- Downloading Neo4j Desktop and creating a local Neo4j instance

#### Claude Desktop

Download Claude Desktop

- navigate to Settings -> Developer -> Edit Config
- Add the Neo4j Memory MCP server to the config file

You may replace the config values with your own database credentials

Preview README.md — ai-engineer-neo4j-memory-demo

(1) claude\_desktop\_config.json    (1) README.md    (1) Preview README.md X

- Navigating to the Neo4j Aura Console to create a cloud hosted Neo4j Aura instance
- Downloading Neo4j Desktop and creating a local Neo4j instance

## Claude Desktop

Download Claude Desktop

- navigate to Settings -> Developer -> Edit Config
- Add the Neo4j Memory MCP server to the config file

You may replace the config values with your own database credentials

```
{  
  "mcpServers": {  
    "neo4j-memory-store": {  
      "command": "uvx",  
      "args": [ "mcp-neo4j-memory@0.1.4" ],  
      "env": {  
        "NEO4J_URL": "bolt://localhost:7687",  
        "NEO4J_USERNAME": "neo4j",  
        "NEO4J_PASSWORD": "password",  
        "NEO4J_DATABASE": "neo4j"  
      }  
    }  
  }  
}
```

- Navigate to Settings -> General -> Claude Settings -> Configure
- In the personal preferences section - paste this text:

Follow these steps for each interaction:

Preview README.md — ai-engineer-neo4j-memory-demo

(1) claude\_desktop\_config.json    (1) README.md    (1) Preview README.md X

```
{  
  "mcpServers": {  
    "neo4j-memory-store": {  
      "command": "uvx",  
      "args": [ "mcp-neo4j-memory@0.1.4" ],  
      "env": {  
        "NEO4J_PASSWORD": "password",  
        "NEO4J_DATABASE": "neo4j"  
      }  
    }  
  }  
}
```

- Navigate to Settings -> General -> Claude Settings -> Configure
- In the personal preferences section - paste this text:

Follow these steps for each interaction:

1. Memory Retrieval:
  - Always begin your chat by saying only "Remembering..." and retrieve all relevant information from your knowledge graph
  - Always refer to your knowledge graph as your "memory"
2. Memory Update:
  - If any new information was gathered during the interaction, update your memory as follows:
    - a) Create entities for recurring organizations, people, and significant events
    - b) Connect them to the current entities using relations
    - c) Store facts about them as observations

This text is a modified version of the system prompt Anthropic recommends for their memory MCP server.

## Demo

Begin a conversation with Claude Desktop.

be automatically logged in the Neo4j database you've configured.

be used between different conversations with Claude and even accessed by other clients such as Cursor or Windsurf, if they are configured with the same Neo4j

**A let's start a company**

Remembering...

R read\_graph

That's exciting! Starting a company is a big adventure. What kind of business are you thinking about? Do you have a specific industry in mind, or maybe a problem you want to solve?

I'd love to hear about your initial ideas - whether it's a product, service, or maybe you're still in the brainstorming phase. What's sparking this entrepreneurial interest for you?

**A let's create a memory service for ai agents**

Remembering...

S search\_nodes

That's a fascinating idea! A memory service for AI agents could be incredibly valuable - there's definitely a growing need for persistent, structured memory systems as AI agents become more sophisticated and deployed in real-world applications.

What's your vision for this service? Are you thinking:

- A cloud-based memory infrastructure that different AI systems could plug into?
- Something focused on personal AI assistants that need to remember user preferences and history?

Reply to Claude...

+ ↻ Research BETA Claude Sonnet 4 ↑

**A let's start a company**

Remembering...

R read\_graph

Request  
{}

Response  
{  
  "entities": [],  
  "relations": []  
}

That's exciting! Starting a company is a big adventure. What kind of business are you thinking about? Do you have a specific industry in mind, or maybe a problem you want to solve?

I'd love to hear about your initial ideas - whether it's a product, service, or maybe you're still in the brainstorming phase. What's sparking this entrepreneurial interest for you?

**A let's create a memory service for ai agents**

Remembering...

S search\_nodes

Reply to Claude...

A let's create a memory service for ai agents

Remembering...

S search\_nodes

That's a fascinating idea! A memory service for AI agents could be incredibly valuable - there's definitely a growing need for persistent, structured memory systems as AI agents become more sophisticated and deployed in real-world applications.

What's your vision for this service? Are you thinking:

- A cloud-based memory infrastructure that different AI systems could plug into?
- Something focused on personal AI assistants that need to remember user preferences and history?
- Enterprise-focused for AI agents that need to maintain context across long-running business processes?
- Or maybe a more general-purpose memory layer that could work across different types of AI applications?

The market timing could be really good for this - as AI agents move beyond simple chat interactions to more complex, stateful applications, the memory problem becomes critical.

What got you thinking about this particular space?

C create\_entities

C create\_relations

Reply to Claude...

A let's create a memory service for ai agents

Remembering...

S search\_nodes



Request

```
{  
  'query': 'memory service ai agents company'  
}
```

Response

```
{  
  "entities": [],  
  "relations": []
```

That's a fascinating idea! A memory service for AI agents could be incredibly valuable - there's definitely a growing need for persistent, structured memory systems as AI agents become more sophisticated and deployed in real-world applications.

What's your vision for this service? Are you thinking:

- A cloud-based memory infrastructure that different AI systems could plug into?
- Something focused on personal AI assistants that need to remember user preferences and history?
- Enterprise-focused for AI agents that need to maintain context across long-running business processes?
- Or maybe a more general-purpose memory layer that could work across different types of AI applications?

Reply to Claude...



Research

BETA

Claude Sonnet 4



more sophisticated and deployed in real-world applications.

What's your vision for this service? Are you thinking:

- A cloud-based memory infrastructure that different AI systems could plug into?
- Something focused on personal AI assistants that need to remember user preferences and history?
- Enterprise-focused for AI agents that need to maintain context across long-running business processes?
- Or maybe a more general-purpose memory layer that could work across different types of AI applications?

The market timing could be really good for this - as AI agents move beyond simple chat interactions to more complex, stateful applications, the memory problem becomes critical. What got you thinking about this particular space?

```
C create_entities
  type: 'business_idea',
  'observations': [
    'Company concept focused on providing memory infrastructure for AI
    agents',
    'Addresses the need for persistent, structured memory systems for AI
    agents'
  ]
C create_relations
```

✖️ 🗑️ ⏪ Retry ↴

- Ⓐ it will need to easily integrate with orchestration frameworks such as langgraph and should have an mcp offering

Reply to Claude...

What's your vision for this service? Are you thinking:

- A cloud-based memory infrastructure that different AI systems could plug into?
- Something focused on personal AI assistants that need to remember user preferences and history?
- Enterprise-focused for AI agents that need to maintain context across long-running business processes?
- Or maybe a more general-purpose memory layer that could work across different types of AI applications?

The market timing could be really good for this - as AI agents move beyond simple chat interactions to more complex, stateful applications, the memory problem becomes critical. What got you thinking about this particular space?

```
C create_entities
  Request
  {
    'entities': [
      {
        'name': 'Memory Service for AI Agents',
        'type': 'Business Idea',
        'observations': [
          'Company concept focused on providing memory infrastructure for AI
          agents',
          'Addresses the need for persistent, structured memory systems for AI
          agents'
        ]
      }
    ]
  }
C create_relations
```

- Ⓐ it will need to easily integrate with orchestration frameworks such as langgraph and should have an mcp offering

Reply to Claude...

+ ⌂ 🔎 Research BETA

Claude Sonnet 4



What's your vision for this service? Are you thinking:

- A cloud-based memory infrastructure that different AI systems could plug into?
- Something focused on personal AI assistants that need to remember user preferences and history?
- Enterprise-focused for AI agents that need to maintain context across long-running business processes?
- Or maybe a more general-purpose memory layer that could work across different types of AI applications?

The market timing could be really good for this - as AI agents move beyond simple chat interactions to more complex, stateful applications, the memory problem becomes critical.  
What got you thinking about this particular space?

```
C create_entities
  {
    'type': 'Business Idea',
    'observations': [
      'Company concept focused on providing memory infrastructure for AI agents',
      'Addresses the need for persistent, structured memory systems for AI applications',
      'Potential market opportunity as AI agents become more sophisticated and stateful'
    ]
  },
C create_relations
```

Copy Like Comment Retry ▾

💡 It will need to easily integrate with existing AI frameworks such as TensorFlow and should

Reply to Claude...

What's your vision for this service? Are you thinking:

- A cloud-based memory infrastructure that different AI systems could plug into?
- Something focused on personal AI assistants that need to remember user preferences and history?
- Enterprise-focused for AI agents that need to maintain context across long-running business processes?
- Or maybe a more general-purpose memory layer that could work across different types of AI applications?

The market timing could be really good for this - as AI agents move beyond simple chat interactions to more complex, stateful applications, the memory problem becomes critical.  
What got you thinking about this particular space?

```
C create_entities
C create_relations
{
  'relations': [
    {
      'source': 'User',
      'target': 'Memory Service for AI Agents',
      'relationType': 'proposed'
    }
  ]
}
```

Copy Like Comment Retry ▾

💡 It will need to easily integrate with existing AI frameworks such as TensorFlow and should

Reply to Claude...

+ ≡ ⚙️ Research **BETA**

Claude Sonnet 4



```
        'target': 'Memory Service for AI Agents',
        'relationType': 'proposed'
    }
}
```

A it will need to easily integrate with orchestration frameworks such as langgraph and should have an mcp offering

Remembering...

S search\_nodes

Excellent strategic thinking! Focusing on integration with orchestration frameworks like LangGraph and offering MCP (Model Context Protocol) support shows you understand the ecosystem well.

This positioning makes a lot of sense:

**LangGraph Integration:** Since LangGraph is becoming a go-to for building complex, stateful AI workflows, having native integration would be huge. You could offer memory nodes that slot right into their graph architectures, making it seamless for developers to add persistent memory to their agent workflows.

**MCP Offering:** Smart move - MCP is gaining traction as the standard for AI tool integration. An MCP server for memory operations would let any MCP-compatible AI system (Claude, others) easily access your memory service. This could be a major distribution channel.

Are you thinking about the technical architecture? Like:

Reply to Claude...



Claude Sonnet 4



Are you thinking about the technical architecture? Like:

- Graph-based memory storage (fitting well with LangGraph's paradigm)
- Vector embeddings for semantic search
- Structured knowledge graphs for relationship tracking
- Maybe hybrid approaches?

And for the business model - usage-based pricing per memory operation, or more of a SaaS tier approach based on memory capacity and features?

The developer experience will be critical here. Having clean SDKs and plug-and-play integrations could really differentiate you in this space.

A add\_observations



C create\_entities

C create\_relations



A yea let's build our SDK in rust with bindings for popular languages

Remembering...

S search\_nodes

Rust for the core SDK is a brilliant choice! That's going to give you some serious advantages:

**Performance:** Memory operations need to be fast, especially when you're dealing with large knowledge graphs or high-frequency agent interactions. Rust's zero-cost abstractions will help here.

Reply to Claude...



Claude Sonnet 4



- All SDK features + MCP server
- Basic LangGraph integrations
- Email support
- Usage analytics dashboard

### Growth Tier - \$199/month

- 5M memory operations/month
- 1GB graph storage
- Advanced graph algorithms
- Priority support
- Custom MCP integrations
- Team collaboration features

### Scale Tier - \$799/month

- 50M memory operations/month
- 10GB graph storage
- White-label options
- Dedicated support channel
- SLA guarantees
- Advanced security features

### Enterprise - Custom

- Unlimited operations
- On-premise deployment options
- Custom Neo4j cluster sizing

Reply to Claude...

A add\_observations

 let's review what we have so far

Remembering...

```
R read_graph
{
  "entities": [
    {
      "name": "User",
      "type": "Person",
      "observations": [
        "Interested in starting a company",
        "Proposed creating a memory service for AI agents",
        "Entrepreneurial mindset"
      ]
    }
  ]
}
```

Here's what we've built so far for our AI Memory Service company:

### Core Business Concept

A memory service for AI agents that provides persistent, structured memory systems to enable more sophisticated and stateful AI applications.

### Technical Architecture

- **Database Foundation:** Built on Neo4j for optimal graph-based relationship modeling
- **Core SDK:** Written in Rust for performance and memory safety

Reply to Claude...

+ ↗ 🔍 Research BETA

Claude Sonnet 4



A add\_observations

A let's review what we have so far

Remembering...

R read\_graph

```
"observations": [
    "Interested in starting a company",
    "Proposed creating a memory service for AI agents",
    "Entrepreneurial mindset"
],
{
    "name": "Rust SDK",
    "type": "Software Component",
    "observations": [
        "Observation 1",
        "Observation 2"
    ]
}
```

Here's what we've built so far for our AI Memory Service company:

### Core Business Concept

A memory service for AI agents that provides persistent, structured memory systems to enable more sophisticated and stateful AI applications.

### Technical Architecture

- **Database Foundation:** Built on Neo4j for optimal graph-based relationship modeling
- **Core SDK:** Written in Rust for performance and memory safety

Reply to Claude...

- **Core SDK:** Written in Rust for performance and memory safety
- **Language Support:** Bindings for Python, JavaScript/Node.js, and other popular languages
- **Integration Strategy:** Native support for LangGraph orchestration and MCP (Model Context Protocol)

### Memory Types Supported

- **Episodic:** Event-based memories (conversations, meetings, interactions)
- **Semantic:** Factual knowledge and concepts (definitions, policies, facts)
- **Procedural:** How-to knowledge and processes (workflows, procedures)
- **Temporal:** Time-based patterns and scheduling (recurring events, deadlines)

### Key Innovation: Automated Classification

Smart SDK that automatically classifies memory types using content analysis, with confidence scoring and learning capabilities to reduce developer friction.

### Pricing Structure

- **Developer:** Free (10K operations, 1MB storage)
- **Startup:** \$49/mo (500K operations, 100MB storage)
- **Growth:** \$199/mo (5M operations, 1GB storage)
- **Scale:** \$799/mo (50M operations, 10GB storage)
- **Enterprise:** Custom (unlimited, on-premise options)

### Strategic Positioning

Targeting the growing market of sophisticated AI agents with focus on developer ecosystem integration and ease of use.

Pretty solid foundation! What aspect should we dive deeper into next?

Reply to Claude...

## Memory Types Supported

- **Episodic:** Event-based memories (conversations, meetings, interactions)
- **Semantic:** Factual knowledge and concepts (definitions, policies, facts)
- **Procedural:** How-to knowledge and processes (workflows, procedures)
- **Temporal:** Time-based patterns and scheduling (recurring events, deadlines)

## Key Innovation: Automated Classification

Smart SDK that automatically classifies memory types using content analysis, with confidence scoring and learning capabilities to reduce developer friction.

## Pricing Structure

- **Developer:** Free (10K operations, 1MB storage)
- **Startup:** \$49/mo (500K operations, 100MB storage)
- **Growth:** \$199/mo (5M operations, 1GB storage)
- **Scale:** \$799/mo (50M operations, 10GB storage)
- **Enterprise:** Custom (unlimited, on-premise options)

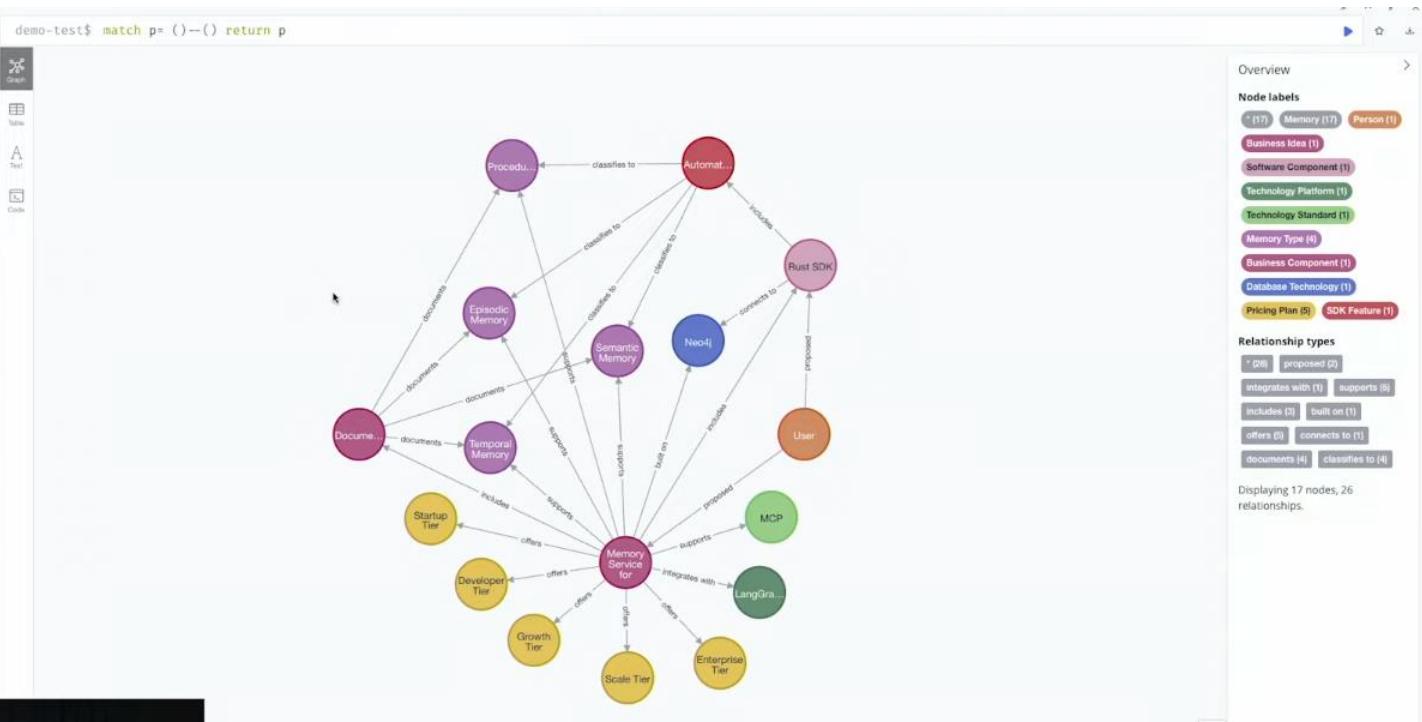
## Strategic Positioning

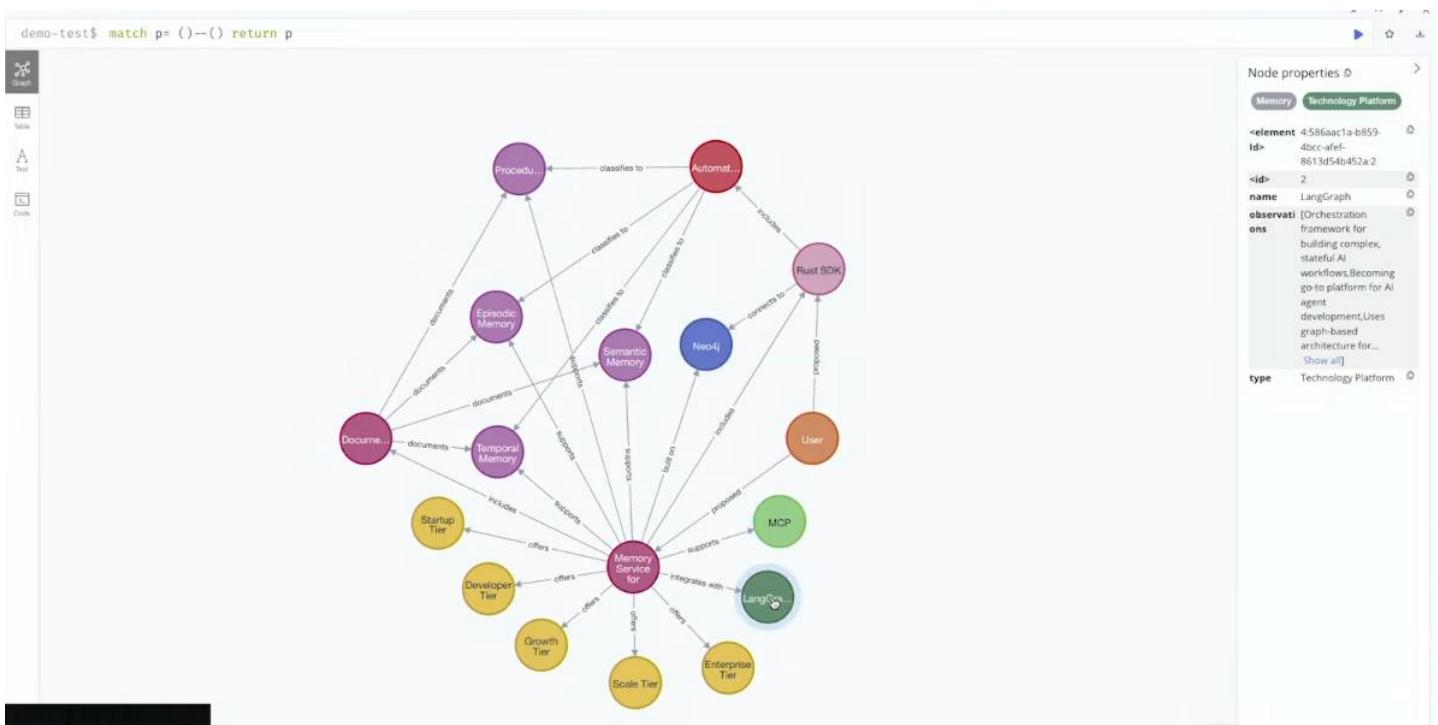
Targeting the growing market of sophisticated AI agents with focus on developer ecosystem integration and ease of use.

Pretty solid foundation! What aspect should we dive deeper into next?



Claude can make mistakes. Please double-check responses.







+++

# One Size Fits None

Why Memory Must Reflect Your Business Domain

RIGHT © ZEP SOFTWARE, INC. 2025

## PROBLEM

### When the Media Assistant Remembers Everything Except Your Listening Habits...

- Daniel likes jazz music.
- Daniel plays NPR podcast.
- Daniel wakes up at 7am.
- Daniel went to the gym.
- Daniel listened to The Daily.
- Daniel prefers mornings.
- Daniel's dog is named Melody.
- Daniel turned up the volume.
- Daniel asked about vegan recipes.
- Daniel paused Taylor Swift song.

Irrelevant facts that are semantically similar pollute memory; you need facts relevant to your business domain

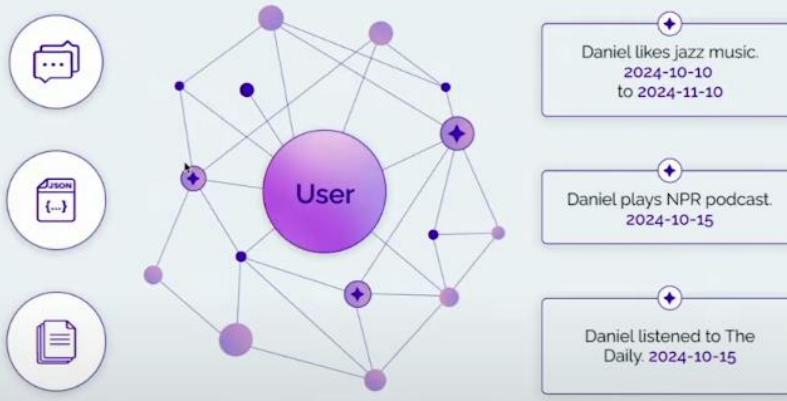
## PROBLEM

# ...Semantic Recall for Your Favorite Tunes is Problematic

## Semantic similarity ≠ Business Relevance

### THE KNOWLEDGE GRAPH

# Zep Creates a Unified, Continuously Evolving User Graph from Every Interaction



I want to buy a house. Help me plan for it.

Buying a house is an exciting goal! To help you plan effectively, could you share a bit more about your current situation? For example:

- Do you have a target price or down payment amount in mind?
- What's your current monthly income and expenses like?
- Do you have any debts or savings right now?

This will help me tailor a plan that fits your unique situation.

I earn \$150 a year

Thanks for sharing your income—that's a solid foundation for your house goal! To get a clearer picture, could you tell me about your current monthly expenses and any debts you might have? Knowing where your money goes each month and if you're carrying any debt will help us figure out how much you can realistically save toward your down payment.

i have student debt of \$75K

Thanks for sharing that. A \$75K student debt consider as you plan for your house purchase. To help you best, could you tell me the interest rate on your student loan ~~~~ your current monthly payment? This will help us understand how much of your income is going

Write a message...

Memory Context

- Financial Snapshot
- Financial Goals (1)
- buy a house
- Debts (1)
  - student\_debt
- Income Sources (1)
  - salary

FACTS and ENTITIES represent relevant context to the current conversation. # These are the most relevant facts and their valid date ranges # format: FACT (Date range: from - to) <FACTS> - jen\_james5 has student debt of \$75K (2025-06-03 03:29:09 - present) - jen\_james5 (Jen James who is a user) wants to buy a house. (2025-06-03 03:29:09 - present)

Updated: 6/2/2025, 9:08:00 PM

This app is a finance coach that I told I want to buy a house. Let us see what is stored in Zep's memory relating to relevant business objects

Finance Coach > jen\_james5 - Session 2315

Buying a house is an exciting goal! To help you plan effectively, could you share a bit more about your current situation? For example:

- Do you have a target price or down payment amount in mind?
- What's your current monthly income and expenses like?
- Do you have any debts or savings right now?

This will help me tailor a plan that fits your unique situation.

I earn \$150K a year

Thanks for sharing your income—that's a solid foundation for your house goal! To get a clearer picture, could you tell me about your current monthly expenses and any debts you might have? Knowing where your money goes each month and if you're carrying any debt will help us figure out how much you can realistically save toward your down payment.

i have student debt of \$75K

Thanks for sharing that. A \$75K student debt is a significant factor to consider as you plan for your house purchase. To help you best, could you tell me the interest rate on your student loan and your current monthly payment? This will help us understand how much of your income is going toward debt and how aggressively you might be saving for a house.

02:00 04:22

Memory Context

Live data from Zep memory and financial services 0s ago

Financial Snapshot

Financial Goals (1)

buy a house medium

Debts (1)

student\_joan %  
Balance: \$75,000

Income Sources (1)

salary guaranteed  
Monthly: \$12,500

Memory Context

FACTS and ENTITIES represent relevant context to the current conversation. # These are the most relevant facts and their valid date ranges # format: FACT (Date range: from - to) <FACTS> - jen\_james5 has student debt of \$75K (2025-06-03 03:29:09 - present) - jen\_james5 (jen James who is a user) wants to buy a house. (2025-06-03 03:29:09 - present)

Updated: 6/2/2025, 9:08:05 PM Live

ively, could you

mind?

?

1.

I earn \$150K a year

for your house

current

where your

ill help us figure

I earn \$150K a year

02:04 04:22

Memory Context

Live data from Zep memory and financial services 0s ago

Financial Snapshot

Financial Goals (1)

buy a house medium

Debts (1)

student\_joan %  
Balance: \$75,000

Income Sources (1)

salary guaranteed  
Monthly: \$12,500

Context

FACTS and ENTITIES represent relevant context to the current conversation. # These are the most relevant facts and their valid date ranges # format: FACT (Date range: from - to) <FACTS> - jen\_james5 has student debt of \$75K (2025-06-03 03:29:09 - present) - jen\_james5 (jen James who is a user) wants to buy a house. (2025-06-03 03:29:09 - present)



```
 3 | entityFields,
 4 | EntityType,
 5 | } from "@getzep/zep-cloud/wrapper/ontology";
 6 |
 7 | export const financialGoalSchema: EntityType = {
 8 |   description: "A specific financial objective the user wants to achieve.",
 9 |   fields: {
10 |     goal_type: entityFields.text(
11 |       "Type of financial goal (emergency_fund, house_down_payment, retirement, vacation, debt_"
12 |     ),
13 |     target_amount: entityFields.float("Target dollar amount for the goal"),
14 |     priority: entityFields.text("Goal priority level (high, medium, low)"),
15 |   },
16 | };
17 |
18 | export const expenseCategorySchema: EntityType = {
19 |   description: "A spending category showing where the user's money goes.",
20 |   fields: {
21 |     category_name: entityFields.text(
22 |       "Name of expense category (housing, food, transportation, entertainment, utilities, shopping, etc.)"
23 |     ),
24 |     monthly_spend: entityFields.float(
25 |       "Average monthly spending amount in this category"
26 |     ),
27 |     optimization_potential: entityFields.text(
28 |       "Potential for reducing spending in this category (high, medium, low, none)"
29 |     ),
30 |   },
31 | };
32 |
33 | export const debtAccountSchema: EntityType = {
34 |   description: "A debt obligation that impacts the user's financial health.",
35 |   fields: {
36 |     debt_type: entityFields.text(
37 |       "Type of debt account (mortgage, credit card, student loan, auto loan, etc.)"
38 |     ),
39 |     balance: entityFields.float("Current balance of the debt account"),
40 |     interest_rate: entityFields.float("Annual interest rate on the debt account"),
41 |     due_date: entityFields.date("Next due date for the debt payment"),
42 |     payment_plan: entityFields.text("Current payment plan for the debt")
43 |   }
44 | };

 0:00 02:28 04:22
```

```
■ zep-domain-enums-ai-engineer.mp4
  3 | entityFields,
  4 | EntityType,
  5 | } from "@getzep/zep-cloud/wrapper/ontology";
  6 |
  7 | export const financialGoalSchema: EntityType = {
  8 |   description: "A specific financial objective the user wants to achieve.",
  9 |   fields: {
10 |     goal_type: entityFields.text(
11 |       "Type of financial goal (emergency_fund, house_down_payment, retirement, vacation, debt_"
12 |     ),
13 |     target_amount: entityFields.float("Target dollar amount for the goal"),
14 |     priority: entityFields.text("Goal priority level (high, medium, low)"),
15 |   },
16 | };
17 |
18 | export const expenseCategorySchema: EntityType = {
19 |   description: "A spending category showing where the user's money goes.",
20 |   fields: {
21 |     category_name: entityFields.text(
22 |       "Name of expense category (housing, food, transportation, entertainment, utilities, shopping)"
23 |     ),
24 |     monthly_spend: entityFields.float(
25 |       "Average monthly spending amount in this category"
26 |     ),
27 |     optimization_potential: entityFields.text(
28 |       "Potential for reducing spending in this category (high, medium, low, none)"
29 |     ),
30 |   };
31 |
32 |
33 | export const debtAccountSchema: EntityType = {
34 |   description: "A debt obligation that impacts the user's financial health.",
35 |   fields: {
36 |     debt_type: entityFields.text(
```



```
■ zep-domain-enums-ai-engineer.mp4
  36 | debt_type: entityFields.text(
  37 |   "Type of debt (credit_card, student_loan, mortgage, personal_loan, auto_loan, medical_debt)"
  38 | ),
  39 | current_balance: entityFields.float(
  40 |   "Current outstanding balance in dollars"
  41 | ),
  42 | interest_rate: entityFields.float(
  43 |   "Annual interest rate as percentage (e.g., 18.5 for 18.5%)"
  44 | ),
  45 | );
  46 |
  47 |
  48 | export const assetSchema: EntityType = {
  49 |   description:
  50 |     "An asset or resource that contributes to the user's financial health.",
  51 |   fields: {
  52 |     asset_type: entityFields.text(
  53 |       "Type of asset (checking_account, savings_account, investment_account, retirement_401k, home)"
  54 |     ),
  55 |     current_value: entityFields.float("Current value or balance in dollars"),
  56 |     liquidity: entityFields.text(
  57 |       "How easily the asset can be accessed (high, medium, low)"
  58 |     ),
  59 |   },
  60 | };
  61 |
  62 | export const incomeSchema: EntityType = {
  63 |   description:
  64 |     "A source of income that contributes to the user's financial resources.",
  65 |   fields: {
  66 |     income_type: entityFields.text(
  67 |       "Type of income (salary, freelance, business, investment, rental, benefits, other)"
  68 |     ),
  69 |     monthly_amount: entityFields.float(
```



```
zep-domain-enriches-an-engineer.mp4
49  description:
50    "An asset or resource that contributes to the user's financial health.",
51  fields: {
52    asset_type: entityFields.text(
53      "Type of asset (checking_account, savings_account, investment_account, retirement_401k, h
54    ),
55    current_value: entityFields.float("Current value or balance in dollars"),
56    liquidity: entityFields.text(
57      "How easily the asset can be accessed (high, medium, low)"
58    ),
59  );
60};
61
62 export const incomeSourceSchema: EntityType = {
63  description:
64    "A source of income that contributes to the user's financial resources.",
65  fields: {
66    income_type: entityFields.text(
67      "Type of income (salary, freelance, business, investment, rental, benefits, pension)"
68    ),
69    monthly_amount: entityFields.float(
70      "Average monthly income from this source in dollars"
71    ),
72    reliability: entityFields.text(
73      "Income reliability (guaranteed, stable, variable, seasonal)"
74    ),
75  };
76};
77
78 // Type definitions for the entities
79 type FinancialGoal = EntityData<typeof financialGoalSchema>;
80 type ExpenseCategory = EntityData<typeof expenseCategorySchema>;
81 type DebtAccount = EntityData<typeof debtAccountSchema>;
82 type Asset = EntityData<typeof assetSchema>;
83
```



```
zep-domain-enriches-an-engineer.mp4
91  const [goals, expenses, debts, assets, income] = await Promise.all([
92    this.zepClient.graph.search({
93      userId,
94      query: "financial goals savings targets priorities",
95      scope: "nodes",
96      searchFilters: { nodeLabels: ["FinancialGoal"] },
97    }),
98    this.zepClient.graph.search({
99      userId,
100     query: "spending expenses monthly budget categories",
101     scope: "nodes",
102     searchFilters: { nodeLabels: ["ExpenseCategory"] },
103   }),
104   this.zepClient.graph.search({
105     userId,
106     query: "debt loans credit cards balance interest",
107     scope: "nodes",
108     searchFilters: { nodeLabels: ["DebtAccount"] },
109   }),
110   this.zepClient.graph.search({
111     userId,
112     query: "assets savings accounts investments emergency fund",
113     scope: "nodes",
114     searchFilters: { nodeLabels: ["Asset"] },
115   }),
116   this.zepClient.graph.search({
117     userId,
118     query: "monthly earnings",
119     scope: "nodes",
120     searchFilters: { nodeLabels: ["IncomeSource"] },
121   });
122 };
123
```



```
82 type Asset = EntityData<typeof assetSchema>;
83 type IncomeSource = EntityData<typeof incomeSourceSchema>;
84
85 // Streamlined FinancialCoachService with only the method used by our 3 essential tools
86 export class FinancialCoachService {
87   constructor(private zepClient: any) {}
88
89   async getFinancialSnapshot(userId: string) {
90     // Get all financial entities for complete picture
91     const [goals, expenses, debts, assets, income] = await Promise.all([
92       this.zepClient.graph.search({
93         userId,
94         query: "financial goals savings targets priorities",
95         scope: "nodes",
96         searchFilters: { nodeLabels: ["FinancialGoal"] },
97       }),
98       this.zepClient.graph.search({
99         userId,
100        query: "spending expenses monthly budget categories",
101        scope: "nodes",
102        searchFilters: { nodeLabels: ["ExpenseCategory"] },
103      }),
104      this.zepClient.graph.search({
105        userId,
106        query: "debt loans credit cards balance interest",
107        scope: "nodes"
108        searchFilters: { nodeLabels: ["Debt"] },
109      }),
110      this.zepClient.graph.search({
111        userId,
112        query: "assets savings accounts investments emergency fund",
113        scope: "nodes",
114        searchFilters: { nodeLabels: ["Asset"] },
115      })
116    ];
117  }
118}
```



```
3   entityFields,
4   EntityType,
5 } from "@getzep/zep-cloud/wrapper/ontology";
6
7 export const financialGoalSchema: EntityType = {
8   description: "A specific financial objective the user wants to achieve.",
9   fields: {
10     goal_type: entityFields.text(
11       "Type of financial goal (emergency_fund, house_down_payment, retirement, vacation, debt_payoff, car_purc
12     ),
13     target_amount: entityFields.float("Target dollar amount for the goal"),
14     priority: entityFields.text("Goal priority level (high, medium, low)"),
15   },
16 };
17
18 export const expenseCategorySchema: EntityType = {
19   description: "A spending category showing where the user's money goes.",
20   fields: {
21     category_name: entityFields.text(
22       "Name of expense category (housing, food, transportation, entertainment, utilities, shopping, healthcare
23     ),
24     monthly_spend: entityFields.float(
25       "Average monthly spending amount in this category"
26     ),
27     optimization_potential: entityFields.text(
28       "Potential for reducing spending in this category (high, medium, low, none)"
29     ),
30   },
31 };
32
33 export const debtAccountSchema: EntityType = {
34   description: "A debt obligation that impacts the user's financial health.",
35   fields: {
36 }
```



```
18  export const expenseCategorySchema: EntityType = {
19    fields: {
20      category_name: entityFields.text(
21        ),
22        monthly_spend: entityFields.float(
23          "Average monthly spending amount in this category"
24        ),
25        optimization_potential: entityFields.text(
26          "Potential for reducing spending in this category (high, medium, low, none)"
27        ),
28      },
29    },
30  };
31
32
33  export const debtAccountSchema: EntityType = {
34    description: "A debt obligation that impacts the user's financial health.",
35    fields: {
36      debt_type: entityFields.text(
37        "Type of debt (credit_card, student_loan, mortgage, personal_loan, auto_loan, medical_debt)"
38      ),
39      current_balance: entityFields.float(
40        "Current outstanding balance in dollars"
41      ),
42      interest_rate: entityFields.float(
43        "Annual interest rate as a percentage (e.g., 18.5 for 18.5%)"
44      ),
45    },
46  };
47
48  export const assetSchema: EntityType = {
49    description:
50      "An asset or resource that contributes to the user's financial health.",
51    fields: {
52      asset_type: entityFields.text(
53        "Type of asset (savings_account, investment, emergency_fund, real_estate, vehicle, etc.)"
54      ),
55    },
56  };
57
```



```
86  export class FinancialCoachService {
87    async getFinancialSnapshot(userId: string) {
88      this.zepClient.graph.search({
89        userId,
90        query: "spending expenses monthly budget categories",
91        scope: "nodes",
92        searchFilters: { nodeLabels: ["ExpenseCategory"] },
93      }),
94      this.zepClient.graph.search({
95        userId,
96        query: "debt loans credit cards balance interest",
97        scope: "nodes",
98        searchFilters: { nodeLabels: ["DebtAccount"] },
99      }),
100     this.zepClient.graph.search({
101       userId,
102       query: "assets savings accounts investments emergency fund",
103       scope: "nodes",
104       searchFilters: { nodeLabels: ["Asset"] },
105     }),
106     this.zepClient.graph.search({
107       userId,
108       query: "income salary freelance monthly earnings",
109       scope: "nodes",
110       searchFilters: { nodeLabels: ["IncomeSource"] },
111     }),
112   );
113   return { goals, expenses, debts, assets, income };
114 }
115
116
117
118
119
120
121
122
123
124
125
126
127
128 // Essential utility functions used by the comprehensive financial advice tool
129 export function debtVsSavingsAdvice()
```





```

5 import { SYSTEM_PROMPT } from "@lib/system-prompt";
6 import {
7   financialGoalSchema,
8   expenseCategorySchema,
9   debtAccountSchema,
10  assetSchema,
11  incomeSourceSchema,
12  FinancialCoachService,
13 } from "@/lib/models";
14
15 const zep = new ZepClient({ apiKey: process.env.ZEP_API_KEY });
16 await zep.graph.setEntityTypes([
17   FinancialGoal: financialGoalSchema,
18   ExpenseCategory: expenseCategorySchema,
19   DebtAccount: debtAccountSchema,
20   Asset: assetSchema,
21   IncomeSource: incomeSourceSchema,
22 ]);
23
24 const financialCoachService = new FinancialCoachService(zep);
25
26 export const runtime = "edge";
27 export const maxDuration = 30;
28
29 const botName = "FinanceCoach";
30 const botRole = "assistant";
31 const userRole = "user";
32
33 export async function POST(
  q: Request,
  params: { params: Promise<{ userId: string; sessionId: string }> }
) {
  const { messages, system, tools } = await req.json();
  ...
}

```

[play button]

03:12 ————— 04:22



Finance Coach > jen\_james5 - Session 2315

Buying a house is a big decision. To help you plan effectively, could you share a bit more about your current situation? For example:

- Do you have a target price or down payment amount in mind?
- What's your current monthly income and expenses like?
- Do you have any debts or savings right now?

This will help me tailor a plan that fits your unique situation.

I earn \$150K a year

Thanks for sharing your income—that's a solid foundation for your house goal! To get a clearer picture, could you tell me about your current monthly expenses and any debts you might have? Knowing where your money goes each month and if you're carrying any debt will help us figure out how much you can realistically save toward your down payment.

I have student debt of \$75K

Thanks for sharing that. A \$75K student debt is a significant factor to consider as you plan for your house purchase. To help you best, could you tell me the interest rate on your student loan and your current monthly payment? This will help us understand how much of your income is going toward debt and how aggressively you might be saving for a house.

[play button]

03:15 ————— 04:22

**Memory Context**

Live data from Zep memory and financial services 0s ago

- Financial Snapshot
- Financial Goals (1)
- buy a house
- Debts (1)
- student\_loan Balance: \$75,000
- Income Sources (1)
- salary Monthly: \$12,500 guaranteed

FACTS and ENTITIES represent relevant context to the current conversation. # These are the most relevant facts and their valid date ranges # format: FACT (Date range: from - to) <FACTS> ~ jen\_james5 has student debt of \$75K (2025-06-03 03:29:09 - present) - jen\_james5 (Jen James who is a user) wants to buy a house. (2025-06-03 03:29:09 - present)

Updated: 6/3/2025, 9:11:10 PM Live



Finance Coach > jen\_james5 - Session 2315

money goes each month and if you're carrying any debt will help us figure out how much you can realistically save toward your down payment.

i have student debt of \$75K

Thanks for sharing that. A \$75K student debt is a significant factor to consider as you plan for your house purchase. To help you best, could you tell me the interest rate on your student loan and your current monthly payment? This will help us understand how much of your income is going toward debt and how aggressively you might want to tackle it alongside saving for a house.

\$5000 a minth

Thanks for letting me know your student loan payment is \$5,000 a month—that's quite a substantial amount. Given your \$150K annual income (about \$12,500 per month before taxes), \$5,000 toward debt is a big chunk of your cash flow.

To help you plan for buying a house, I want to understand a bit more:

- What is your monthly take-home pay after taxes?
- What are your other monthly expenses like (rent, utilities, food, transportation, etc.)?

This will help us see how



Write a message...

Memory Context

Live data from Zep memory and financial services 0s ago

Financial Snapshot

Financial Goals (1)

buy a house medium

Debts (1)

student\_loan Balance: \$75,000 %

Income Sources (1)

salary Monthly: \$12,500 guaranteed

Memory Context

FACTS and ENTITIES represent relevant context to the current conversation. # These are the most relevant facts and their valid date ranges # format: FACT (Date range: from - to) <FACTS> - jen\_james5 has student debt of \$75K (2025-06-03 03:29:09 - present) - jen\_james5 (Jen James who is a user) wants to buy a house. (2025-06-03 03:29:09 - present)

Updated: 6/2/2025, 9:11:20 PM Live



current monthly

ur income is going

ickle it alongside

\$5000 a minth

is \$5,000 a month

annual income

debt is a big

nd a bit more:

ilities, food,

\$5000 a month

Monthly: \$5,000 medium

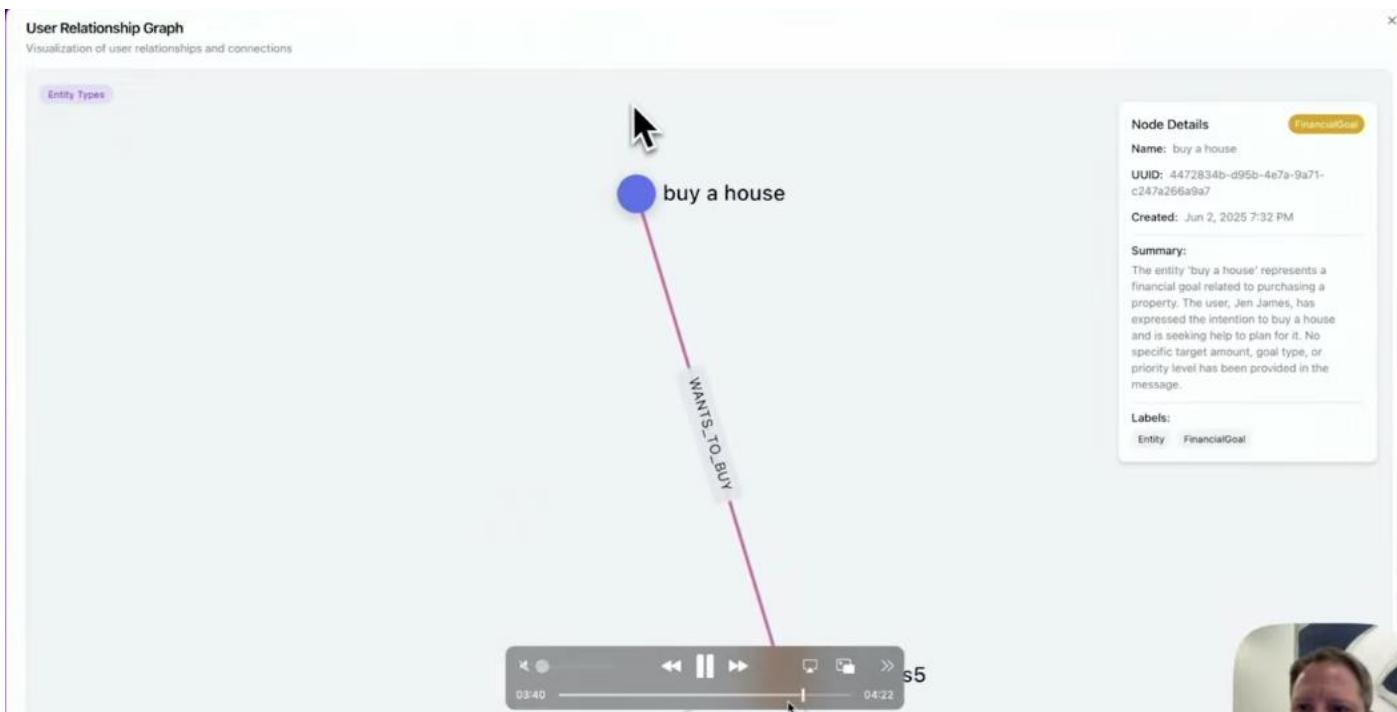
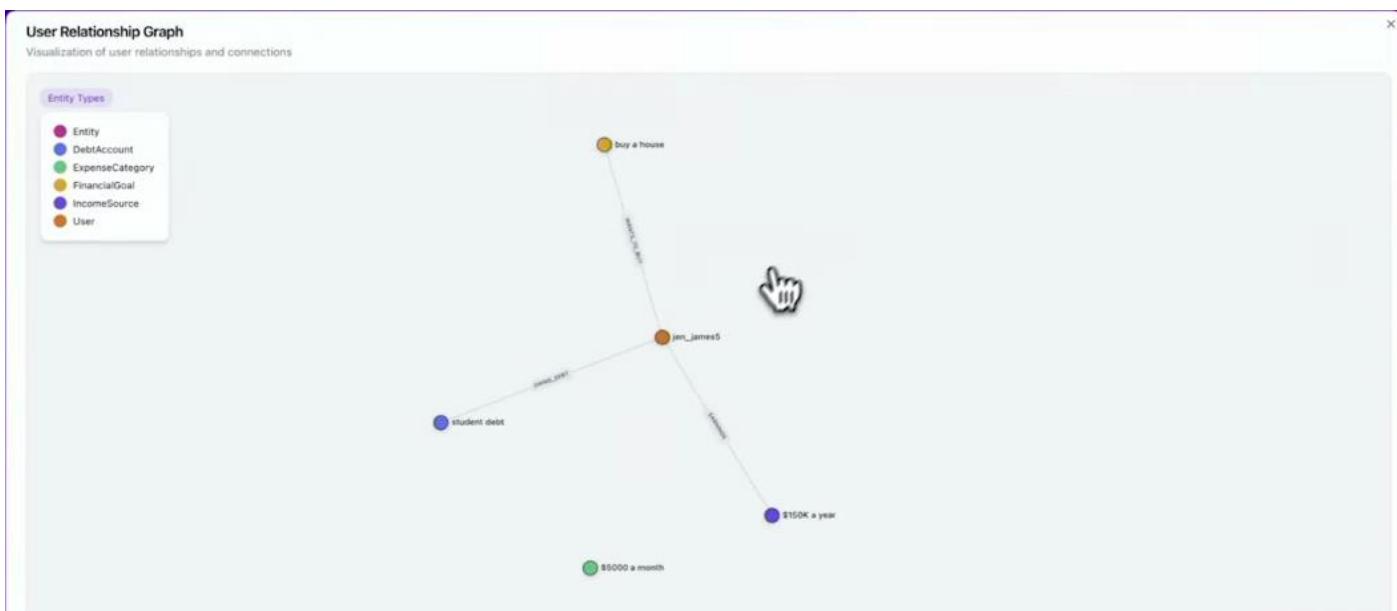
Memory C

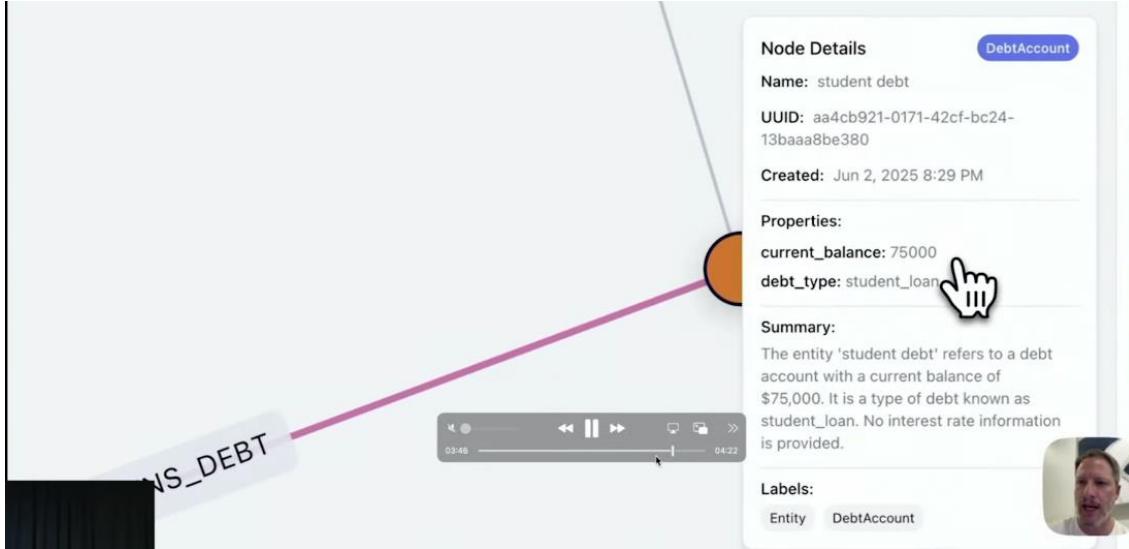
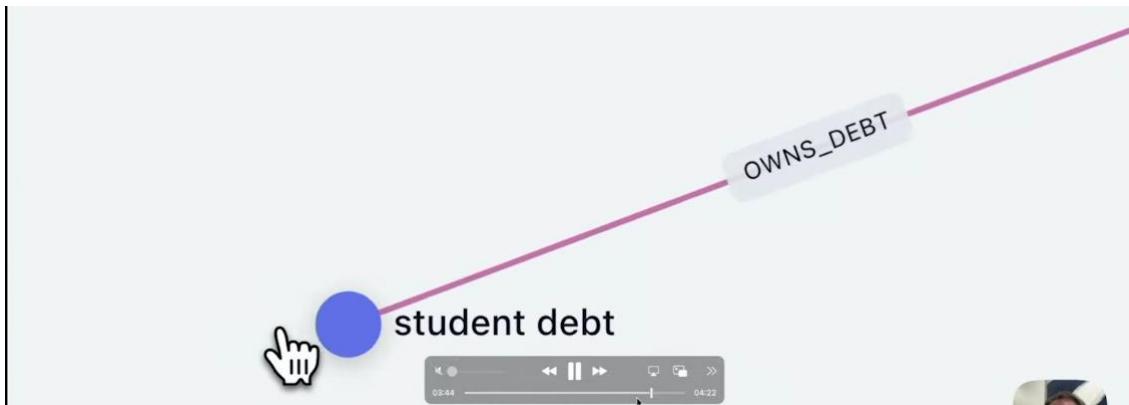


FACTS and ENTITIES represent relevant context to the current conversation. # These are the most relevant facts and their valid date ranges # format: FACT (Date range: from - to) <FACTS> - jen\_james5 has student debt of \$75K (2025-06-03 03:29:09 - present) - jen\_james5 (Jen James who is a user) wants to buy a house. (2025-06-03 03:29:09 - present)

Updated: 6/2/2025, 9:11:30 PM Live







# Zep crushes benchmarks

## LongMemEval Accuracy



Zep excels on benchmarks designed to replicate **complex real-world enterprise scenarios**.

† Full-context baseline places the entire dataset into the context window.

## ZEP: A TEMPORAL KNOWLEDGE GRAPH ARCHITECTURE FOR AGENT MEMORY

Preston Rasmussen  
Zep AI  
preston@getzep.com

Pavlo Palichuk  
Zep AI  
paul@getzep.com

Travis Beauvais  
Zep AI  
travis@getzep.com

Jack Ryan  
Zep AI  
jack@getzep.com

Daniel Chalef  
Zep AI  
daniel@getzep.com

### ABSTRACT

We introduce Zep, a novel memory layer service for AI agents that outperforms the current state-of-the-art system, MemGPT, in the Deep Memory Retrieval (DMR) benchmark. Additionally, Zep excels in more comprehensive and challenging evaluations than DMR that better reflect real-world enterprise use cases. While existing retrieval-augmented generation (RAG) frameworks for large language model (LLM)-based agents are limited to static document retrieval, enterprise applications demand dynamic knowledge integration from diverse sources including ongoing conversations and business data. Zep addresses this fundamental limitation through its core component Graphiti—a temporally-aware knowledge graph engine that dynamically synthesizes both unstructured conversational data and structured business data while maintaining historical relationships. In the DMR benchmark, which the MemGPT team established as their primary evaluation metric, Zep demonstrates superior performance (94.8% vs 93.4%). Beyond DMR, Zep's capabilities are further validated through the more challenging LongMemEval benchmark, which better reflects enterprise use cases through complex temporal reasoning tasks. In this evaluation, Zep achieves substantial results with accuracy improvements of up to 18.5% while simultaneously reducing response latency by 90% compared to baseline implementations. These results are particularly pronounced in enterprise-critical tasks such as cross-session information synthesis and long-term context maintenance, demonstrating Zep's effectiveness for deployment in real-world applications.

<https://zep.link/sota-paper>

+++

```
ainers/Projects/aius/aius(mark)$ show me running docker cont
↳ sudo docker ps
[sudo] password for mark:
CONTAINER ID   IMAGE          COMMAND                  CREATED        STATUS       PORTS
RTS
      NAMES
0660e5e7fc16   neo4j:latest  "tini -g -- /startup..."   4 days ago   Up 9 hours   0.
0.0.0:7474->7474/tcp, [::]:7474->7474/tcp, 7473/tcp, 0.0.0.0:7687->7687/tcp, [::]:7
  'tcp_neo4j-apoc
  ↳ objects/aius/aius(mark)$
```

```
aius:~/Projects/aius/aius(mark)$ fix
↳ echo "The command executed 3 commands ago was: 'sudo docker ps --format \"{{.Image}} {{.Status}}\"'"
The command executed 3 commands ago was: 'sudo docker ps --format "{{.Image}} {{.Status}}"'"
if apache is runnings/aius(mark)$ show the command we dic
Press Ctrl+C again to exit.
show me cpus and gpus
  ↳
```

```
Terminal      Terminal      Terminal      Terminal      Terminal      Terminal
aius:~/Projects/aius/aius(mark)$
NUMA node0 CPU(s):          0-15
Vulnerability Itlb multihit: Not affected
Vulnerability L1tf:         Not affected
Vulnerability Mds:          Not affected
Vulnerability Meltdown:     Not affected
Vulnerability Mmio stale data: Not affected
Vulnerability Retbleed:      Mitigation; untrained return thunk; SMT enabled with STIBP protection
Vulnerability Spec store bypass: Mitigation; Speculative Store Bypass disabled via prctl and seccomp
Vulnerability Spectre v1:    Mitigation; usercopy/swapgs barriers and __user pointer sanitization
Vulnerability Spectre v2:    Mitigation; Retpolines, IBPB conditional, STIBP always-on, RSB filling, PBRSB-eIBRS
Vulnerability Srbds:         Not affected
Vulnerability Tsx async abort: Not affected
01:00.0 VGA compatible controller: NVIDIA Corporation TU117M [GeForce GTX 1650 Ti Mobile] (rev a1)
05:00.0 VGA compatible controller: Advanced Micro Devices, Inc. [AMD/ATI] Renoir (r
  ↳ objects/aius/aius(mark)$
```

+++

# workshop-graphrag-chat

Unlock your team's entire message and file history **PRO**  
Upgrade to a paid plan for unlimited access to past decisions and discussions.

Learn More

Today

Mark Bain 10:39 AM joined #workshop-graphrag-chat. Also, Daniel Chalef and 4 others joined.

Mark Bain 12:14 PM Hi everyone,  
We're excited to see you and to discuss the progression in agentic memory together.  
Welcome to our Agentic Memory and GraphRAG workshop!

Follow along the workshop walkthrough: <https://docs.google.com/document/d/1UQhOGzG6sucro0DNYPisUjwly8qCzOx02RHhYY5dA78/edit?tab=t.0>

Only visible to you

Slackbot 12:14 PM That looks like a Google Drive link – do you want to connect Google Drive to Slack? You'll be able to preview files, manage access to documents, get notified about and reply to comments, and see when new files are shared with you.

Connect Not now Don't ask again

Anita Shen 1:02 PM joined #workshop-graphrag-chat. Also, Cindy Z and 11 others joined.

Alex Gilmore 1:32 PM The set up document referenced during the Neo4j memory MCP server demo  
<https://github.com/a-s-g93/ai-engineer-neo4j-memory-demo>

B I S O E M D C

Message #workshop-graphrag-chat

+ Aa @ #

File Edit View History Bookmarks Tools Help

workshop-gr Make Your A Make Your neo4j@neo4j AI Engineer Cognee Docs Python SDK Installation graphrag/gra Agent Devel Agent Devel i squared - Po Asymmetries

File Edit View Insert Format Tools Extensions Help

Make Your Agents Remember What They Do

Document tabs +

Tab 1

Make Your Agents Remember ...

Steps:

References

Hi everyone,

We're excited to see you and to discuss the progression in agentic memory together.

Welcome to our Agentic Memory and GraphRAG workshop!

We'll start with a quick overview of the repository: [GraphRAG Chat](#) (specifically this [Fork](#)). Think of it as a precursor to an agentic memory arena.

**Steps:**

Join the [slack channel: #workshop-graphrag-chat](#)

Clone the fork:

```
git clone git@github.com:markmbain/graphrag-graphrag-chat.git
graphrag-chat && cd graphrag-chat
```

Update your credentials to neo4j and provide OPENAI\_API\_KEY, optionally GOOGLE\_API\_KEY and GEMINI\_API\_KEY

```
cp .env.template .env
```

File Edit View History Bookmarks Tools Help

graphrag / graphrag-chat

Code Issues Pull requests Actions Projects Security Insights

graphrag-chat Public

main 1 Branch 0 Tags Go to file Add file Code About

akollegger added general notes about the project to readme 19dd8c5 · last month 7 Commits

graphrag\_chat converted in\_memory to a singleton, added a tool for ag... last month

.env.template initial draft of a mediated multi-agent chat last month

.gitignore initial draft of a mediated multi-agent chat last month

.python-version initial draft of a mediated multi-agent chat last month

.tool-versions initial draft of a mediated multi-agent chat last month

README.md added general notes about the project to readme last month

main.py initial draft of a mediated multi-agent chat last month

neo4j.example.json initial draft of a mediated multi-agent chat last month

pyproject.toml initial draft of a mediated multi-agent chat last month

uv.lock initial draft of a mediated multi-agent chat last month

README

## GraphRAG Chat

Moderated multi-agent group chat.

- top-level agent routes questions to sub-agents by name or speciality

File Edit View History Bookmarks Tools Help

graphrag / graphrag-chat

README Languages

## GraphRAG Chat

Moderated multi-agent group chat.

- top-level agent routes questions to sub-agents by name or speciality
- top-level agent includes multi-memory delegation
- all memories saved to each memory implementation
- sub-agents can select which memory implementation to read from
- a Neo4j graph catalog is available for read/write to multiple knowledge graphs

Agents:

```
graphrag_chat_agent_v1 (root agent)
├── cypher_agent: direct read/write access to available knowledge graphs
├── agent_smith: Neo4j product specialist
│   └── (memory: in-memory service. TBD migrate to simple KG memory)
└── new_agent: Give them a fun name and topical expertise
    └── (memory: note the memory impl)
```

## Developer Notes

### 1. Initialize the Python environment with uv

- Create a virtual environment:  
uv venv
- Activate the virtual environment:  
source .venv/bin/activate # select the appropriate activate.\* for your shell

The screenshot shows a browser window with the URL [google.github.io/adk-docs/sessions/](https://google.github.io/adk-docs/sessions/). The page title is "Agent Development Kit". The left sidebar contains a navigation menu with sections like Home, Get Started, Tutorials, Agents, Tools, Running Agents, Deploy, Sessions & Memory, Session, State, Memory, Callbacks, Artifacts, Events, Context, Evaluate, MCP, Bidi-streaming (live), Safety and Security, Agent2Agent (A2A) Protocol, Community Resources, Contributing Guide, and API Reference. The main content area has a heading "Introduction to Conversational Context: Session, State, and Memory" and a sub-section "Why Context Matters" which explains that meaningful, multi-turn conversations require agents to understand context, similar to humans. It mentions the ADK provides structured ways to manage context through Session, State, and Memory. Below this is a section "Core Concepts" with a definition of conversation threads and long-term knowledge. Two numbered lists follow: 1. Session : The Current Conversation Thread (explaining it's a single, ongoing interaction with chronological sequence of messages and actions) and 2. State (session.state) : Data Within the Current Conversation (explaining it holds temporary data relevant only to the conversation). The top right of the page shows a "Table of contents" with links to Why Context Matters, Core Concepts, Managing Context: Services, and What's Next?.

The screenshot shows a complex network graph on the left and a chat interface on the right. The graph has nodes represented by green ovals and rounded rectangles, connected by arrows indicating relationships. Nodes include 'agent\_cypher\_v1', 'graph\_is\_ready', 'get\_physical\_schema', 'read\_neo4j\_cypher', 'write\_neo4j\_cypher', 'agent\_neo\_v1', 'search\_memory', 'cognee\_add', 'agent\_cognee\_v1', 'cognee\_search', 'agent\_graphiti\_v1', 'graphiti\_add', 'agent\_mem0\_v1', 'graphiti\_search', 'mem0\_add', 'mem0\_search', and 'set\_spaker'. The right side features a chat window with a message from 'Agent Neo' and a detailed tooltip about Neo4j's contributions to open source.

SESSION ID: b6f95935-8f8f-447b-9bbb-7dc1a879f63b

Event 11 of 33

Event Request Response

graph\_is\_ready

agent\_cypher\_v1

get\_physical\_schema

read\_neo4j\_cypher

write\_neo4j\_cypher

agent\_neo\_v1

search\_memory

cognee\_add

agent\_cognee\_v1

cognee\_search

agent\_graphiti\_v1

graphiti\_add

agent\_mem0\_v1

graphiti\_search

mem0\_add

mem0\_search

set\_spaker

SESSION ID: b6f95935-8f8f-447b-9bbb-7dc1a879f63b

90%

Token Streaming | New Session

Hi

Hello! How can I assist you today?  
-- Gary

can I talk with neo

transfer\_to\_agent

✓ transfer\_to\_agent

Hi I'm Agent Neo, a product specialist for Neo4j. How can I help you today?  
-- Agent Neo

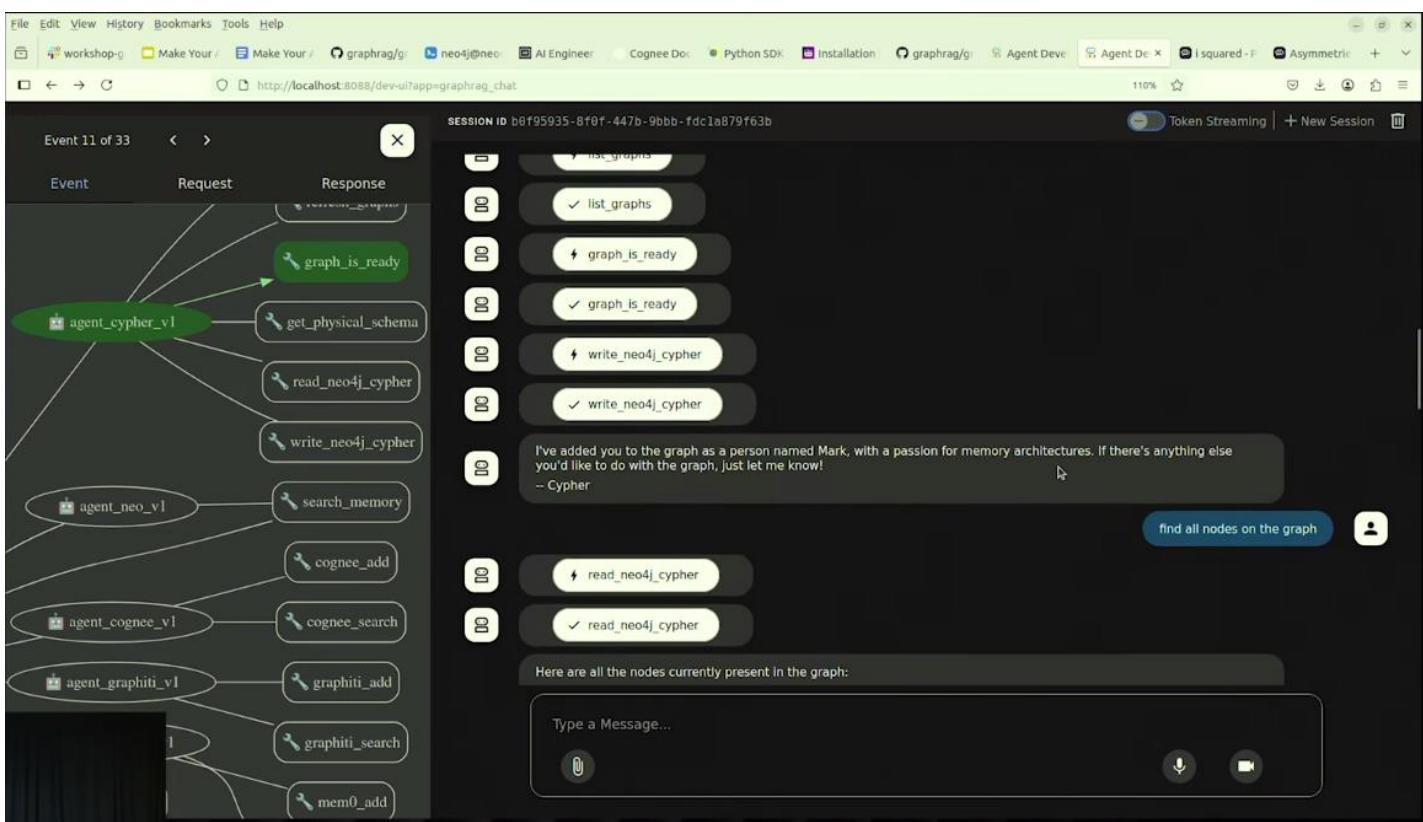
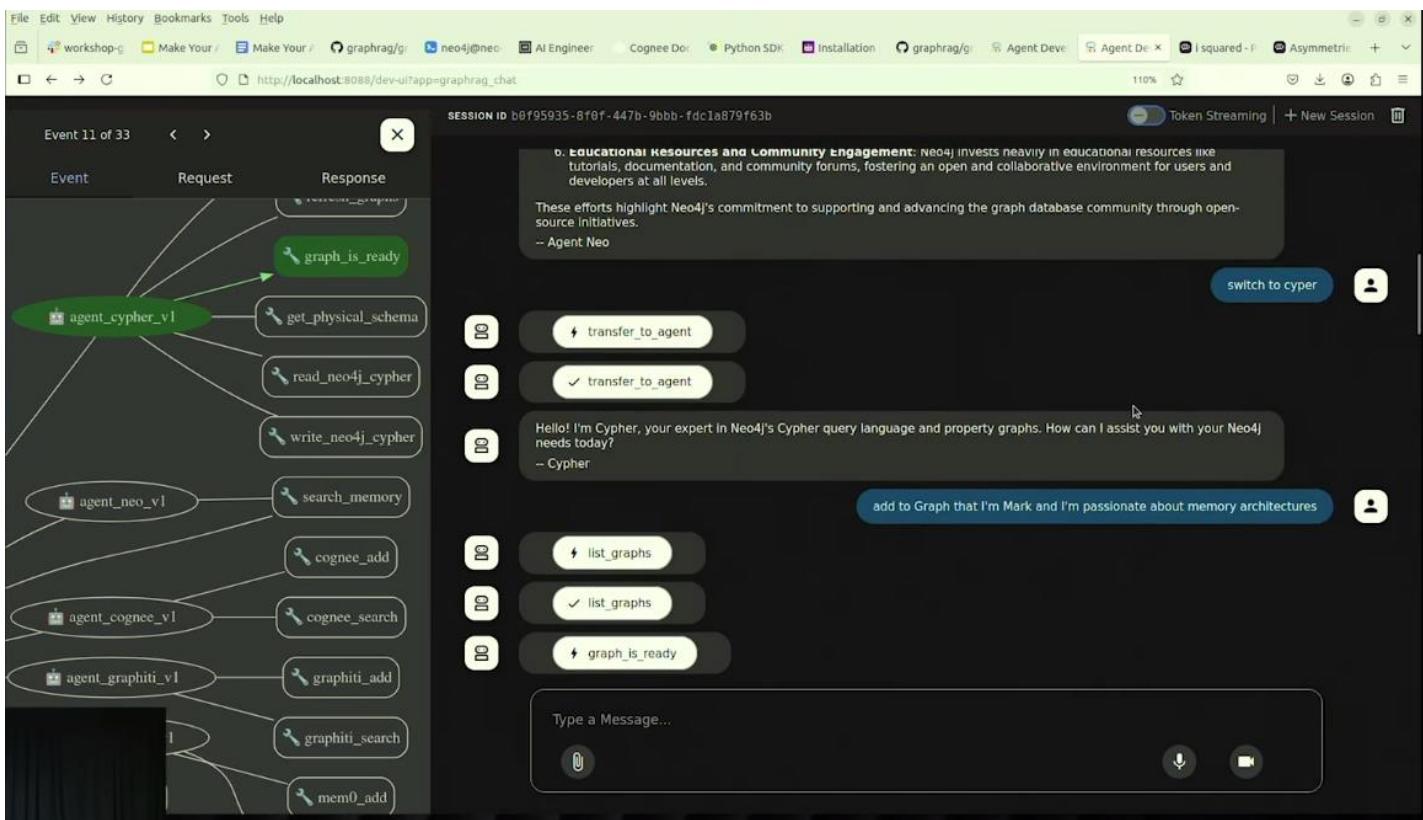
what's neo4j's contribution to open source

Neo4j is a significant contributor to the open-source community in several ways:

1. **Core Product:** The Community Edition of Neo4j is available as an open-source product, allowing developers to use the core features of Neo4j for free. This encourages learning and experimentation within the graph database community.
2. **Extensions and Tools:** Neo4j has open-sourced several extensions and tools that integrate with its database. This includes tools for data import/export, visualization, and integration with other technologies.
3. **Cypher Query Language:** Cypher, the query language developed by Neo4j for graph databases, has been made available as open-source. This has helped in standardizing a way to query graph data.
4. **Community Contributions:** Neo4j actively encourages contributions from its community. Many enhancements and features of Neo4j have been influenced by open collaborations with community developers.
5. **Graph Algorithms and GDS:** The Graph Data Science library offers an extensive collection of graph algorithms and is a crucial open-source tool for data scientists who want to leverage graph analytics.
6. **Educational Resources and Community Engagement:** Neo4j invests heavily in educational resources like tutorials, documentation, and community forums, fostering an open and collaborative environment for users and developers at all levels.

These efforts highlight Neo4j's commitment to supporting and advancing the graph database community through open-source initiatives.  
-- Agent Neo

Type a Message...



File Edit View History Bookmarks Tools Help

SESSION ID b8f95935-8f0f-447b-9bbb-fd1a879f63b

Event 11 of 33 < > X

Event Request Response

you'd like to do with the graph, just let me know!

-- Cypher

find all nodes on the graph

read\_neo4j\_cypher

read\_neo4j\_cypher

Here are all the nodes currently present in the graph:

1. Person: Mark, passionate about memory architectures.
2. DocumentChunk: "Kiran has a cat and three dogs. Betty, Catty, and Dotty."
3. Entity:
  - dotty (The third dog owned by Kiran)
  - animal
  - betty (One of the three dogs owned by Kiran)
  - kiran (A person who has pets)
  - cat (A domestic cat owned by Kiran)
  - catty (Another dog owned by Kiran)
  - clay bavor (Co-founder of the AI Engineer World's Fair, speaker on AI Architectures)
  - dani grant (CEO and speaker on funding strategies for AI Engineers)
  - ai engineer worlds fair 2025 (The largest technical conference for engineers working in AI, held in San Francisco)
  - ai architects track (A dedicated track at the AI Engineer World's Fair focusing on AI architectures and related topics)
  - natural language processing (A subfield of computer science focused on the interaction between computers and human language)
  - computer science (The study of computers and computational systems)
4. EntityType: Entity, EntityType, event, concept, field, session track, person
5. TextDocument:
  - text.b8122bf4c463a0d17ef0704f17ccaa69
  - text.96d81046997835a9f2d3898bb2c514
  - text.4dd4b1ccc64682c225674f36ba9d3ac
6. DocumentChunk: Details about AI Engineer World's Fair 2025
7. TextSummary:
  - "Kiran owns a feline and a trio of canines."
  - "AI Engineer World's Fair 2025 is a leading technical conference for AI professionals..."
  - "NLP is a branch of computer science."
8. EntityType: Dog

Type a Message...

File Edit View History Bookmarks Tools Help

SESSION ID b8f95935-8f0f-447b-9bbb-fd1a879f63b

Event 11 of 33 < > X

Event Request Response

Let me know if you need more information or need to perform other operations!

-- Cypher

Add to the graphVasilije will show how integrating cognee memory enables agents to store and retrieve data they need to optimize their tasks in different memory layers, and how these layers can be used as dynamic system prompts, empowering LLM Apps and Agentic frameworks.

write\_neo4j\_cypher

Type a Message...

