

## Pi0 - generalist Vision Language Action policy for robots (VLA Series Ep.2)



Ilia  
3.63K subscribers



350



Share

Ask

Save

...

8,708 views Sep 16, 2025 #Robotics #transformers #AI

⌚ The second video in the series about Visual Language Action policies for robotics!

This time, I will focus on one of the most notable open-source VLA models, Pi0 from Physical Intelligence.

🕒 What You'll Learn in This Episode:

- ✓ The model architecture and the overview
- ✓ Fine-tuning on your own robot
- ✓ FAST tokenization vs. Flow Matching
- ✓ LoRA vs Full fine-tuning
- ✓ Practical demo with LeKiwi and LeRobot

🔥 Coming Up in This Series:

- Pi-0.5 - the more advanced version of Pi 0
- SmolVLA from HuggingFace LeRobot
- Gr0ot N1.5 from NVIDIA
- And other related topics

0:00 Intro

3:44 Model Architecture

8:21 Action Expert Details

22:30 Fine-tuning with OpenPI

52:01 Experiments with My Robot

1:14:12 Outro

🎬 Series Playlist: VLA policies in robotics

🔗 Resources & Code:

- Pi-0 Blog post: <https://www.physicalintelligence.com...>
- Pi-0 Paper: <https://www.physicalintelligence.com...>
- OpenPI Github: <https://github.com/Physical-Intellige...>

← → ✐ https://www.physicalintelligence.company/blog/pi0



## Physical Intelligence ( $\pi$ )

[Home](#) [Research](#) [Join Us](#)

# $\pi_0$ : Our First Generalist Policy

Published

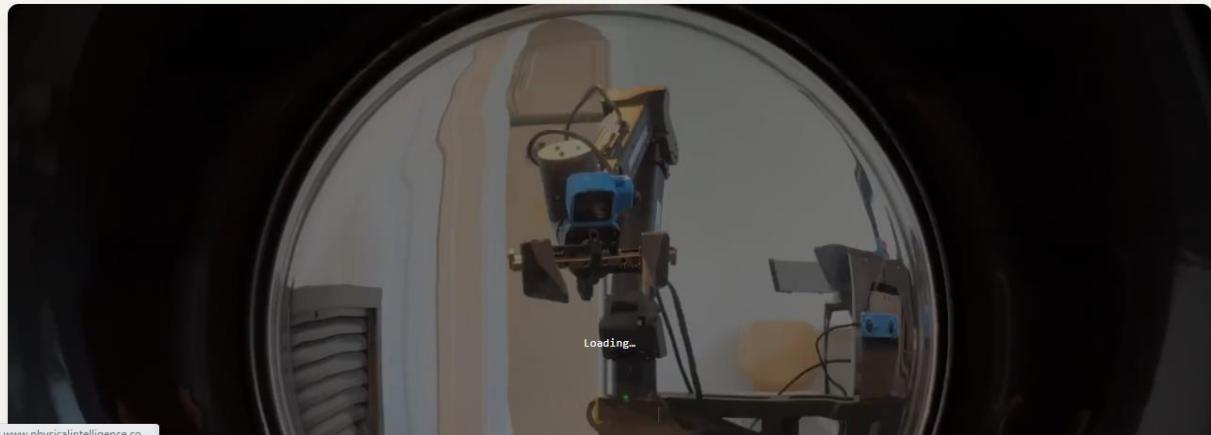
October 31, 2024

Email

[research@physicalintelligence.company](mailto:research@physicalintelligence.company)

Paper

[π\\_0.pdf](#)



Waiting for www.physicalintelligence.co...

Physical Intelligence / openpi (Public)

[Code](#) [Issues 178](#) [Pull requests 32](#) [Discussions](#) [Actions](#) [Projects](#) [Security](#) [Insights](#)

[main](#) [6 Branches](#) [0 Tags](#)

[kpertsch Multi-RLDS Datasets \(Cotraining\) + PolaRIS DROID jointpos policy conf...](#) [bd70b8f · 3 days ago](#) [186 Commits](#)

File	Commit Message	Date
.github	Upgrade LeRobot	7 months ago
.vscode	Initial commit	11 months ago
docs	broken link fix	2 months ago
examples	Fix typo in comment for read_camera method	2 weeks ago
packages/openpi-client	feat: add wss support	2 months ago
scripts	Fix typo in comment about sudo requirement	2 weeks ago
src/openpi	fix ruff formatting	last week
third_party	Initial commit	11 months ago
.dockerignore	Initial commit	11 months ago

[About](#)  
*No description, website, or topics provided.*

[Readme](#) [Apache-2.0 license](#) [Contributing](#) [Activity](#) [Custom properties](#) [9.4k stars](#) [78 watching](#) [1.3k forks](#) [Report repository](#)

[Releases](#)  
No releases published

[Backups](#)

Physical Intelligence ( $\pi$ ) [Home](#) [Research](#) [Join Us](#)

physicalintelligence.company

Physical Intelligence is bringing general-purpose AI into the physical world. We are a group of engineers, scientists, roboticists, and company builders developing foundation models and learning algorithms to power the robots of today and the physically-actuated devices of the future.

- Real-Time Action Chunking with Large Models** June 9, 2025  
A real-time system for large VLAs that maintains precision and speed in the face of high latency.
- VLAs that Train Fast, Run Fast, and Generalize Better** May 28, 2025  
A method to train vision-language-action models that train quickly, maintain internet-scale knowledge, have high quality inference properties, and generalize well.
- #0.5: a VLA with Open-World Generalization** April 22, 2025  
Our latest generalist policy, #0.5, extends m0 and enables open-world generalization. Our new model can control a mobile manipulator to clean up an entirely new kitchen or bedroom.
- Teaching Robots to Listen and Think Harder** February 26, 2025  
A method for robots to think through complex tasks step by step, incorporating human-in-the-loop feedback.
- Open Sourcing #0** February 4, 2025  
We are releasing the weights and code for m0 as well as our new m0-FAST autoregressive model.
- FAST: Efficient Robot Action Tokenization** January 16, 2025  
A new robot action tokenizer that allows us to train generalist policies 5x faster than previous models.



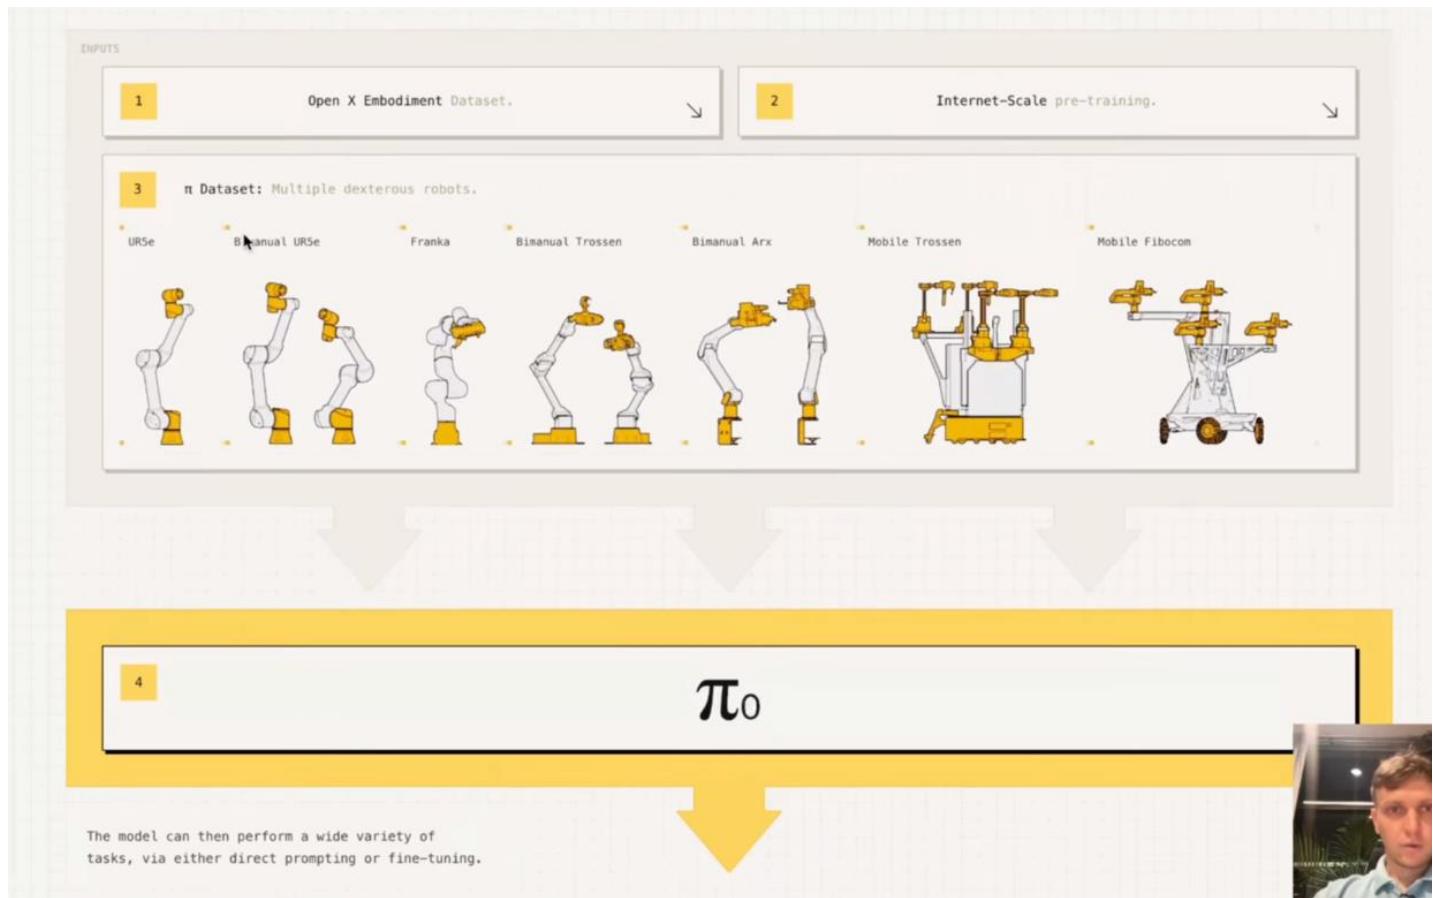
# $\pi_0$ : Our First Generalist Policy

Published  
Email  
Paper

October 31, 2024  
[research@physicalintelligence.company](mailto:research@physicalintelligence.company)  
[Download PDF](#)



<https://www.physicalintelligence.company/download/pi0.pdf>

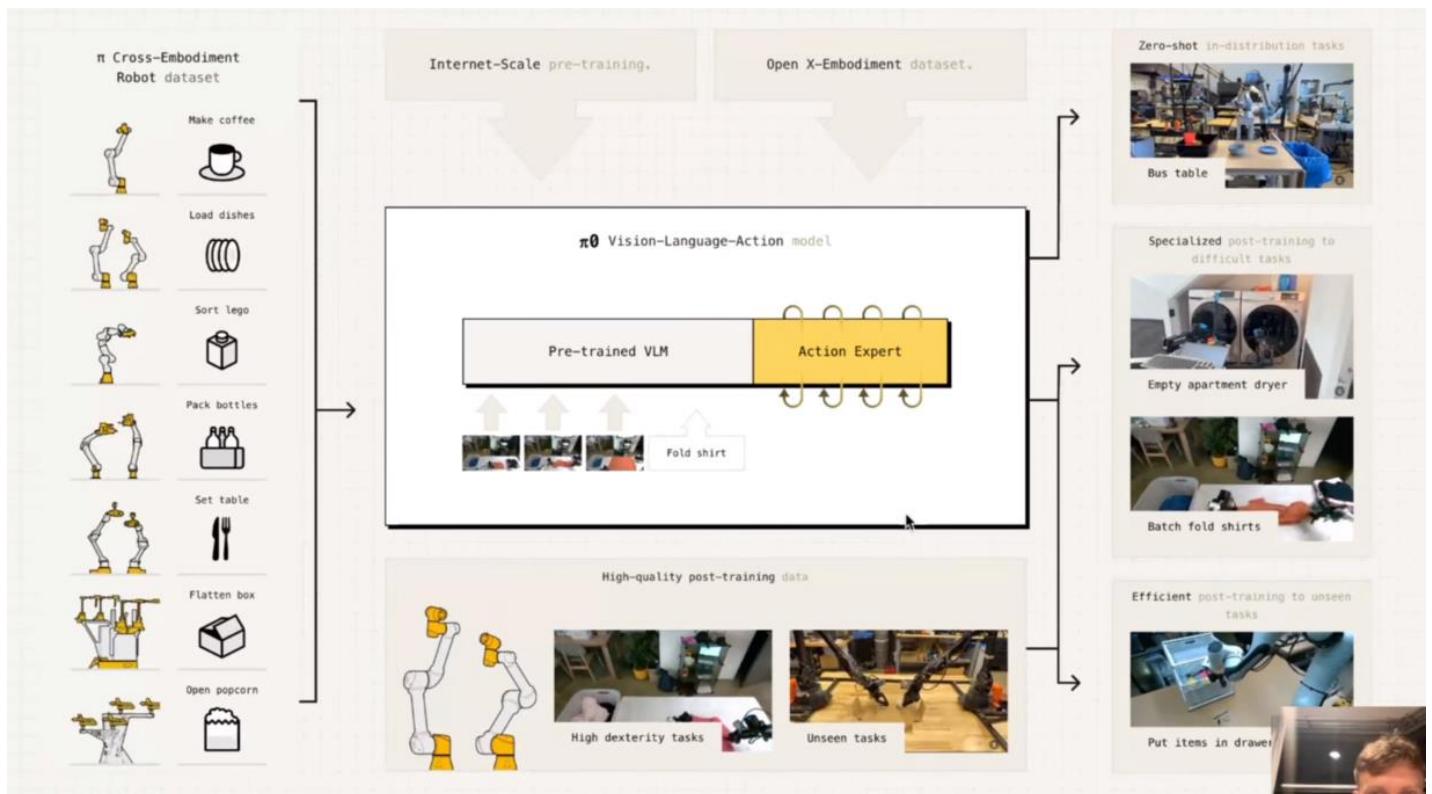


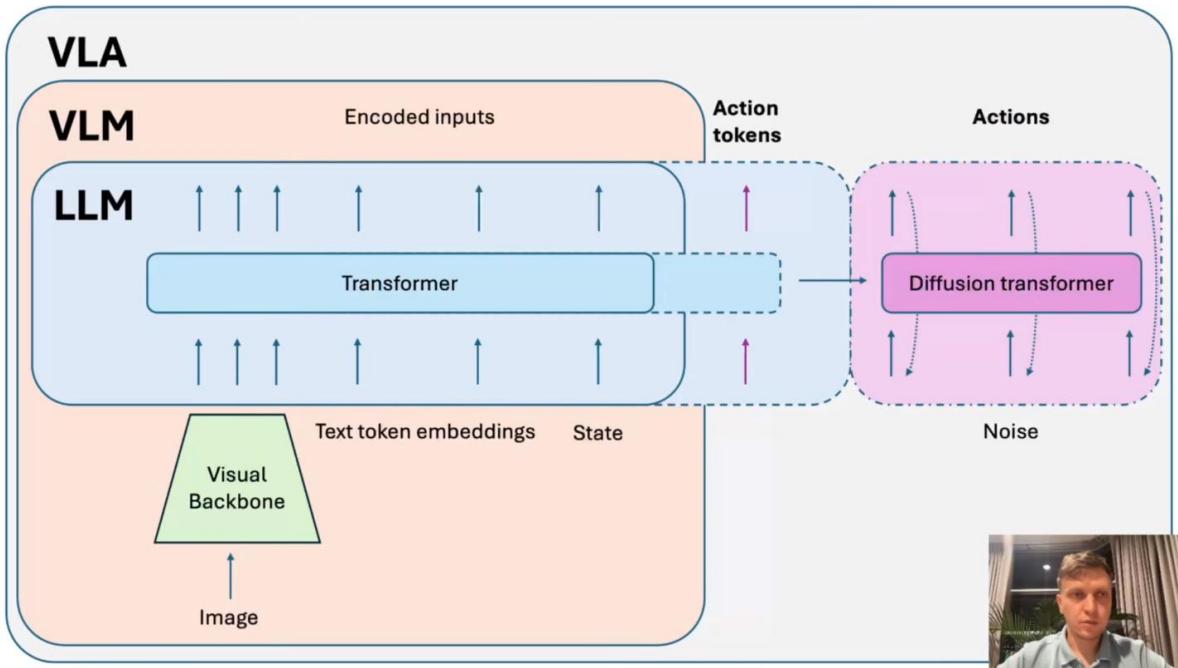
# $\pi_0$

The model can then perform a wide variety of tasks, via either direct prompting or fine-tuning.

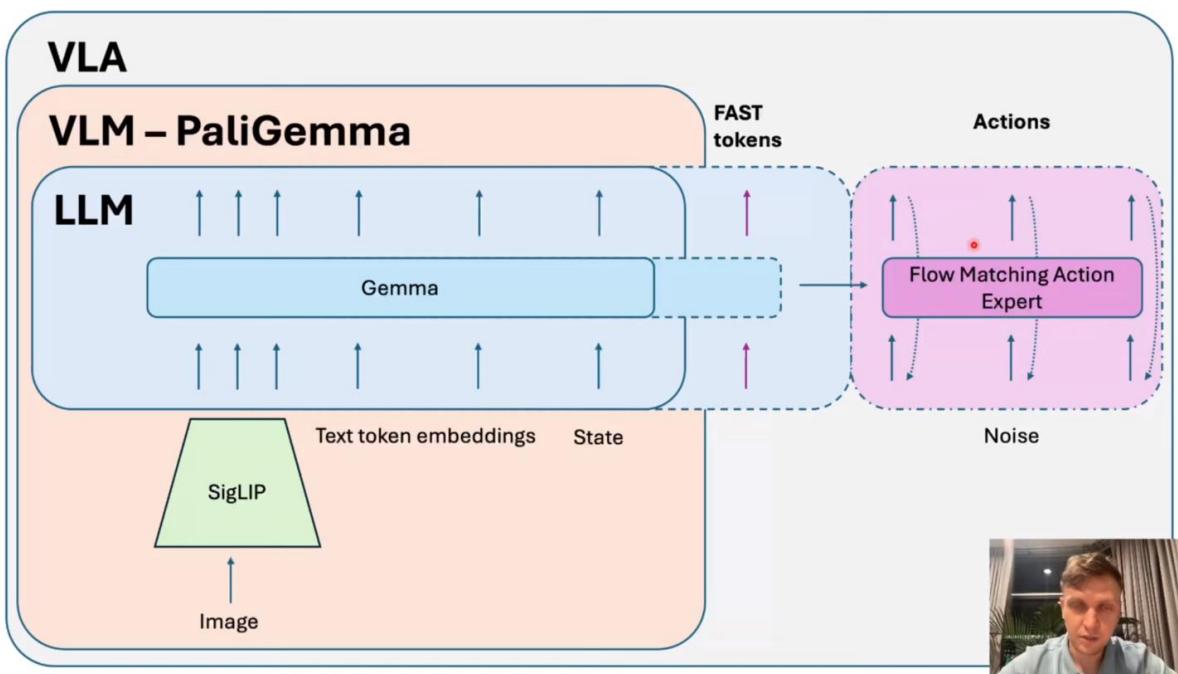


Our first prototype generalist robot policy is trained on the largest robot interaction dataset to date. The full training mixture includes both open-source data and a large and diverse dataset of dexterous tasks that we collected across 8 distinct robots.

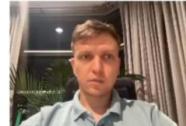
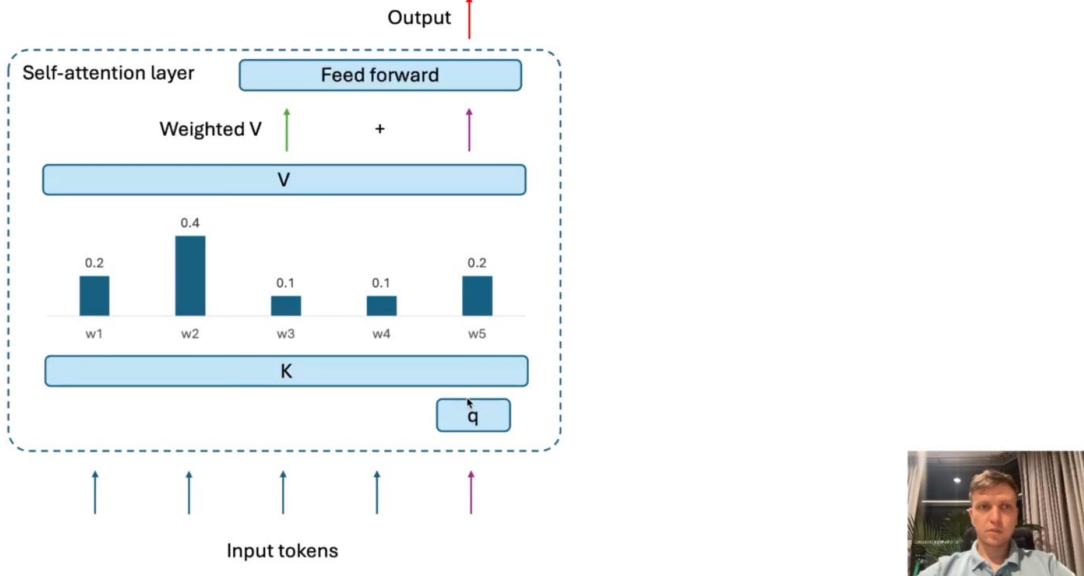




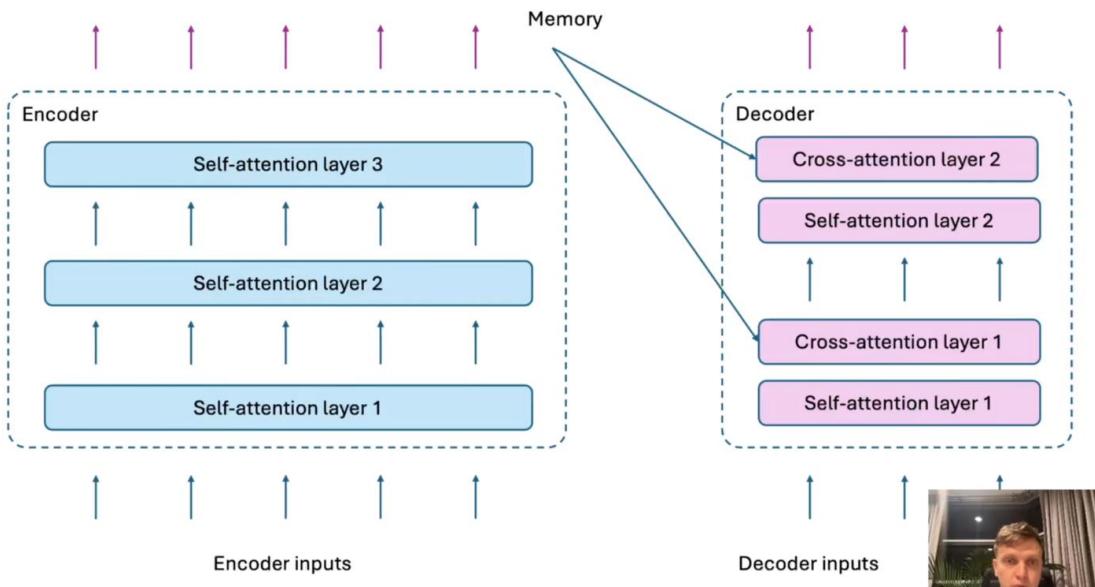
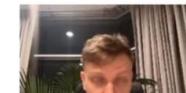
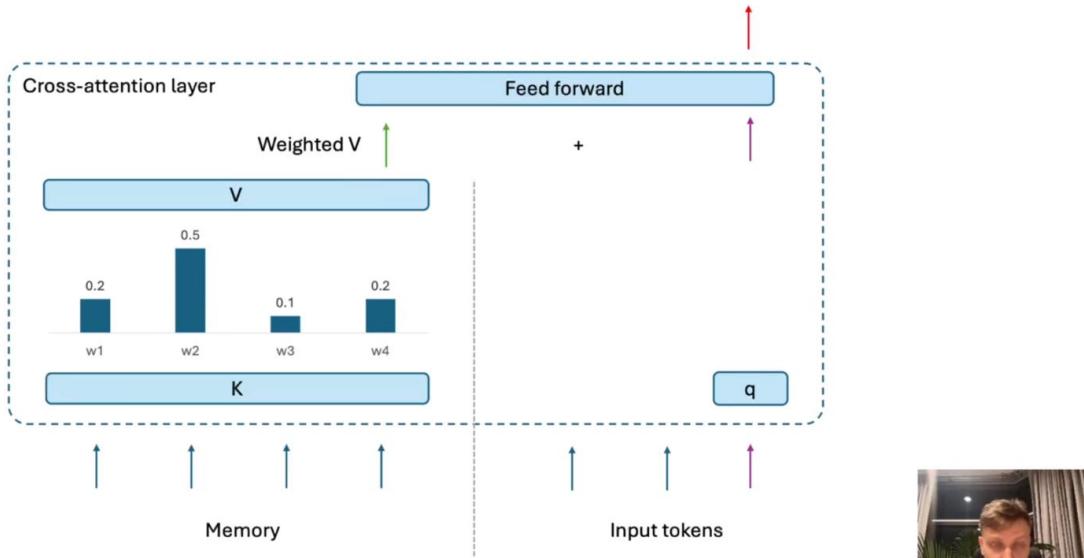
We can now identify the specific models used in the **Pi0 VLA architecture** below

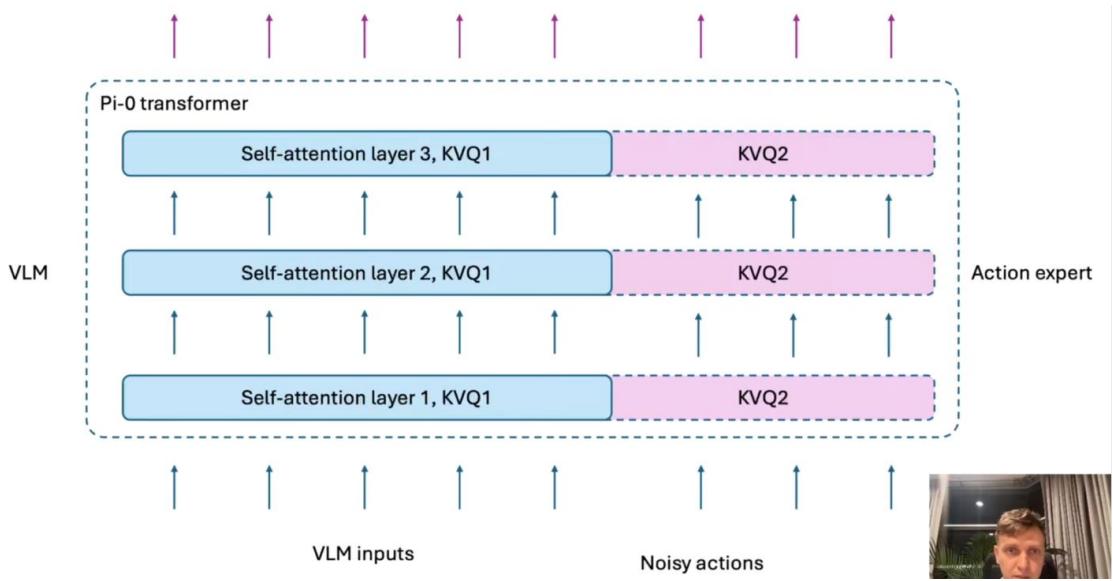


Gemma + SigLiP = PaliGemma model. FAST is the way of predicting Action tokens directly from the VLM developed by Physical Intelligence using image compression principles.

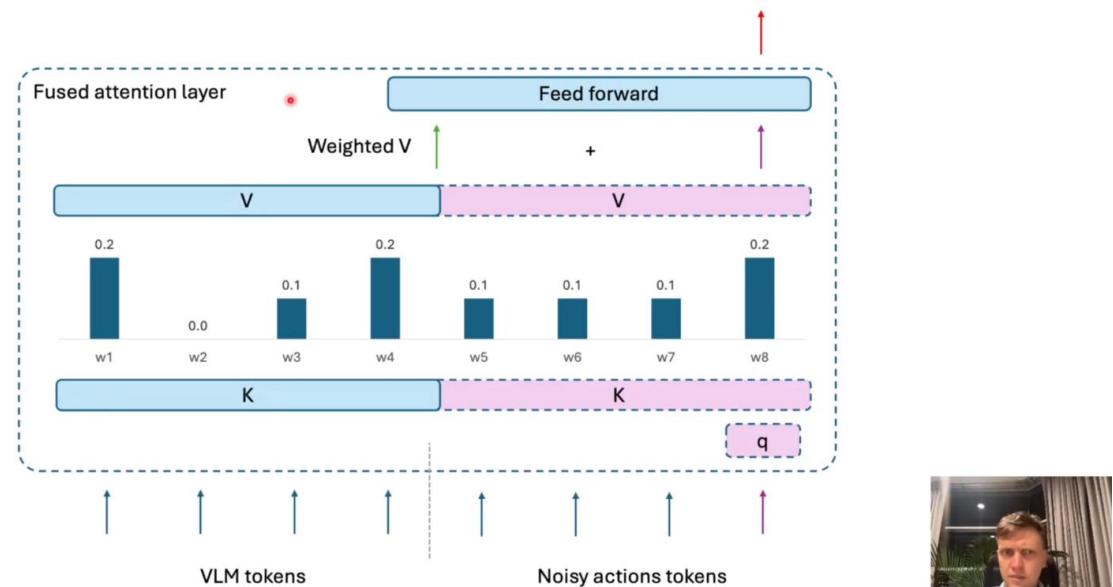


Let us see how Pi0 generates Action tokens from the Query, Key, Attention Weights, Values and Outputs. The Self-Attention layer looks at all the **Input data** tokens while the Cross-Attention only looks at the **Memory** tokens.





This is what we actually have in the PiO transformer, it uses only self-attention layers and no cross-attention layers at all. The VLM part is blue and the Action expert part is pink that tries to denoise using flow matching approach.



# uses Internet-scale vision-language pre-training, open-source robot manipulation datasets, and our own datasets consisting of dexterous tasks from 8 distinct robots as inputs.

The diagram illustrates the X Embodiment pipeline:

- INPUTS**: Open X Embodiment Dataset.
- Internet-Scale pre-training.**
- Dataset**: Multiple dexterous robots. The robots shown are:
  - UR5e
  - Bimanual UR5e
  - Franka
  - Bimanual Trossen
  - Bimanual Arx
  - Mobile Trossen
  - Mobile Fibocom

4

# $\pi_0$

The model can then perform a wide variety of tasks, via either direct prompting or fine-tuning.



This policy doesn't have any robot-specific action tokens, meaning it is robust enough to work on any robot as long as the model was fine-tuned for the robot at startup.

The currently existing VLAs are not supposed to be used Zero-shot, you are expected to fine-tune the model. You take the pretrained Pi0 model, then collect some small dataset using your particular robot from when your robot completes the particular tasks that you need, like pick and place some objects. Then do some extra training of the model of the policy on your dataset like 15 minutes of data for 55 examples in total.

Physical-Intelligence / openpi

Code Issues Pull requests Discussions Actions Projects Security Insights

openpi Public

Watch 69 Fork 729 Starred 6.1k

main 6 Branches 0 Tags Go to file Add file Code

kvablock	Update README	5bfff19b · 18 hours ago	156 Commits
.github	Upgrade LeRobot	4 months ago	
.vscode	Initial commit	7 months ago	
docs	Update references to GCP assets	3 months ago	
examples	DROID readme updates	3 days ago	
packages/openpi-client	kept the tree.map_structure call and perform the check in...	3 months ago	
scripts	Fix PyTorch assets save path	2 days ago	
src/openpi	Pin PyTorch and fix tests	3 days ago	
third_party	Initial commit	7 months ago	
.dockerignore	Initial commit	7 months ago	

About

No description, website, or topics provided.

Readme Apache-2.0 license Contributing Activity Custom properties 6.1k stars 69 watching 729 forks Report repository

Releases

No releases published

## openpi

openpi holds open-source models and packages for robotics, published by the [Physical Intelligence team](#).

Currently, this repo contains three types of models:

- the  [\$\pi\_0\$  model](#), a flow-based vision-language-action model (VLA).
- the  [\$\pi\_0\$ -FAST model](#), an autoregressive VLA, based on the FAST action tokenizer.
- the  [\$\pi\_{0.5}\$  model](#), an upgraded version of  $\pi_0$  with better open-world generalization trained with [knowledge insulation](#). Note that, in this repository, we currently only support the flow matching head for both  $\pi_{0.5}$  training and inference.

For all models, we provide *base model* checkpoints, pre-trained on 10k+ hours of robot data, and examples for using them out of the box or fine-tuning them to your own datasets.

This is an experiment:  $\pi_0$  was developed for our own robots, which differ from the widely used platforms such as [ALOHA](#) and [DROID](#), and though we are optimistic that researchers and practitioners will be able to run creative new experiments adapting  $\pi_0$  to their own platforms, we do not expect every such attempt to be successful. All this is to say:  $\pi_0$  may or may not work for you, but you are welcome to try it and see!

### Updates

- [Sept 2025] We released PyTorch support in openpi.
- [Sept 2025] We released pi05, an upgraded version of pi0 with better open-world generalization.
- [Sept 2025]: We have added an [improved idle filter](#) for DROID training.
- [Jun 2025]: We have added [instructions](#) for using `openpi` to train VLAs on the full [DROID dataset](#). This is an approximate open-source implementation of the training pipeline used to train pi0-FAST-DROID.



your own data. We will explain three steps:

1. Convert your data to a LeRobot dataset (which we use for training)
2. Defining training configs and running training
3. Spinning up a policy server and running inference

#### 1. Convert your data to a LeRobot dataset

We provide a minimal example script for converting LIBERO data to a LeRobot dataset in [examples/libero/convert\\_libero\\_data\\_to\\_lerobot.py](#). You can easily modify it to convert your own data! You can download the raw LIBERO dataset from [here](#), and run the script with:

```
uv run examples/libero/convert_libero_data_to_lerobot.py --data_dir /path/to/your/libero/data
```

Note: if you just want to fine-tune on LIBERO, you can skip this step, because our LIBERO fine-tuning configs point to a pre-converted LIBERO dataset. This step is merely an example that you can adapt to your own data.

#### 2. Defining training configs and running training

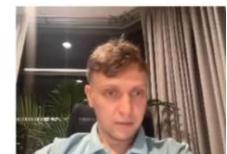
To fine-tune a base model on your own data, you need to define configs for data processing and training. We provide example configs with detailed comments for LIBERO below, which you can modify for your own dataset:

- [LiberoInputs](#) and [LiberoOutputs](#): Defines the data mapping from the LIBERO environment to the model and vice versa. Will be used for both, training and inference.
- [LeRobotLiberoDataConfig](#): Defines how to process raw LIBERO data from LeRobot dataset for training.
- [TrainConfig](#): Defines fine-tuning hyperparameters, data config, and weight loader.

We provide example fine-tuning configs for  [\$\pi\_0\$](#) ,  [\$\pi\_0\$ -FAST](#), and  [\$\pi\_{0.5}\$](#)  on LIBERO data.

Before we can run training, we need to compute the normalization statistics for the training data. Run the script below with the name of your training config:

```
uv run scripts/compute_norm_stats.py --config-name pi05_libero
```



You can follow the instructions to finetune the model.

## Requirements

To run the models in this repository, you will need an NVIDIA GPU with at least the following specifications. These estimations assume a single GPU, but you can also use multiple GPUs with model parallelism to reduce per-GPU memory requirements by configuring `fsdp_devices` in the training config. Please also note that the current training script does not yet support multi-node training.

Mode	Memory Required	Example GPU
Inference	> 8 GB	RTX 4090
Fine-Tuning (LoRA)	> 22.5 GB	RTX 4090
Fine-Tuning (Full)	> 70 GB	A100 (80GB) / H100

The repo has been tested with Ubuntu 22.04, we do not currently support other operating systems.

## Installation

When cloning this repo, make sure to update submodules:

```
git clone --recurse-submodules git@github.com:Physical-Intelligence/openpi.git
# Or if you already cloned the repo:
git submodule update --init --recursive
```

We use [uv](#) to manage Python dependencies. See the [uv installation instructions](#) to set it up. Once uv is installed, run the following to set up the environment:

```
GIT_LFS_SKIP_SMUDGE=1 uv sync
GIT_LFS_SKIP_SMUDGE=1 uv pip install -e .
```

NOTE: `GIT_LFS_SKIP_SMUDGE=1` is needed to pull LeRobot as a dependency.



Screenshot of a GitHub code review interface for a file named `setup_remote.sh`. The file contains a bash script for setting up a remote machine. The script installs various dependencies, including apt-get updates, specific libraries (libgl1-mesa-glx, libglfw3, libglew-dev, xorg-dev), and git-lfs. It also checks for conda availability and installs miniconda3 if not found. The commit message is "Added setup script".

```
#!/usr/bin/env bash
# I use this script to setup a remote machine to use LeRobot and other tools
# It installs different dependencies that can be missing
# The list can be redundant in some cases and not complete in others
# But it covers most of the cases I faced when worked with LeRobot
# It can take around 10 minutes to install all the dependencies
# Just run 'bash setup_remote.sh' to install all the dependencies
set -euo pipefail
# 1) Apt deps
# Some dependencies can be redundant but I included all that caused me issues in the past
$SUDO=""
if [ "${EUID:-$(id -u)}" -ne 0 ]; then $SUDO=sudo; fi
export DEBIAN_FRONTEND=noninteractive
$SUDO apt-get update
$SUDO apt-get install -y \
curl htop git-lfs gcc build-essential python3-dev linux-libc-dev \
libgl1-mesa-glx libgl1-mesa-dev libglfw3-dev libglew-dev xorg-dev \
libsolv1-dev
git lfs install || true
# 2) Ensure conda is available and usable in this script
if ! command -v conda >/dev/null 2>&1; then
echo "Conda not found, installing Miniconda to $HOME/miniconda3..."
tmpd=$(mktemp -d)
curl -fSSL https://repo.anaconda.com/miniconda/Miniconda3-latest-Linux-x86_64.sh -o "$tmpd/miniconda.sh"
bash "$tmpd/miniconda.sh" -b -n "miniconda3"
```

You can use this script if you want to train your model.



## Installation

When cloning this repo, make sure to update submodules:

```
git clone --recurse-submodules git@github.com:Physical-Intelligence/openpi.git  
# Or if you already cloned the repo:  
git submodule update --init --recursive
```

We use [uv](#) to manage Python dependencies. See the [uv installation instructions](#) to set it up. Once uv is installed, run the following to set up the environment:

```
GIT_LFS_SKIP_SMUDGE=1 uv sync  
GIT_LFS_SKIP_SMUDGE=1 uv pip install -e .
```

NOTE: `GIT_LFS_SKIP_SMUDGE=1` is needed to pull LeRobot as a dependency.

**Docker:** As an alternative to uv installation, we provide instructions for installing openpi using Docker. If you encounter issues with your system setup, consider using Docker to simplify installation. See [Docker Setup](#) for more details.

## Model Checkpoints

### Base Models

We provide multiple base VLA model checkpoints. These checkpoints have been pre-trained on 10k+ hours of robot data, and can be used for fine-tuning.

Model	Use Case	Description	Checkpoint Path
$\pi_0$	Fine-Tuning	Base <a href="#"><math>\pi_0</math> model</a> for fine-tuning	<code>gs://openpi-assets/checkpoints/pi0_base</code>



After setting up the machine, you need to set up OpenPi using `git clone` and install `uv`, then follow the instructions.

## Model Checkpoints

### Base Models

We provide multiple base VLA model checkpoints. These checkpoints have been pre-trained on 10k+ hours of robot data, and can be used for fine-tuning.

Model	Use Case	Description	Checkpoint Path
$\pi_0$	Fine-Tuning	Base <a href="#"><math>\pi_0</math> model</a> for fine-tuning	<code>gs://openpi-assets/checkpoints/pi0_base</code>
$\pi_0$ -FAST	Fine-Tuning	Base autoregressive <a href="#"><math>\pi_0</math>-FAST model</a> for fine-tuning	<code>gs://openpi-assets/checkpoints/pi0_fast_base</code>
$\pi_{0.5}$	Fine-Tuning	Base <a href="#"><math>\pi_{0.5}</math> model</a> for fine-tuning	<code>gs://openpi-assets/checkpoints/pi05_base</code>

### Fine-Tuned Models

We also provide "expert" checkpoints for various robot platforms and tasks. These models are fine-tuned from the base models above and intended to run directly on the target robot. These may or may not work on your particular robot. Since these checkpoints were fine-tuned on relatively small datasets collected with more widely available robots, such as ALOHA and the DROID Franka setup, they might not generalize to your particular setup, though we found some of these, especially the DROID checkpoint, to generalize quite broadly in practice.

Model	Use Case	Description	Checkpoint Path
$\pi_0$ -FAST-DROID	Inference	$\pi_0$ -FAST model fine-tuned on the <a href="#">DROID dataset</a> ; can perform a wide range of simple table-top manipulation tasks 0-shot in new scenes on the DROID	<code>gs://openpi-assets/checkpoints/pi0_fast_droid</code>



Model	Use Case	Description	Checkpoint Path
$\pi_0$ -FAST-DROID	Inference	$\pi_0$ -FAST model fine-tuned on the <a href="#">DROID dataset</a> : can perform a wide range of simple table-top manipulation tasks 0-shot in new scenes on the DROID robot platform	<a href="gs://openpi-assets/checkpoints/pi0_fast_droid">gs://openpi-assets/checkpoints/pi0_fast_droid</a>
$\pi_0$ -DROID	Fine-Tuning	$\pi_0$ model fine-tuned on the <a href="#">DROID dataset</a> : faster inference than $\pi_0$ -FAST-DROID, but may not follow language commands as well	<a href="gs://openpi-assets/checkpoints/pi0_droid">gs://openpi-assets/checkpoints/pi0_droid</a>
$\pi_0$ -ALOHA-towel	Inference	$\pi_0$ model fine-tuned on internal <a href="#">ALOHA</a> data: can fold diverse towels 0-shot on ALOHA robot platforms	<a href="gs://openpi-assets/checkpoints/pi0_aloha_towel">gs://openpi-assets/checkpoints/pi0_aloha_towel</a>
$\pi_0$ -ALOHA-tupperware	Inference	$\pi_0$ model fine-tuned on internal <a href="#">ALOHA</a> data: can unpack food from a tupperware container	<a href="gs://openpi-assets/checkpoints/pi0_aloha_tupperware">gs://openpi-assets/checkpoints/pi0_aloha_tupperware</a>
$\pi_0$ -ALOHA-pen-uncap	Inference	$\pi_0$ model fine-tuned on public <a href="#">ALOHA</a> data: can uncap a pen	<a href="gs://openpi-assets/checkpoints/pi0_aloha_pen_uncap">gs://openpi-assets/checkpoints/pi0_aloha_pen_uncap</a>
$\pi_{0.5}$ -LIBERO	Inference	$\pi_{0.5}$ model fine-tuned for the <a href="#">LIBERO</a> benchmark: gets state-of-the-art performance (see <a href="#">LIBERO README</a> )	<a href="gs://openpi-assets/checkpoints/pi05_libero">gs://openpi-assets/checkpoints/pi05_libero</a>
		$\pi_{0.5}$ model fine-tuned on the	



You can also use a finetuned model if you have one of the above robots

## Fine-Tuning Base Models on Your Own Data

We will fine-tune the  $\pi_{0.5}$  model on the [LIBERO dataset](#) as a running example for how to fine-tune a base model on your own data. We will explain three steps:

1. Convert your data to a LeRobot dataset (which we use for training)
2. Defining training configs and running training
3. Spinning up a policy server and running inference

### 1. Convert your data to a LeRobot dataset

We provide a minimal example script for converting LIBERO data to a LeRobot dataset in [examples/libero/convert\\_libero\\_data\\_to\\_lerobot.py](examples/libero/convert_libero_data_to_lerobot.py). You can easily modify it to convert your own data! You can download the raw LIBERO dataset from [here](#), and run the script with:

```
uv run examples/libero/convert_libero_data_to_lerobot.py --data_dir /path/to/your/libero/data
```

**Note:** If you just want to fine-tune on LIBERO, you can skip this step, because our LIBERO fine-tuning configs point to a pre-converted LIBERO dataset. This step is merely an example that you can adapt to your own data.

### 2. Defining training configs and running training

To fine-tune a base model on your own data, you need to define configs for data processing and training. We provide example configs with detailed comments for LIBERO below, which you can modify for your own dataset:

- [LiberoInputs and LiberoOutputs](#) : Defines the data mapping from the LIBERO environment to the model and vice versa. Will be used for both, training and inference.
- [LeRobotLiberoDataConfig](#) : Defines how to process raw LIBERO data from LeRobot dataset for training.
- [TrainConfig](#) : Defines fine-tuning hyperparameters, data config, and weight loader.

We provide example fine-tuning configs for  $\pi_0$ ,  $\pi_0$ -FAST, and  $\pi_{0.5}$  on LIBERO data.

Before we can run training, we need to compute the normalization statistics for the training data. Run the script



## 2. Defining training configs and running training

To fine-tune a base model on your own data, you need to define configs for data processing and training. We provide example configs with detailed comments for LIBERO below, which you can modify for your own dataset:

- [LiberoInputs and LiberoOutputs](#) : Defines the data mapping from the LIBERO environment to the model and vice versa. Will be used for both, training and inference.
- [LeRobotLiberoDataConfig](#) : Defines how to process raw LIBERO data from LeRobot dataset for training.
- [TrainConfig](#) : Defines fine-tuning hyperparameters, data config, and weight loader.

We provide example fine-tuning configs for [π₀](#), [π₀-FAST](#), and [π₀-s](#) on LIBERO data.

Before we can run training, we need to compute the normalization statistics for the training data. Run the script below with the name of your training config:

```
uv run scripts/compute_norm_stats.py --config-name pi05_libero
```



Now we can kick off training with the following command (the `--overwrite` flag is used to overwrite existing checkpoints if you rerun fine-tuning with the same config):

```
XLA_PYTHON_CLIENT_MEM_FRACTION=0.9 uv run scripts/train.py pi05_libero --exp-name=my_experimen
```



The command will log training progress to the console and save checkpoints to the `checkpoints` directory. You can also monitor training progress on the Weights & Biases dashboard. For maximally using the GPU memory, set `XLA_PYTHON_CLIENT_MEM_FRACTION=0.9` before running training -- this enables JAX to use up to 90% of the GPU memory (vs. the default of 75%).

**Note:** We provide functionality for *reloading* normalization statistics for state / action normalization from pre-training. This can be beneficial if you are fine-tuning to a new task on a robot that was part of our pre-training mixture. For more details on how to reload normalization statistics, see the [norm\\_stats.md](#) file.



## 3. Spinning up a policy server and running inference

Physical-Intelligence / openpi

Code Issues 79 Pull requests 19 Discussions Actions Projects Security Insights

main openpi / src / openpi / policies / libero\_policy.py

kvablick Fix normalize/padding ordering 59d1371 · 3 days ago History

Code Blame 100 lines (81 loc) · 4.22 KB

```

1 import dataclasses
2
3 import einops
4 import numpy as np
5
6 from openpi import transforms
7 from openpi.models import model as _model
8
9
10 def make_libero_example() -> dict:
11     """Creates a random input example for the Libero policy."""
12     return {
13         "observation/state": np.random.rand(8),
14         "observation/image": np.random.randint(256, size=(224, 224, 3), dtype=np.uint8),
15         "observation/wrist_image": np.random.randint(256, size=(224, 224, 3), dtype=np.uint8),
16         "prompt": "do something",
17     }
18
19
20 def _parse_image(image) -> np.ndarray:
21     image = np.asarray(image)
22     if np.issubdtype(image.dtype, np.floating):
23         image = (255 * image).astype(np.uint8)
24     if image.shape[0] == 3:
25         image = einops.rearrange(image, "c h w -> h w c")
26     return image
27

```

Raw

Symbols

Find definitions and references for functions and other symbols in this file by clicking a symbol below or in the code.

Filter symbols

- func make\_libero\_example
- func \_\_parse\_image
- class LiberoInputs
  - const model\_type
  - func \_\_call\_\_
- class LiberoOutputs
  - func \_\_call\_\_

Create this config file for your robot if it doesn't exist yet. Read the comments to know what is going on.

IliaLarchenko / lerobot\_random

Code Issues Pull requests Actions Projects Wiki Security Insights Settings

main / lerobot\_random / via / pi / lekiwi\_policy.py

IliaLarchenko Added useful files for pi0 training and inference 4df3d2e · 20 minutes ago History

**Code** Blame 92 lines (72 loc) · 3.28 KB

```

1  # Based on https://github.com/Physical-Intelligence/openpi/src/openpi/policies/libero_policy.py
2
3  import dataclasses
4
5  import einops
6  import numpy as np
7
8  from openpi import transforms
9  from openpi.models import model as _model
10
11
12 def make_libero_example() -> dict:
13     """Creates a random input example for the Libero policy."""
14     return {
15         "observation/state": np.random.rand(9),
16         "observation/images/front": np.random.randint(256, size=(640, 360, 3), dtype=np.uint8),
17         "observation/images/wrist": np.random.randint(256, size=(640, 360, 3), dtype=np.uint8),
18         "observation/images/top": np.random.randint(256, size=(640, 360, 3), dtype=np.uint8),
19         "prompt": "do something",
20     }
21
22
23 def _parse_image(image) -> np.ndarray:
24     image = np.asarray(image)
25     if np.issubdtype(image.dtype, np.floating):
26         image = (255 * image).astype(np.uint8)
27     if image.shape[0] == 3:
28         image = einops.rearrange(image, "c h w -> h w c")
29

```

This is my own example policy file.

main / lerobot\_random / via / pi / lekiwi\_policy.py

**Code** Blame 92 lines (72 loc) · 3.28 KB

```

33 class LeKiwiInputs(transforms.DataTransformFn):
34     action_dim: int
35
36     # Determines which model will be used.
37     # Do not change this for your own dataset.
38     model_type: _model.ModelType = _model.ModelType.PI0
39
40     def __call__(self, data: dict) -> dict:
41         state = transforms.pad_to_dim(data["observation/state"], self.action_dim)
42
43         ## I have 3 cameras and it is not perfectly fit the type of cameras Pi0 uses
44         ## So I pass the front image instead of the right wrist image - it still works
45         base_image = _parse_image(data["observation/images/top"])
46         wrist_image = _parse_image(data["observation/images/wrist"])
47         front_image = _parse_image(data["observation/images/front"])
48
49         # Create inputs dict. Do not change the keys in the dict below.
50         inputs = {
51             "state": state,
52             "image": {
53                 "base_0_rgb": base_image,
54                 "left_wrist_0_rgb": wrist_image,
55                 "right_wrist_0_rgb": front_image,
56             },
57             "image_mask": {
58                 "base_0_rgb": np.True_,
59                 "left_wrist_0_rgb": np.True_,
60                 "right_wrist_0_rgb": np.True_,
61             },
62         }
63
64         # Pad actions to the model action dimension. Keep this for your own dataset.
65         # Actions are only available during training.
66         if "action" in data:
67             # We are padding to the model action dim.
68             actions = transforms.pad_to_dim(data["action"], self.action_dim)
69             inputs["actions"] = actions
70
71
72
73
74
75

```

**All Symbols**

**front\_image**

2 References Search

In this file

52 `front_image = _parse_image(data["observation/images/front"])`

60 `"right_wrist_0_rgb": front_image,`

Search for this symbol

The trickiest part is aligning the names of the 1 or 3 cameras.

Physical-Intelligence / openpi

<> Code Issues 79 Pull requests 19 Discussions Actions Projects Security Insights

main openpi / src / openpi / training / config.py

kvablock update paths 2d70d96 · 3 days ago History

**Code** Blame 981 lines (890 loc) · 43.2 KB

```

1  """See _CONFIGS for the list of available configs."""
2
3  import abc
4  from collections.abc import Sequence
5  import dataclasses
6  import difflib
7  import logging
8  import pathlib
9  from typing import Any, Literal, Protocol, TypeAlias
10
11 import etils.epath as epath
12 import flax.nnx as nnx
13 from typing_extensions import override
14 import tyro
15
16 import openpi.models.model as _model
17 import openpi.models.pi0_config as pi0_config
18 import openpi.models.pi0_fast as pi0_fast
19 import openpi.models.tokenizer as _tokenizer
20 import openpi.policies.aloha_policy as aloha_policy
21 import openpi.policies.droid_policy as droid_policy
22 import openpi.policies.libero_policy as libero_policy
23 import openpi.shared.download as _download
24 import openpi.shared.normalize as _normalize
25 import openpi.training.droid_rlds_dataset as droid_rlds_dataset
26 import openpi.training.misc.roboarena_config as roboarena_config
27 import openpi.training.optimizer as _optimizer
28 import openpi.training.weight_loaders as weight_loaders

```



You need to make some changes in this file. Read the comments to understand better

main openpi / src / openpi / training / config.py

**Code** Blame 981 lines (890 loc) · 43.2 KB

```

212 class SimpleDataConfig(DataConfigFactory):
213     def create(self, assets_dirs: pathlib.Path, model_config: _model.BaseModelConfig) -> DataConfig:
214         )
215
216
217     @dataclasses.dataclass(frozen=True)
218     class LeRobotAlohaDataConfig(DataConfigFactory):
219         # If true, will convert joint dimensions to deltas with respect to the current state before passing to the model.
220         # Gripper dimensions will remain in absolute values.
221         use_delta_joint_actions: bool = True
222
223         # If provided, will be injected into the input data if the "prompt" key is not present.
224         default_prompt: str | None = None
225
226         # If true, this will convert the joint and gripper values from the standard Aloha space to
227         # the space used by the pi internal runtime which was used to train the base model. People who
228         # use standard Aloha data should set this to true.
229         adapt_to_pi: bool = True
230
231
232         # Repack transforms.
233         repack_transforms: tyro.conf.Suppress[_transforms.Group] = dataclasses.field(
234             default=_transforms.Group(
235                 inputs=[
236                     _transforms.RepackTransform(
237                         [
238                             "images": {"cam_high": "observation.images.top"},
239                             "state": "observation.state",
240                             "actions": "action",
241                         ]
242                     )
243                 ]
244             )
245         )
246
247         # Action keys that will be used to read the action sequence from the dataset.
248         action_sequence_keys: Sequence[str] = ("action",)
249
250         @override
251         def create(self, assets_dirs: pathlib.Path, model_config: _model.BaseModelConfig) -> DataConfig:
252             data_transforms = _transforms.Group(
253                 inputs=[aloha_policy.AlohaInputs(adapt_to_pi=self.adapt_to_pi)],
254

```



main openpi / src / openpi / training / config.py ↑ Top

Code Blame 981 lines (890 loc) · 43.2 KB · ⌂

```
358 class RLSDroidDataConfig(DataConfigFactory):
359     def __init__(self, assets_dirs: pathlib.Path, model_config: _model.BaseModelConfig) -> DataConfig:
360         ...
361
362     @dataclasses.dataclass(frozen=True)
363     class LeRobotDROIDDataConfig(DataConfigFactory):
364         """
365             Example data config for custom DROID dataset in LeRobot format.
366             To convert your custom DROID dataset (<10s of hours) to LeRobot format, see examples/droid/convert_droid_data_to_lerobot.py
367         """
368
369         @override
370         def __init__(self, assets_dirs: pathlib.Path, model_config: _model.BaseModelConfig) -> DataConfig:
371             repack_transform = _transforms.Group(
372                 inputs=[
373                     _transforms.RepackTransform(
374                         {
375                             "observation/exterior_image_1_left": "exterior_image_1_left",
376                             "observation/exterior_image_2_left": "exterior_image_2_left",
377                             "observation/wrist_image_left": "wrist_image_left",
378                             "observation/joint_position": "joint_position",
379                             "observation/gripper_position": "gripper_position",
380                             "actions": "actions",
381                             "prompt": "prompt",
382                         }
383                     )
384                 ]
385             )
386
387             # We assume joint *velocity* actions, so we should *not* apply an additional delta transform.
388             data_transforms = _transforms.Group(
389                 inputs=[droid_policy.DroidInputs(model_type=model_config.model_type)],
390                 outputs=[droid_policy.DroidOutputs()],
391             )
392
393             model_transforms = ModelTransformFactory()(model_config)
394
395             return dataclasses.replace(
396                 self.create_base_config(assets_dirs, model_config),
397                 repack_transform=repack_transform,
398                 data_transforms=data_transforms,
399                 model_transforms=model_transforms
400             )
401
402     
```





main · openpi / src / openpi / training / config.py · Top

Code Blame 981 lines (890 loc) · 43.2 KB · ⌂

```
551 _CONFIGS = [
552     },
553     #
554     # Debugging configs.
555     #
556     TrainConfig(
557         name="debug",
558         data=FakeDataConfig(),
559         batch_size=2,
560         model=pi0_config.Pi0Config(pi0gemma_variant="dummy", action_expert_variant="dummy"),
561         save_interval=100,
562         overwrite=True,
563         exp_name="debug",
564         num_train_steps=10,
565         wandb_enabled=False,
566     ),
567     TrainConfig(
568         name="debug_restore",
569         data=FakeDataConfig(),
570         batch_size=2,
571         model=pi0_config.Pi0Config(pi0gemma_variant="dummy", action_expert_variant="dummy"),
572         weight_loader=weight_loaders.CheckpointWeightLoader("./checkpoints/debug/debug/9/params"),
573         overwrite=True,
574         exp_name="debug",
575         num_train_steps=10,
576         wandb_enabled=False,
577     ),
578     TrainConfig(
579         name="debug_pi05",
580         model=pi0_config.Pi0Config(pi05=True, pi0gemma_variant="dummy", action_expert_variant="dummy"),
581         data=FakeDataConfig(),
582         batch_size=2,
583         num_train_steps=10,
584         overwrite=True,
585         exp_name="debug_pi05",
586         wandb_enabled=False,
587     ),
588 ]
```



IliaLarchenko / lerobot\_random

<> Code Issues Pull requests Actions Projects Wiki Security Insights Settings

main / lerobot\_random / via / pi / config.py

IliaLarchenko Added useful files for pi0 training and inference 4df3d2e · 22 minutes ago History

**Code** Blame 170 lines (131 loc) · 6.71 KB

```

1 # This code is based on and should be added to the https://github.com/Physical-Intelligence/openpi/blob/main/src/openpi/training/config.py
2
3
4 # ...
5
6 import openpi.policies.lekiwi_policy as lekiwi_policy
7
8 # ...
9
10 @dataclasses.dataclass(frozen=True)
11 class LeRobotLeKiwiDataConfig(DataConfigFactory):
12     """
13     This config is used to configure transforms that are applied at various parts of the data pipeline.
14     """
15
16     action_sequence_keys: Sequence[str] = ("action",)
17
18     @override
19     def create(self, assets_dirs: pathlib.Path, model_config: _model.BaseModelConfig) -> DataConfig:
20         """
21         # The repack transform simply remaps key names here.
22         # I keep it here as an example and to align naming with the original notation
23         repack_transform = _transforms.Group(
24             inputs=[
25                 _transforms.RepackTransform(
26                     {
27                         "observation/images/top": "observation.images.top",
28                         "observation/images/wrist": "observation.images.wrist",
29                         "observation/images/front": "observation.images.front",
30                         "observation/state": "observation.state",
31                         "action": "action",
32                         "prompt": "prompt",
33                     }
34                 )
35             ]
36
37             # The data transforms are applied to the data coming from the dataset *and* during inference.
38             data_transforms = _transforms.Group(
39                 inputs=[lekiwi_policy.LeKiwiInputs(action_dim=model_config.action_dim, model_type=model_config.model_type)],
40                 outputs=[lekiwi_policy.LeKiwiOutputs()],
41             )
42
43             # My LeKiwi actions dim is 9, first 5 are joints that should be converted to delta actions.
44             # 6th is gripper that should be left unchanged.
45             # 7-9 are mobile based velocities and they are 0 in my case anyways.
46             delta_action_mask = _transforms.make_bool_mask(5, -4)
47             data_transforms = data_transforms.push(
48                 inputs=[_transforms.DeltaActions(delta_action_mask)],

```

All Symbols

lekiwi\_policy

4 References Search

In this file

6 lekiwi\_policy as lekiwi\_policy

6 as lekiwi\_policy

39 inputs=[lekiwi\_policy.LeKiwiInputs(action\_dim=model\_config.action\_dim, model\_type=model\_config.model\_type)],

40 outputs=[lekiwi\_policy.LeKiwiOutputs()],

Search for this symbol



This is my example config file used for the training

main / lerobot\_random / via / pi / config.py

**Code** Blame 170 lines (131 loc) · 6.71 KB

```

11 class LeRobotLeKiwiDataConfig(DataConfigFactory):
12
13     """
14     This config is used to configure transforms that are applied at various parts of the data pipeline.
15     """
16
17     action_sequence_keys: Sequence[str] = ("action",)
18
19     @override
20     def create(self, assets_dirs: pathlib.Path, model_config: _model.BaseModelConfig) -> DataConfig:
21         """
22         # The repack transform simply remaps key names here.
23         # I keep it here as an example and to align naming with the original notation
24         repack_transform = _transforms.Group(
25             inputs=[
26                 _transforms.RepackTransform(
27                     {
28                         "observation/images/top": "observation.images.top",
29                         "observation/images/wrist": "observation.images.wrist",
30                         "observation/images/front": "observation.images.front",
31                         "observation/state": "observation.state",
32                         "action": "action",
33                         "prompt": "prompt",
34                     }
35                 )
36             ]
37
38             # The data transforms are applied to the data coming from the dataset *and* during inference.
39             data_transforms = _transforms.Group(
40                 inputs=[lekiwi_policy.LeKiwiInputs(action_dim=model_config.action_dim, model_type=model_config.model_type)],
41                 outputs=[lekiwi_policy.LeKiwiOutputs()],
42             )
43
44             # My LeKiwi actions dim is 9, first 5 are joints that should be converted to delta actions.
45             # 6th is gripper that should be left unchanged.
46             # 7-9 are mobile based velocities and they are 0 in my case anyways.
47             delta_action_mask = _transforms.make_bool_mask(5, -4)
48             data_transforms = data_transforms.push(
49                 inputs=[_transforms.DeltaActions(delta_action_mask)],

```



Align the naming of the inputs and outputs in your dataset.

Code Blame 170 lines (131 loc) - 6.71 KB

```

11     class LeRobotLeKiwiDataConfig(DataConfigFactory):
12         def create(self, assets_dirs: pathlib.Path, model_config: _model.BaseModelConfig) -> DataConfig:
13             "observation/state": "observation.state",
14             "action": "action",
15             "prompt": "prompt",
16         )
17     ]
18   )
19   ]
20
21   # The data transforms are applied to the data coming from the dataset *and* during inference.
22   data_transforms = _transforms.Group(
23       inputs=[lekiwi_policy.LeKiwiIInputs(action_dim=model_config.action_dim, model_type=model_config.model_type)],
24       outputs=[lekiwi_policy.LeKiwiOOutputs()],
25   )
26
27   # My LeKiwi actions dim is 9, first 5 are joints that should be converted to delta actions.
28   # 6th is gripper that should be left unchanged.
29   # 7-9 are mobile based velocities and they are 0 in my case anyways.
30   delta_action_mask = _transforms.make_bool_mask(5, -4)
31   data_transforms = data_transforms.push(
32       inputs=[_transforms.DeltaActions(delta_action_mask)],
33       outputs=[_transforms.AbsoluteActions(delta_action_mask)],
34   )
35
36   # Model transforms include things like tokenizing the prompt and action targets
37   # You do not need to change anything here for your own dataset.
38   model_transforms = ModelTransformFactory()(model_config)
39
40   # We return all data transforms for training and inference. No need to change anything here.
41   return dataclasses.replace(
42       self.create_base_config(assets_dirs, model_config),
43       repack_transforms=repack_transform,
44       data_transforms=data_transforms,
45       model_transforms=model_transforms,
46       action_sequence_keys=self.action_sequence_keys,
47   )
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88

```



Transform the joint actions into delta actions.

Code Blame 170 lines (131 loc) - 6.71 KB

```

11     class LeRobotLeKiwiDataConfig(DataConfigFactory):
12         def create(self, assets_dirs: pathlib.Path, model_config: _model.BaseModelConfig) -> DataConfig:
13             # You do not need to change anything here for your own dataset.
14             model_transforms = ModelTransformFactory()(model_config)
15
16             # We return all data transforms for training and inference. No need to change anything here.
17             return dataclasses.replace(
18                 self.create_base_config(assets_dirs, model_config),
19                 repack_transforms=repack_transform,
20                 data_transforms=data_transforms,
21                 model_transforms=model_transforms,
22                 action_sequence_keys=self.action_sequence_keys,
23             )
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
39
40
41
42
43
44
45
46
47
48
49
49
50
51
52
53
54
55
56
57
58
59
59
60
61
62
63
64
65
66
66
67
68
69
70
71
72
73
74
75
76
77
78
79
79
80
81
82
83
84
85
86
87
88

```



Finally, you need to create your training config. I created multiple configs in my file above.

Code Blame 178 lines (131 loc) · 6.71 KB

```

63     )
64
65
66     _CONFIGS = [
67         # ...
68         TrainConfig(
69             # Change the name to reflect your model and dataset.
70             name="pi0_lekiwi_lora",
71             model=pi0_config.Pi0Config(paligemma_variant="gemma_2b_lora", action_expert_variant="gemma_300m_lora", action_horizon=30),
72
73             data=LeRobotLeKiwiDataConfig(
74                 repo_id="IliaLarchenko/vla_demo",
75                 base_config>DataConfig(prompt_from_task=True),
76             ),
77
78             weight_loader=weight_loaders.CheckpointWeightLoader("gs://openpi-assets/checkpoints/pi0_base/params"),
79
80             # In my case, 5k steps was enough
81             num_train_steps=5_000,
82
83             freeze_filter=pi0_config.Pi0Config(
84                 paligemma_variant="gemma_2b_lora", action_expert_variant="gemma_300m_lora", action_horizon=30
85             ).get_freeze_filter(),
86
87             # Turn off EMA for LoRA finetuning.
88             ema_decay=None,
89         ),
90
91         TrainConfig(
92             # Change the name to reflect your model and dataset.
93             name="pi0_lekiwi",
94             model=pi0_config.Pi0Config(paligemma_variant="gemma_2b", action_expert_variant="gemma_300m", action_horizon=30),
95
96             data=LeRobotLeKiwiDataConfig(
97                 repo_id="IliaLarchenko/vla_demo",
98                 base_config>DataConfig(prompt_from_task=True),
99             ),
100
101

```



[README](#) [Contributing](#) [Apache-2.0 license](#)

## 2. Defining training configs and running training

To fine-tune a base model on your own data, you need to define configs for data processing and training. We provide example configs with detailed comments for LIBERO below, which you can modify for your own dataset:

- [LiberoInputs and LiberoOutputs](#) : Defines the data mapping from the LIBERO environment to the model and vice versa. Will be used for both, training and inference.
- [LeRobotLiberoDataConfig](#) : Defines how to process raw LIBERO data from LeRobot dataset for training.
- [TrainConfig](#) : Defines fine-tuning hyperparameters, data config, and weight loader.

We provide example fine-tuning configs for [π<sub>θ</sub>](#), [π<sub>θ</sub>-FAST](#), and [π<sub>θ</sub>-S](#) on LIBERO data.

Before we can run training, we need to compute the normalization statistics for the training data. Run the script below with the name of your training config:

```
uv run scripts/compute_norm_stats.py --config-name pi05_libero
```

Now we can kick off training with the following command (the `--overwrite` flag is used to overwrite existing checkpoints if you rerun fine-tuning with the same config):

```
XLA_PYTHON_CLIENT_MEM_FRACTION=0.9 uv run scripts/train.py pi05_libero --exp-name=my_experimen
```

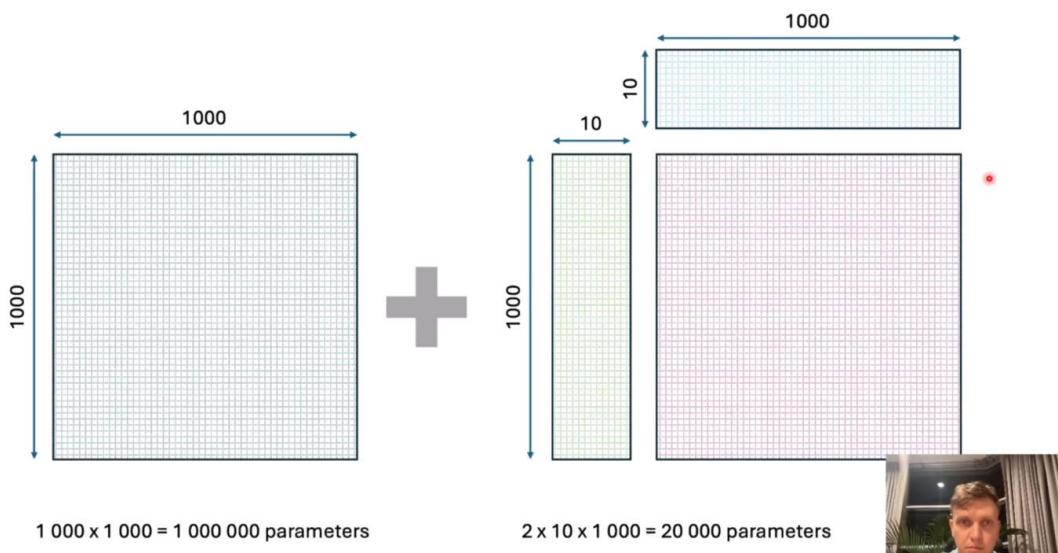
The command will log training progress to the console and save checkpoints to the `checkpoints` directory. You can also monitor training progress on the Weights & Biases dashboard. For maximally using the GPU memory, set `XLA_PYTHON_CLIENT_MEM_FRACTION=0.9` before running training -- this enables JAX to use up to 90% of the GPU memory (vs. the default of 75%).

**Note:** We provide functionality for *reloading* normalization statistics for state / action normalization from pre-training. This can be beneficial if you are fine-tuning to a new task on a robot that was part of our pre-training mixture. For more details on how to reload normalization statistics, see the [norm\\_stats.md](#) file.

## 3. Spinning up a policy server and running inference



Next, let us see how LoRA finetuning works below



[README](#) [Contributing](#) [Apache-2.0 license](#)

## Running Inference for a Pre-Trained Model

Our pre-trained model checkpoints can be run with a few lines of code (here our  $\pi_0$ -FAST-DROID model):

```
from openpi.training import config as _config
from openpi.policies import policy_config
from openpi.shared import download

config = _config.get_config("pi05_droid")
checkpoint_dir = download.maybe_download("gs://openpi-assets/checkpoints/pi05_droid")

# Create a trained policy.
policy = policy_config.create_trained_policy(config, checkpoint_dir)

# Run inference on a dummy example.
example = {
    "observation/exterior_image_1_left": ...,
    "observation/wrist_image_left": ...,
    ...
    "prompt": "pick up the fork"
}
action_chunk = policy.infer(example)["actions"]
```

You can also test this out in the [example notebook](#).

We provide detailed step-by-step examples for running inference of our pre-trained checkpoints on [DROID](#) and [ALOHA](#) robots.

**Remote Inference:** We provide [examples and code](#) for running inference of our models **remotely**: the model can run on a different server and stream actions to the robot via a websocket connection. This makes it easy to use more powerful GPUs off-robot and keep robot and policy environments separate.

**Test inference without a robot:** We provide a [script](#) for testing inference without a robot. This script will generate a random observation and run inference with the model. See [here](#) for more details.



After the model is finetuned, you can now start running inference on it

IliaLarchenko / lerobot\_random

Code Issues Pull requests Actions Projects Wiki Security Insights Settings

main / lerobot\_random / via / pi / evaluate\_pi0.py

Go to file

IliaLarchenko Added useful files for pi0 training and inference

4df3d2e · 34 minutes ago History

Code Blame 108 lines (83 loc) · 3.7 KB

Raw

```

1 #!/usr/bin/env python3
2
3 # Adapted from https://github.com/huggingface/lerobot/blob/main/examples/lekiwi/evaluate.py
4 # And other examples from LeRobot repo
5 # LeRobot and OpenPi should be installed
6 # Run 'python evaluate_pi0.py' to run the model
7
8 import os
9 # Configure JAX memory allocation before importing JAX-related modules
10 # Is needed to resolve possible issues with memory allocation
11 os.environ["XLA_PYTHON_CLIENT_MEM_FRACTION"] = "0.9"
12
13 import time
14 import numpy as np
15
16 from lerobot.robots.lekiwi import LeKiwiClient, LeKiwiClientConfig
17 from lerobot.utils.robot_utils import busy_wait
18
19 # Import Pi0 model from openpi
20 from openpi.training import config as pi0_config
21 from openpi.policies import policy_config
22 from huggingface_hub import snapshot_download
23
24 # Configuration
25 FPS = 30
26 TASK_DESCRIPTION = "Pick the orange ball and put it in the green cup"
27 ACTIONS_TO_EXECUTE = 20 # Execute this many actions from each predicted chunk
28

```



main / lerobot\_random / via / pi / evaluate\_pi0.py

Top

Code Blame 108 lines (83 loc) · 3.7 KB

Raw

```

22 from huggingface_hub import snapshot_download
23
24 # Configuration
25 FPS = 30
26 TASK_DESCRIPTION = "Pick the orange ball and put it in the green cup"
27 ACTIONS_TO_EXECUTE = 20 # Execute this many actions from each predicted chunk
28
29 # Load Pi0 model
30 print("Loading Pi0 model...")
31
32 # Use your robot Pi0 config - in my case "pi0_lekiwi_fast", "pi0_lekiwi" or "pi05_lekiwi"
33 config = pi0_config.get_config("pi0_lekiwi")
34
35 # Use your trained policy HF directory
36 # You should upload your model (assets and params directories from the checkpoint) to Hugging Face to use it here
37 checkpoint_dir = snapshot_download(repo_id="YOUR_HF_REPO_PATH")
38 pi0_policy = policy_config.create_trained_policy(config, checkpoint_dir)
39 print("Pi0 model loaded successfully")
40
41 # Connect to robot
42 # Use your robot IP and ID
43 robot_config = LeKiwiClientConfig(remote_ip="YOUR_ROBOT_IP", id="YOUR_ROBOT_ID")
44 robot = LeKiwiClient(robot_config)
45 robot.connect()
46
47 if not robot.is_connected:
48     raise ValueError("Robot is not connected!")
49
50 print(f"Robot connected. Starting control loop with task: '{TASK_DESCRIPTION}'")
51 print(f"Will execute {ACTIONS_TO_EXECUTE} actions from each predicted chunk")
52
53 # Control loop variables
54 step = 0
55 last_actions = None
56 action_index = 0
57 pred_times = []
58
59 while True:

```



main ▾ lerobot\_random / via / pi / evaluate\_pi0.py ↑ Top

Code Blame 108 lines (83 loc) · 3.7 KB

```
51 print(f"Will execute {ACTIONS_TO_EXECUTE} actions from each predicted chunk")
52
53 # Control loop variables
54 step = 0
55 last_actions = None
56 action_index = 0
57 pred_times = []
58
59 while True:
60     t0 = time.perf_counter()
61
62     # Run prediction when we need new actions (either first time or when we've executed enough actions)
63     if last_actions is None or action_index >= min(ACTIONS_TO_EXECUTE, len(last_actions)):
64         # Get robot observation
65         observation = robot.get_observation()
66
67         # Prepare input for Pi0
68         # Adjust for your robot
69         pi0_input = {
70             "prompt": TASK_DESCRIPTION,
71             "observation/state": observation["observation.state"],
72         }
73
74         for k in ["top", "wrist", "front"]:
75             pi0_input[f"observation/images/{k}"] = observation[k]
76
77         # Run inference
78         t_pred_start = time.perf_counter()
79         output = pi0_policy.infer(pi0_input)
80         t_pred = time.perf_counter() - t_pred_start
81
82         # Keep track of last 10 prediction times
83         pred_times.append(t_pred)
84         pred_times = pred_times[-10:] # Keep last 10
85
86         last_actions = output["actions"]
87         action_index = 0
88
89     print(f"Step {step}: Predicted {len(last_actions)} actions, will execute {min(ACTIONS_TO_EXECUTE, len(last_actions))}")
90     print(f"Prediction took {t_pred:.3f}s (avg over last {len(pred_times)}: {np.mean(pred_times):.3f}s)")
91
92     # Execute action
93     else:
94         # Get current action from the sequence
95         action = last_actions[action_index]
96
97         # Convert action array to robot's expected format
98         action_dict = {}
99         for i, action_name in enumerate(list(robot.action_features.keys())):
100             if i < len(action):
101                 action_dict[action_name] = action[i]
102
103         robot.send_action(action_dict)
104         action_index += 1
105         step += 1
106
107     # Maintain FPS
108     busy_wait(max(1.0 / FPS - (time.perf_counter() - t0), 0.0))
```



main ▾ lerobot\_random / via / pi / evaluate\_pi0.py ↑ Top

Code Blame 108 lines (83 loc) · 3.7 KB

```
73
74     for k in ["top", "wrist", "front"]:
75         pi0_input[f"observation/images/{k}"] = observation[k]
76
77     # Run inference
78     t_pred_start = time.perf_counter()
79     output = pi0_policy.infer(pi0_input)
80     t_pred = time.perf_counter() - t_pred_start
81
82     # Keep track of last 10 prediction times
83     pred_times.append(t_pred)
84     pred_times = pred_times[-10:] # Keep last 10
85
86     last_actions = output["actions"]
87     action_index = 0
88
89     print(f"Step {step}: Predicted {len(last_actions)} actions, will execute {min(ACTIONS_TO_EXECUTE, len(last_actions))}")
90     print(f"Prediction took {t_pred:.3f}s (avg over last {len(pred_times)}: {np.mean(pred_times):.3f}s)")
91
92     # Execute action
93     else:
94         # Get current action from the sequence
95         action = last_actions[action_index]
96
97         # Convert action array to robot's expected format
98         action_dict = {}
99         for i, action_name in enumerate(list(robot.action_features.keys())):
100             if i < len(action):
101                 action_dict[action_name] = action[i]
102
103         robot.send_action(action_dict)
104         action_index += 1
105         step += 1
106
107     # Maintain FPS
108     busy_wait(max(1.0 / FPS - (time.perf_counter() - t0), 0.0))
```



That's it.



*After post-training, the robot can unload the dryer, bring the clothes over to the table, and fold the clothes into a stack. The video is uncut, from a single policy operating fully autonomously.*



**Laundry.** We fine-tuned  $\pi_0$  to fold laundry, using either a mobile robot or a fixed pair of arms. The goal is to get the clothing into a neat stack. This task is exceptionally difficult for robots (...and some humans): while a single t-shirt laid flat on the table can sometimes be folded just by repeating a pre-specified set of motions, a pile of tangled laundry can be crumpled in many different ways, so it is not enough to simply move the arms through the same motion. To our knowledge, no prior robot system has been demonstrated to perform this task at this level of complexity.



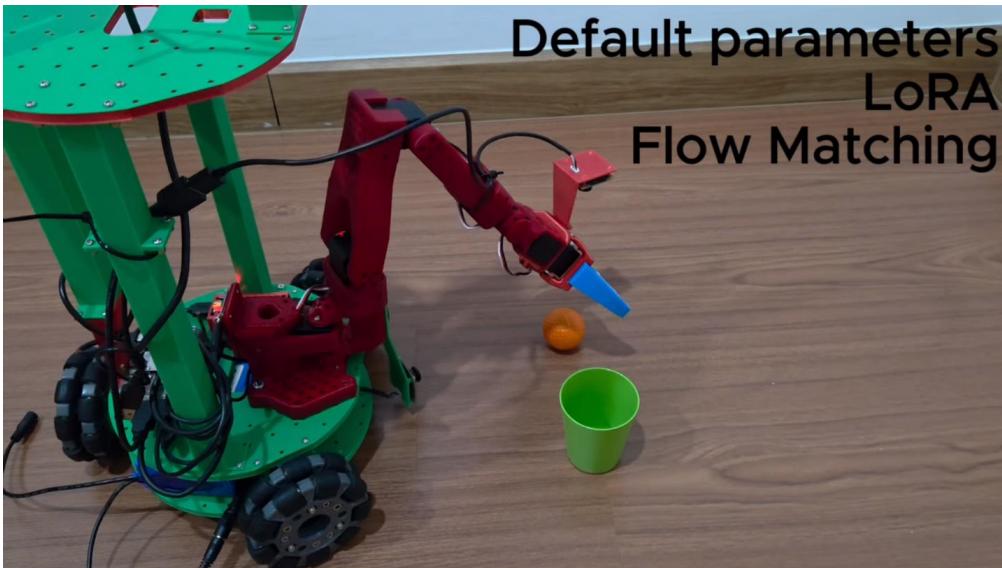
Notably, by training on diverse data, we find that the robot is able to recover when someone tries to intervene in a variety of different ways.



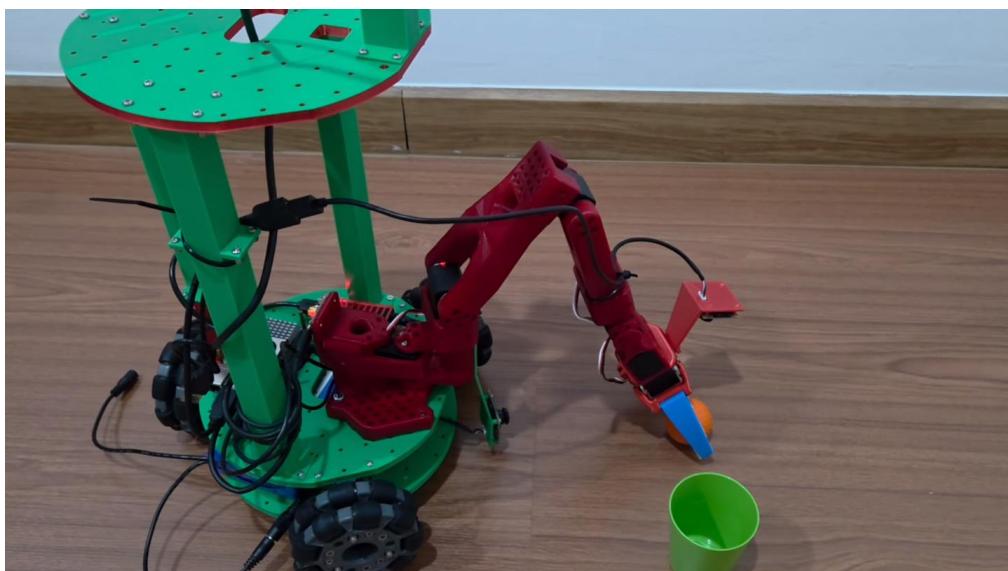
**Table bussing.** We also fine-tuned the model to bus a table. This requires the robot to pick up dishes and trash on the table, putting any dishes, cutlery, or cups into a bussing bin, and putting trash into the trash bin. This task requires the robot to handle a dizzying variety of items. One of the exciting consequences of training  $\pi_0$  on large and diverse datasets was the range of emergent strategies that the robot employed: instead of simply grasping each item in turn, the model could stack multiple dishes to put them into the bin together, or shake off trash from a plate into the garbage before placing the plate into the bussing bin.



Default parameters  
LoRA  
Flow Matching

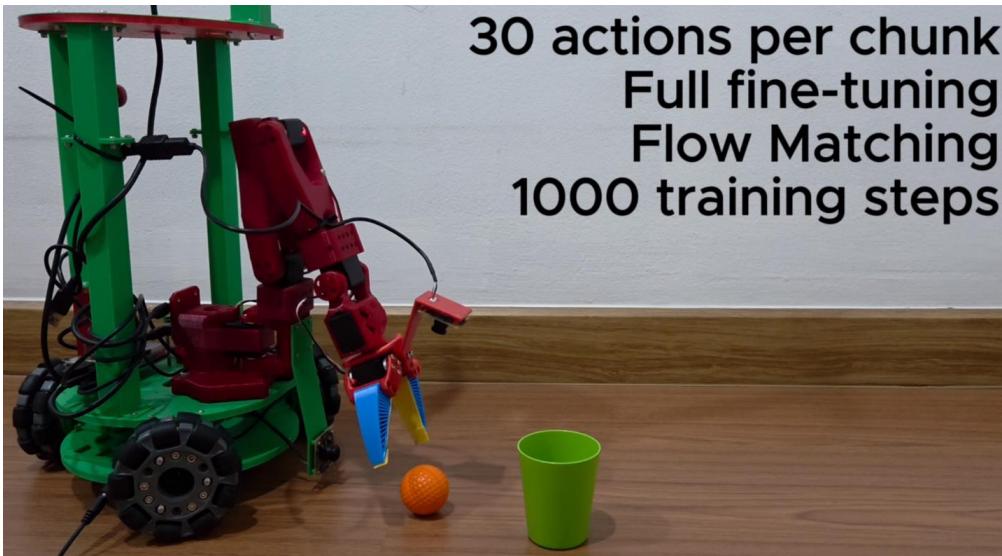


The issue of accuracy is likely the default predicted actions of 50 steps pre chunk



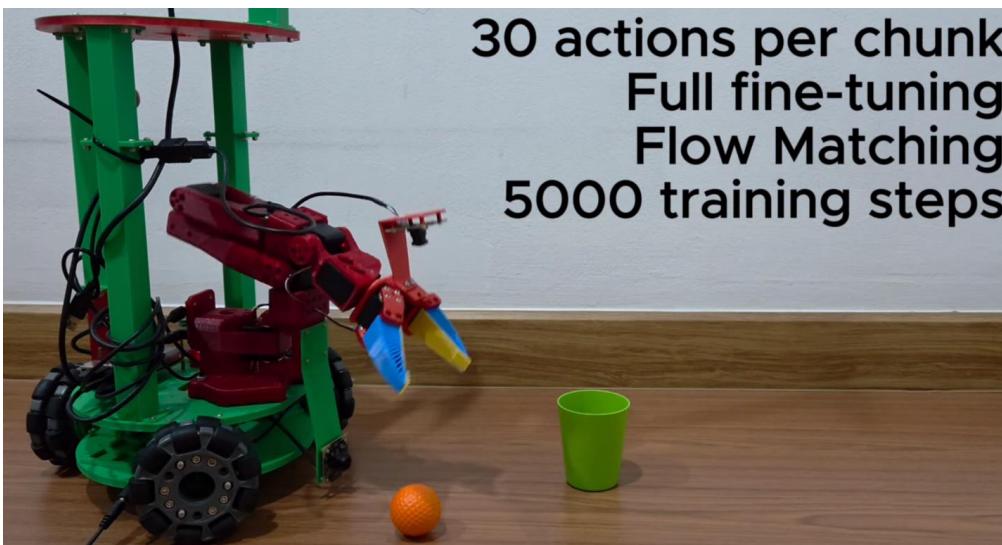
This is better now using 20 predicted chunks for the same model; it completes the tasks correctly now. The wheels are even now moving correctly by transforming the velocity actions into wheel movements.

**30 actions per chunk  
Full fine-tuning  
Flow Matching  
1000 training steps**

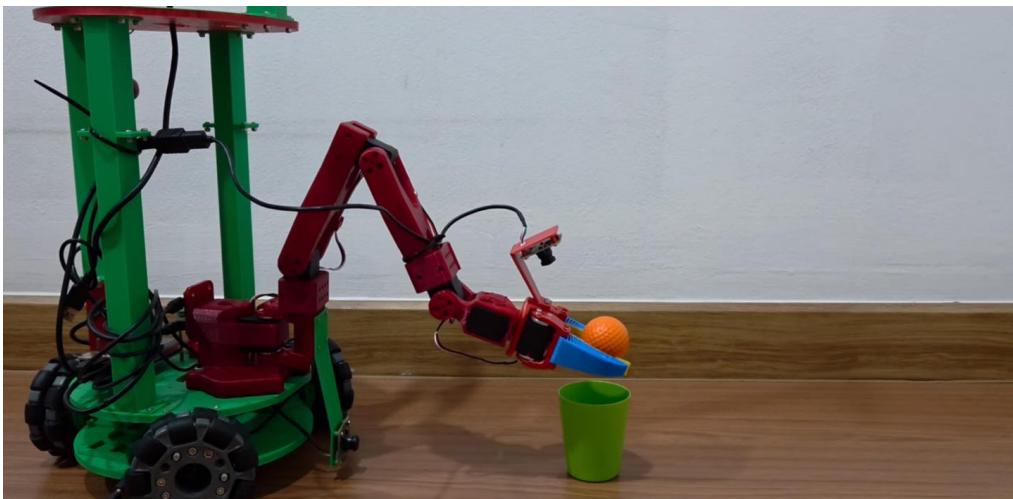


This model achieved the best result and the model converges very fast in training. Based on my inference computer GPU, it currently takes 800ms to make a prediction that contains the 20 next actions, it then runs 20 actions that were just predicted in about 700ms. So, it takes about 1.5s for the robot to execute the predicted actions. A faster GPU inference hardware might speed things up.

**30 actions per chunk  
Full fine-tuning  
Flow Matching  
5000 training steps**



Training longer results in better performance up to a threshold. It recovers from failure and acts faster too.



It predicts 30 actions at a time but the robot executes 20 of the actions per cycle of 1.5s.



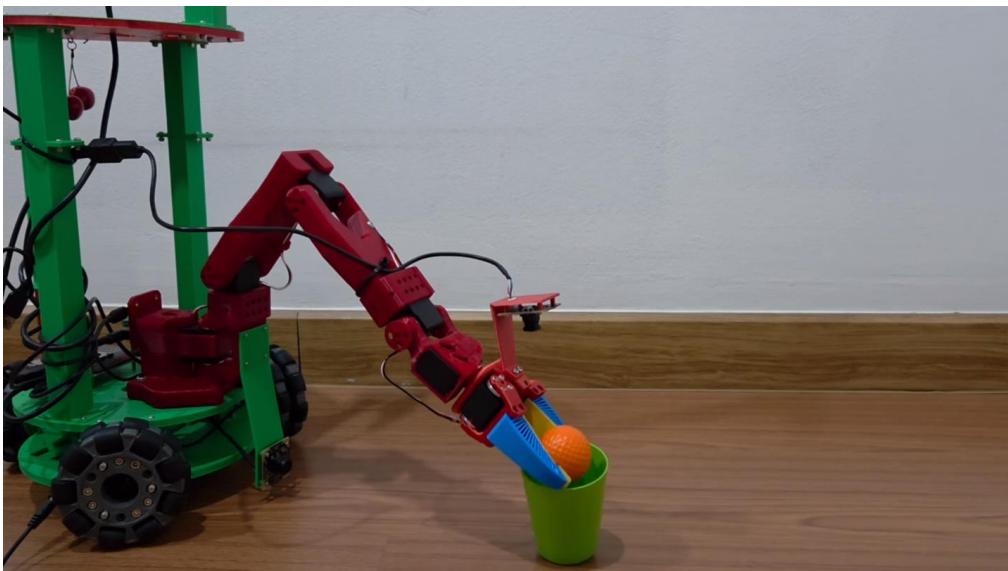
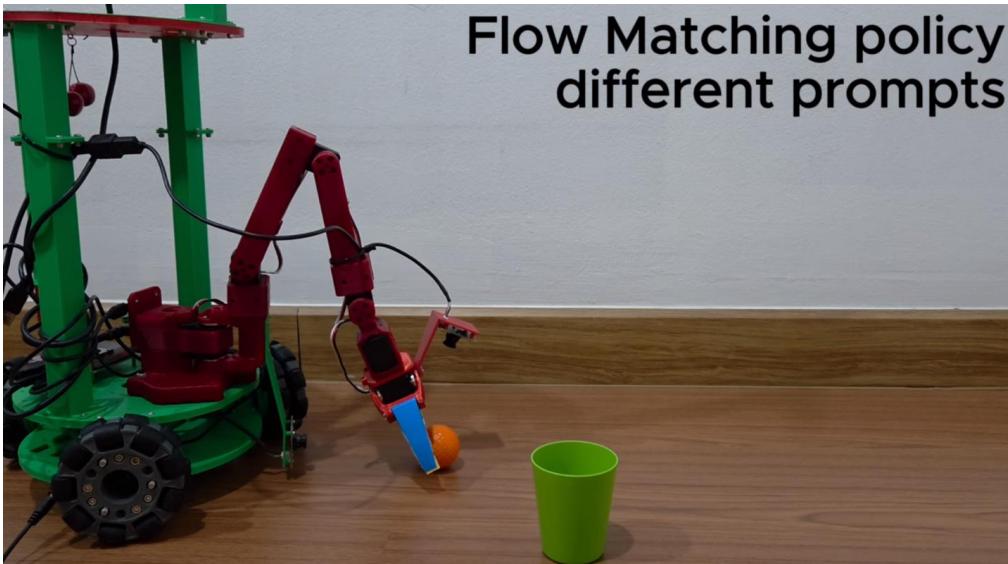
**30 actions per chunk  
Full fine-tuning  
FAST  
5000 training steps**

Instead of using Flow Matching, you can train PiO with a FAST tokenizer that will predict action tokens autoregressively.



FAST has slower inference time of about 1.5s, which is slower than previous Flow Matching, FAST has faster training convergence in the checkpoints and training steps needed, it also has a fewer success rate. LoRA + FAST didn't work.

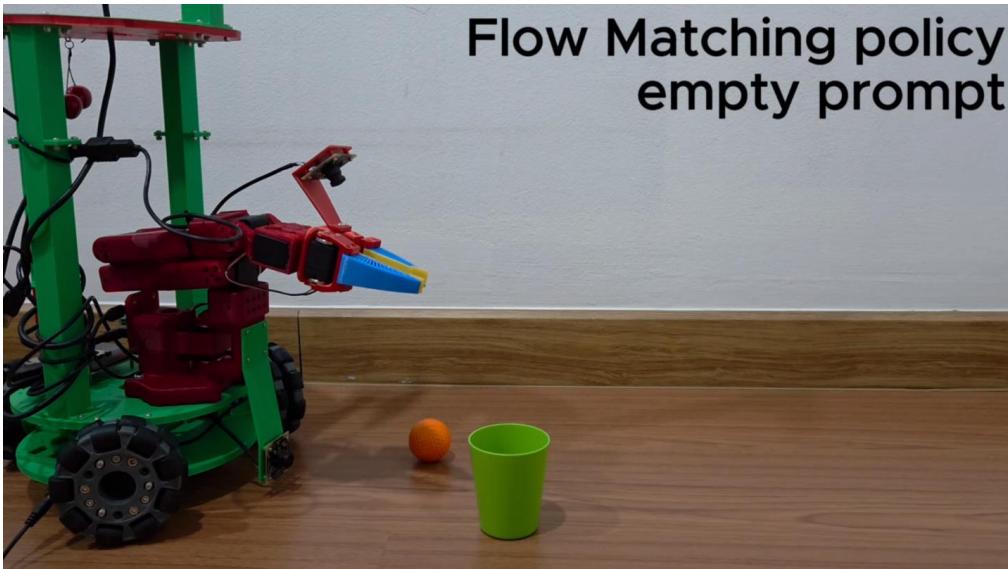
## Flow Matching policy different prompts



We can give text command to the robot, here we give it a wrong prompt to grab the cup instead. The robot just keeps picking up the ball and dropping in the cup.

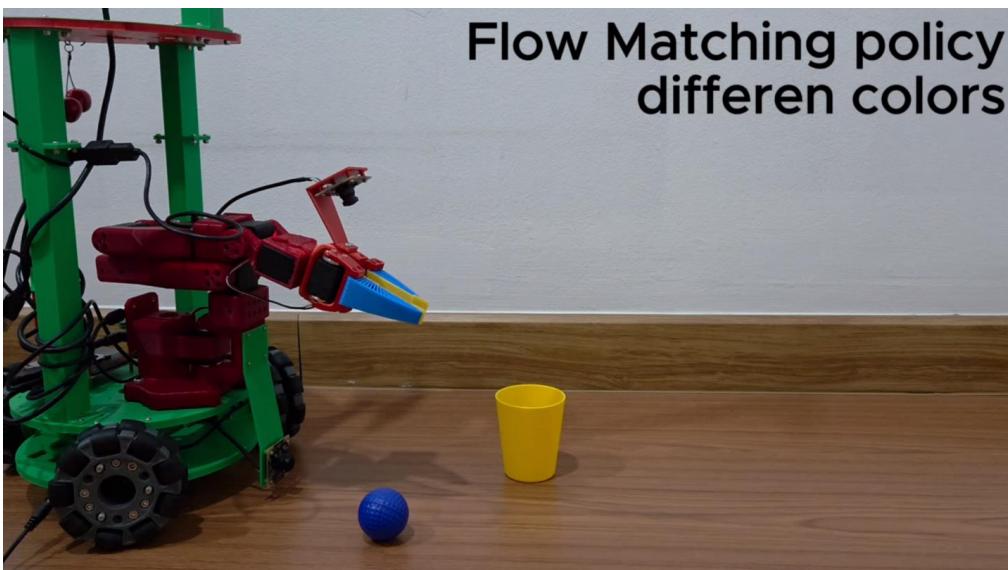


## Flow Matching policy empty prompt

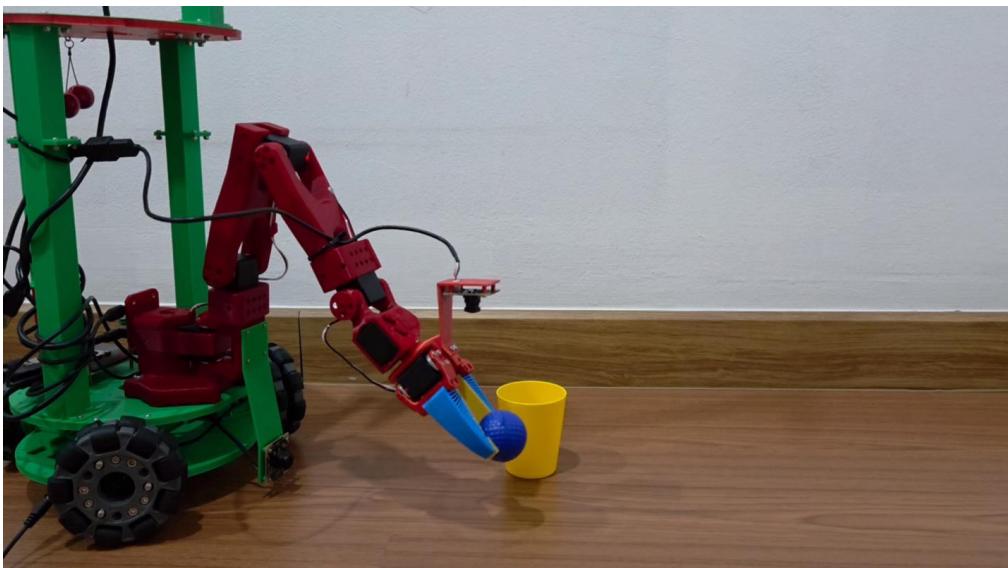


Now I left the prompt empty but the robot keeps grabbing the ball and putting it in the cup, displaying no command following. The prompt is just ignored.

## Flow Matching policy differen colors



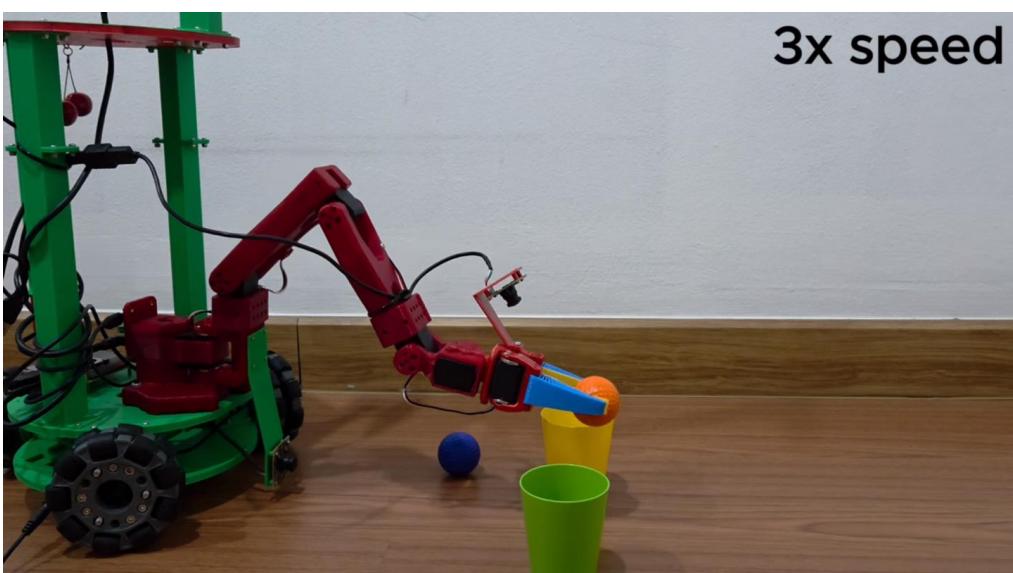
The robot is capable of working with different color balls.





Robot becomes indecisive when giving 2 cups sometimes because it was trained with a few small, single cup samples.

**3x speed**



The ball sometimes tries to pick up fathom balls



It works fine in a different environment; the policy is robust enough.

